

Homework1

February 7, 2026

0.1 K-Means

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

class KMeans:
    def __init__(self, init):
        self.data = np.array([(1, 2), (3, 3), (2, 2), (8, 8), (6, 6), (7, 7),
        ↵(-3, -3), (-2, -4), (-7, -7)])
        self.init = init
        self.old_centroids = []
        self.new_centroids = init.copy()
        self.group = {}

    def cluster(self, display=True):
        i = 0
        while (self.new_centroids != self.old_centroids):
            self.old_centroids = self.new_centroids

            centroids = np.array(self.old_centroids)
            self.group = {}
            for each_data in self.data:
                # Vectorization Optimized
                point = np.array(each_data)
                dist = np.sqrt(np.sum((centroids - point) ** 2, axis=1))
                group_of_each_data = np.argmin(dist)

                if (group_of_each_data not in self.group):
                    self.group[group_of_each_data] = []
                self.group[group_of_each_data].append(each_data)

            # Traditional Way
            # min_distance = 1e9
            # group_of_each_data = -1
            # for i in range(len(self.old_centroids)):
            #     current_distance = ( (self.old_centroids[i][0] -
            ↵each_data[0]) ** 2 + (self.old_centroids[i][1] - each_data[1]) ** 2 ) ** 0.5
```

```

        #     if (current_distance < min_distance):
        #         group_of_each_data = i
        #         min_distance = current_distance

        # if (group_of_each_data not in self.group):
        #     self.group[group_of_each_data] = []

    # self.group[group_of_each_data].append(each_data)

    self.new_centroids = []
    for each in self.group.keys():
        centroid_x = sum(t[0] for t in self.group[each]) / len(self.
        ↪group[each])
        centroid_y = sum(t[1] for t in self.group[each]) / len(self.
        ↪group[each])
        self.new_centroids.append((centroid_x, centroid_y))

    i += 1

    if (display):
        print(f"--- Round {i} ---")
        print(f"Centroids: {self.new_centroids}")

    if (display):
        for cluster_id, points in self.group.items():
            xs, ys = zip(*points)
            plt.scatter(xs, ys, label=f"Cluster {cluster_id}")

    plt.xlabel("X")
    plt.ylabel("Y")
    plt.show()

```

0.1.1 T5. starting points are (3,3), (2,2), and (-3,-3).

[2]: init_t5 = [(3, 3), (2, 2), (-3, -3)]

```
km_t5 = KMeans(init_t5)
km_t5.cluster()
```

```

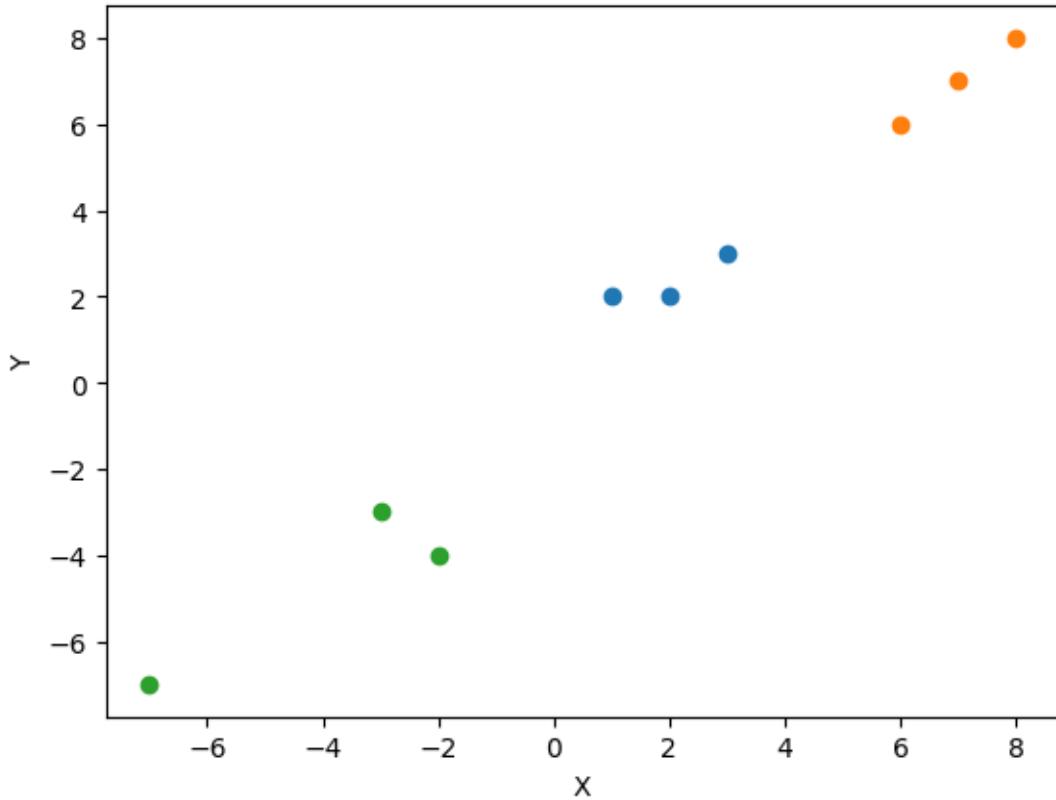
--- Round 1 ---
Centroids: [(np.float64(1.5), np.float64(2.0)), (np.float64(6.0),
np.float64(6.0)), (np.float64(-4.0), np.float64(-4.66666666666667))]

--- Round 2 ---
Centroids: [(np.float64(2.0), np.float64(2.33333333333335)), (np.float64(7.0),
np.float64(7.0)), (np.float64(-4.0), np.float64(-4.66666666666667))]

--- Round 3 ---
Centroids: [(np.float64(2.0), np.float64(2.33333333333335)), (np.float64(7.0),

```

```
np.float64(7.0)), (np.float64(-4.0), np.float64(-4.66666666666667))]
```



0.1.2 T6. starting points are (-3,-3), (2,2), and (-7,-7),

```
[3]: init_t6 = [(-3, -3), (2, 2), (-7, -7)]
```

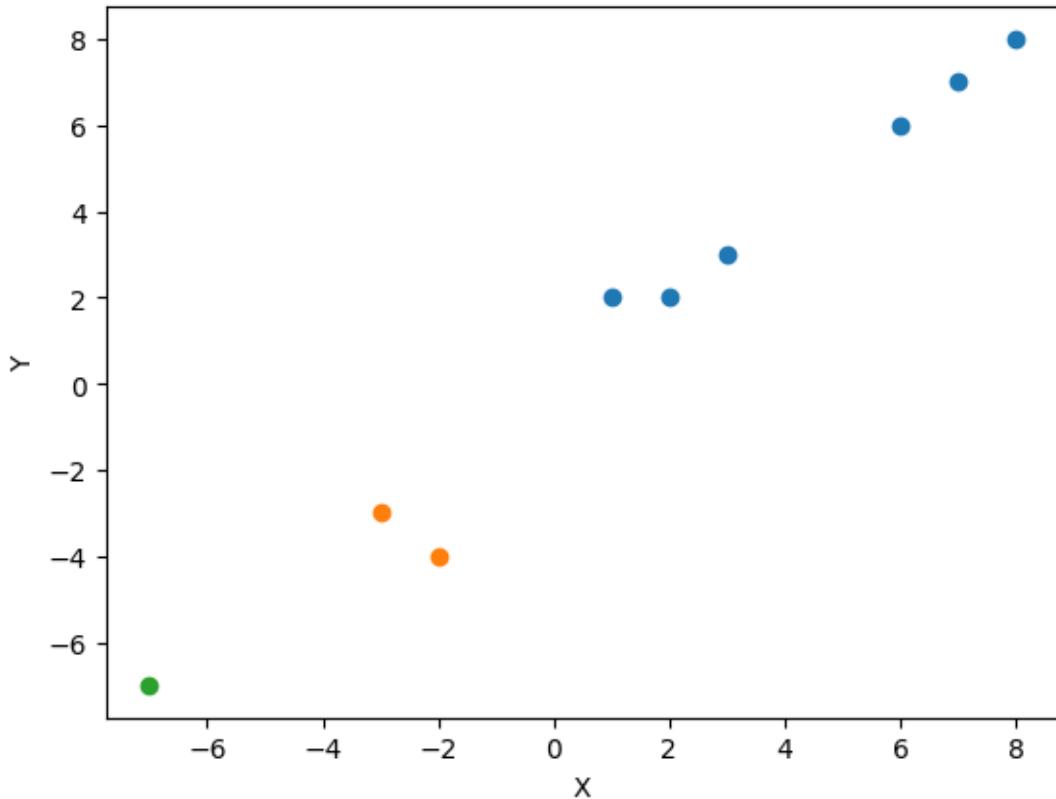
```
km_t6 = KMeans(init_t6)
km_t6.cluster()
```

```
--- Round 1 ---
```

```
Centroids: [(np.float64(4.5), np.float64(4.66666666666667)), (np.float64(-2.5),
np.float64(-3.5)), (np.float64(-7.0), np.float64(-7.0))]
```

```
--- Round 2 ---
```

```
Centroids: [(np.float64(4.5), np.float64(4.66666666666667)), (np.float64(-2.5),
np.float64(-3.5)), (np.float64(-7.0), np.float64(-7.0))]
```



- 0.1.3 T7. Between the two starting set of points in the previous two questions, which one do you think is better? How would you measure the 'goodness' quality of a set of starting points?
- 0.1.4 In general, it is important to try different sets of starting points when doing k-means.

For me, the better one is the first one.
 The 'goodness' may measure with the variance
 in the group

- 0.1.5 OT2. What would be the best K for this question? Describe your reasoning.

For me, it's 4 because the distance between (-7, -7)
 to (-3, -3) and (-2, -4) is too large to be in the
 same group. The distances are estimated equal to
 distances between centroids of each group.

0.2 Regression

```
[4]: train_url = "http://s3.amazonaws.com/assets.datacamp.com/course/Kaggle/train.csv"
train = pd.read_csv(train_url)

test_url = "http://s3.amazonaws.com/assets.datacamp.com/course/Kaggle/test.csv"
test = pd.read_csv(test_url)
```

```
[5]: train.describe()
```

```
[5]:    PassengerId  Survived  Pclass   Age  SibSp \
count  891.000000  891.000000  891.000000  714.000000  891.000000
mean   446.000000  0.383838   2.308642  29.699118  0.523008
std    257.353842  0.486592   0.836071  14.526497  1.102743
min    1.000000  0.000000   1.000000  0.420000  0.000000
25%   223.500000  0.000000   2.000000  20.125000  0.000000
50%   446.000000  0.000000   3.000000  28.000000  0.000000
75%   668.500000  1.000000   3.000000  38.000000  1.000000
max   891.000000  1.000000   3.000000  80.000000  8.000000

          Parch      Fare
count  891.000000  891.000000
mean   0.381594   32.204208
std    0.806057   49.693429
min    0.000000   0.000000
25%   0.000000   7.910400
50%   0.000000  14.454200
75%   0.000000  31.000000
max   6.000000  512.329200
```

0.2.1 T8. Median of the training datasets?

```
[6]: train["Age"].median()
```

```
[6]: np.float64(28.0)
```

28

0.2.2 T9. Mode of Embarked

```
[7]: train["Embarked"].mode().iloc[0]
```

```
[7]: 'S'
```

"S" which is 0

0.2.3 T10.

```
[8]: test
```

```
[8]:      PassengerId  Pclass          Name \
0            892       3   Kelly, Mr. James
1            893       3  Wilkes, Mrs. James (Ellen Needs)
2            894       2   Myles, Mr. Thomas Francis
3            895       3   Wirz, Mr. Albert
4            896       3  Hirvonen, Mrs. Alexander (Helga E Lindqvist)
..           ...
413           ...     ...  Spector, Mr. Woolf
414           1305     3  Oliva y Ocana, Dona. Fermina
415           1307     3  Saether, Mr. Simon Sivertsen
416           1308     3   Ware, Mr. Frederick
417           1309     3  Peter, Master. Michael J

      Sex   Age  SibSp  Parch      Ticket     Fare Cabin Embarked
0  male  34.5      0      0    330911  7.8292   NaN      Q
1 female  47.0      1      0    363272  7.0000   NaN      S
2  male  62.0      0      0    240276  9.6875   NaN      Q
3  male  27.0      0      0    315154  8.6625   NaN      S
4 female  22.0      1      1   3101298  12.2875  NaN      S
..   ...
413  male    NaN      0      0        ...     ...     ...     ...
414 female  39.0      0      0        PC 17758  108.9000  C105      C
415  male  38.5      0      0  SOTON/O.Q.  3101262  7.2500   NaN      S
416  male    NaN      0      0    359309  8.0500   NaN      S
417  male    NaN      1      1        ...     ...     ...     ...

[418 rows x 11 columns]
```

```
[9]: import pandas as pd
import numpy as np
import os

class LogisticRegression:
    def __init__(self, train, test, feature_cols, target_col, learning_rate=0.001, iters=100000, threshold=0.5):
        self.train = train
        self.test = test
        self.feature_cols = feature_cols
        self.target_col = target_col
        self.learning_rate = learning_rate
        self.iters = iters
        self.threshold = threshold
        self.theta = np.zeros((len(self.feature_cols), 1)) # column vector
```

```

def fit(self):
    self.process()
    x = np.array(self.train[self.feature_cols].values)
    y = np.array(self.train[self.target_col].values).reshape(-1, 1)
    for i in range(self.iters):
        self.theta += self.learning_rate / x.shape[0] * (x.T @ (y - self.
        ↪h(x @ self.theta)))

    print(f"trained with learning rate: {self.learning_rate} and iterations:
    ↪{self.iters}")

def predict(self, title="submission", directory="submission"):
    self.test["Survived"] = self.h(np.array(self.test[self.feature_cols].
    ↪values) @ self.theta)

    self.test.loc[self.test["Survived"] < self.threshold, "Survived"] = 0
    self.test.loc[self.test["Survived"] >= self.threshold, "Survived"] = 1
    self.test["Survived"] = self.test["Survived"].astype(int)

    os.makedirs(directory, exist_ok=True)
    self.test[["PassengerId", "Survived"]].to_csv(f"{directory}/{title}.
    ↪csv", index=False)
    print(f"Saved to {directory}/{title}.csv")

def params(self):
    print(self.theta)

def h(self, x):
    return 1 / (1 + np.exp(-x))

def process(self):
    self.train["Age"] = self.train["Age"].fillna(self.train["Age"].median())
    self.test["Age"] = self.test["Age"].fillna(self.train["Age"].median())

    self.train.loc[self.train["Embarked"] == "S", "Embarked"] = 0
    self.train.loc[self.train["Embarked"] == "C", "Embarked"] = 1
    self.train.loc[self.train["Embarked"] == "Q", "Embarked"] = 2

    self.test.loc[self.test["Embarked"] == "S", "Embarked"] = 0
    self.test.loc[self.test["Embarked"] == "C", "Embarked"] = 1
    self.test.loc[self.test["Embarked"] == "Q", "Embarked"] = 2

    self.train["Embarked"] = self.train["Embarked"].fillna(self.
    ↪train["Embarked"].mode().iloc[0])
    self.test["Embarked"] = self.test["Embarked"].fillna(self.
    ↪train["Embarked"].mode().iloc[0])

```

```

        self.train.loc[self.train["Sex"] == "male", "Sex"] = 0
        self.train.loc[self.train["Sex"] == "female", "Sex"] = 1

        self.test.loc[self.test["Sex"] == "male", "Sex"] = 0
        self.test.loc[self.test["Sex"] == "female", "Sex"] = 1
        # print(self.test[self.feature_cols].dtypes)

        self.train[self.feature_cols] = self.train[self.feature_cols].
        ↪astype(float)
        self.test[self.feature_cols] = self.test[self.feature_cols].
        ↪astype(float)

```

[10]:

```

lr = LogisticRegression(train, test, ["Pclass", "Sex", "Age", "Embarked"], ↪
    ↪"Survived", iters=100000)
lr.fit()
lr.params()
lr.predict(title="T10_100k_iters")

```

C:\Users\chyut\AppData\Local\Temp\ipykernel_22544\3690440150.py:53:
 FutureWarning: Downcasting object dtype arrays on .fillna, .ffill, .bfill is
 deprecated and will change in a future version. Call
 result.infer_objects(copy=False) instead. To opt-in to the future behavior, set
`pd.set_option('future.no_silent_downcasting', True)`
 self.train["Embarked"] =
 self.train["Embarked"].fillna(self.train["Embarked"].mode().iloc[0])
C:\Users\chyut\AppData\Local\Temp\ipykernel_22544\3690440150.py:54:
 FutureWarning: Downcasting object dtype arrays on .fillna, .ffill, .bfill is
 deprecated and will change in a future version. Call
result.infer_objects(copy=False) instead. To opt-in to the future behavior, set
`pd.set_option('future.no_silent_downcasting', True)`
 self.test["Embarked"] =
 self.test["Embarked"].fillna(self.train["Embarked"].mode().iloc[0])
trained with learning rate: 0.001 and iterations: 100000
[[-0.70281072]
 [2.59431049]
 [-0.00447441]
 [0.35126993]]
Saved to submission/T10_100k_iters.csv

[11]:

```

lr = LogisticRegression(train, test, ["Pclass", "Sex", "Age", "Embarked"], ↪
    ↪"Survived", iters=2000000)
lr.fit()
lr.params()
lr.predict(title="T10_2m_iters")

```

trained with learning rate: 0.001 and iterations: 2000000
[[-0.71333809]
 [2.65049353]

```
[-0.00447566]
[ 0.35092113]
Saved to submission/T10_2m_iters.csv
```

```
[12]: print(pd.read_csv("submission/T10_2m_iters.csv").set_index("PassengerId"))
```

```
Survived
PassengerId
892          0
893          1
894          0
895          0
896          1
...
1305         0
1306         1
1307         0
1308         0
1309         0
```

```
[418 rows x 1 columns]
```

0.2.4 T11.

Submissions

All	Successful	Errors	Recent ▾
Submission and Description			Public Score ⓘ
T10_2m_iters.csv Complete · now			0.76555

0.2.5 T12.

```
[13]: train_high = train.copy()
train_high["Embarked**2"] = train_high["Embarked"] ** 2
train_high["Age*Pclass"] = train_high["Age"] * train_high["Pclass"]
train_high["Age**2"] = train_high["Age"] ** 2
train_high["Age*Sex"] = train_high["Age"] * train_high["Sex"]

test_high = test.copy()
test_high["Embarked**2"] = test_high["Embarked"] ** 2
test_high["Age*Pclass"] = test_high["Age"] * test_high["Pclass"]
test_high["Age**2"] = test_high["Age"] ** 2
test_high["Age*Sex"] = test_high["Age"] * test_high["Sex"]
```

```
[14]:
```

```

lr = LogisticRegression(train_high, test_high,
    ↪["Pclass", "Sex", "Age", "Embarked", "Embarked**2", "Age*Pclass", "Age**2", ↪
    ↪"Age*Sex"], "Survived", iters=100000)
lr.fit()
lr.params()
lr.predict(title="T12_100k_iters")

```

C:\Users\chyut\AppData\Local\Temp\ipykernel_22544\3690440150.py:39:
RuntimeWarning: overflow encountered in exp
return 1 / (1 + np.exp(-x))
trained with learning rate: 0.001 and iterations: 100000
[[-0.52522612]
[-2.49243163]
[2.86031408]
[0.8270162]
[0.7425793]
[-1.47554321]
[0.01012895]
[7.54272117]]
Saved to submission/T12_100k_iters.csv

Submission and Description	Public Score
T13_100k_iters.csv Complete · 1h ago	0.75358
T12_100k_iters.csv Complete · 1h ago	0.58851

0.2.6 T13.

```
[15]: lr = LogisticRegression(train_high, test_high, ["Age", "Sex"], ↪
    ↪"Survived", iters=100000)
lr.fit()
lr.params()
lr.predict(title="T13_100k_iters")
```

trained with learning rate: 0.001 and iterations: 100000
[[-0.04000972]
[2.21648735]]
Saved to submission/T13_100k_iters.csv

Submission and Description	Public Score
T13_100k_iters.csv Complete · 1h ago	0.75358
T12_100k_iters.csv Complete · 1h ago	0.58851

0.2.7 OT3.

```
[16]: import pandas as pd
import numpy as np
import os

class LinearRegression:
    def __init__(self, train, test, feature_cols, target_col, learning_rate=0.001, iters=100000, threshold=0.5):
        self.train = train
        self.test = test
        self.feature_cols = feature_cols
        self.target_col = target_col
        self.learning_rate = learning_rate
        self.iters = iters
        self.threshold = threshold
        self.theta = np.zeros((len(self.feature_cols), 1)) # column vector

    def fit(self):
        self.process()
        x = np.array(self.train[self.feature_cols].values)
        y = np.array(self.train[self.target_col].values).reshape(-1, 1)
        for i in range(self.iters):
            self.theta += self.learning_rate / x.shape[0] * (x.T @ (y - (x @ self.theta)))

        print(f"trained with learning rate: {self.learning_rate} and iterations: {self.iters}")

    def params(self):
        print(self.theta)
        return self.theta

    def process(self):
        self.train["Age"] = self.train["Age"].fillna(self.train["Age"].median())
        self.test["Age"] = self.test["Age"].fillna(self.train["Age"].median())

        self.train.loc[self.train["Embarked"] == "S", "Embarked"] = 0
        self.train.loc[self.train["Embarked"] == "C", "Embarked"] = 1
        self.train.loc[self.train["Embarked"] == "Q", "Embarked"] = 2

        self.test.loc[self.test["Embarked"] == "S", "Embarked"] = 0
        self.test.loc[self.test["Embarked"] == "C", "Embarked"] = 1
        self.test.loc[self.test["Embarked"] == "Q", "Embarked"] = 2

        self.train["Embarked"] = self.train["Embarked"].fillna(self.train["Embarked"].mode().iloc[0])
```

```

        self.test["Embarked"] = self.test["Embarked"].fillna(self.
    ↪train["Embarked"].mode().iloc[0])

        self.train.loc[self.train["Sex"] == "male", "Sex"] = 0
        self.train.loc[self.train["Sex"] == "female", "Sex"] = 1

        self.test.loc[self.test["Sex"] == "male", "Sex"] = 0
        self.test.loc[self.test["Sex"] == "female", "Sex"] = 1
        # print(self.test[self.feature_cols].dtypes)

        self.train[self.feature_cols] = self.train[self.feature_cols].
    ↪astype(float)
        self.test[self.feature_cols] = self.test[self.feature_cols].
    ↪astype(float)

```

[17]: lr = LinearRegression(train, test, ["Pclass", "Sex", "Age", "Embarked"],
 ↪"Survived", iters=1000000)
 lr.fit()
 theta_ot3 = lr.params()

trained with learning rate: 0.001 and iterations: 1000000
 [[-0.01411427]
 [0.60420619]
 [0.00501483]
 [0.06116326]]

0.2.8 OT4.

[18]: x = np.array(train[["Pclass", "Sex", "Age", "Embarked"]])
 y = np.array(train["Survived"]).reshape((-1, 1))
 theta_ot4 = np.linalg.inv(x.T @ x) @ (x.T @ y)
 theta_ot4

[18]: array([[-0.01411427],
 [0.60420619],
 [0.00501483],
 [0.06116326]])

[19]: mse = np.sum((theta_ot3 - theta_ot4) ** 2) / theta_ot3.shape[1]
 mse

[19]: np.float64(5.471194131828056e-26)