

Homework 1 Clustering and Regression

Instructions

Answer the questions and upload your answers to courseville. Answers can be in Thai or English. Answers can be either typed or handwritten and scanned. the assignment is divided into several small tasks. Each task is weighted equally (marked with **T**). For this assignment, each task is awarded 1 points. There are also optional tasks (marked with **OT**) counts for 0.5 points each.

Metrics

In a population where the amount of cats is equal to the amount of dogs. Considering the following classification results from a classifier.

Model A	Predicted dog	Predicted cat
Actual dog	30	20
Actual cat	10	40

T1. What is the accuracy of Model A? $\frac{30+40}{30+20+10+40} = \frac{70}{100} = 0.70 \text{ #}$

T2. Consider cats as 'class 1' (positive) and dogs as 'class 0' (negative), calculate the precision, recall, and F1. $\text{precision} = \frac{TP}{\# \text{predicted pos}} = \frac{40}{60} = 0.67 \text{ #}$ $\text{recall} = \frac{TP}{\# \text{actual pos}} = \frac{40}{50} = 0.80 \text{ #}$ $F_1 = \frac{2}{\frac{1}{\text{precision}} + \frac{1}{\text{recall}}} = \frac{2(0.67)(0.80)}{0.67+0.80} = 0.73 \text{ #}$

T3. Consider class cat as 'class 0' and class dog as 'class 1', calculate the precision, recall, and F1. $\text{precision} = \frac{TP}{\# \text{predicted pos}} = \frac{20}{90} = 0.75 \text{ #}$ $\text{recall} = \frac{TP}{\# \text{actual pos}} = \frac{20}{50} = 0.60 \text{ #}$ $F_1 = \frac{2(0.75)(0.60)}{0.75+0.60} = 0.67 \text{ #}$

It is important to specify the 'positive' class when you calculate precision, recall, and F1. If there are **more than two classes**, it is usually done in a **one-versus-all** setting where one class is considered positive and the rest of the classes are considered negative.

T4. Now consider a lopsided population where there are 80% cats. What is the accuracy of Model A? Using dog as the positive class, what is the precision, recall, and F1? Explain how and why these numbers change (or does not change) from the previous questions. $\text{Accuracy} = \frac{68}{100} = 0.68 \text{ #}$ $\text{Precision} = \frac{15}{20} = \frac{3}{4} \text{ #}$ $\text{Recall} = \frac{15}{42} = \frac{5}{14} \text{ #}$ $F_1 = \frac{2}{\frac{3}{4} + \frac{5}{14}} = \frac{15}{21} \text{ #}$

OT1. Consider the equations for accuracy and F1 $\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$ $F_1 = \frac{2TP}{2TP + FP + FN}$ change, because the ratio has change

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN} \quad (1)$$

$$F_1 = \frac{2TP}{2TP + FP + FN}$$

When will accuracy be equal, greater, or less than F1?

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN} \quad F_1 = \frac{TP + TN}{TP + TN + FP + FN}$$

let $k_1 = TP + FP + FN$, $k_2 = TN$

$$\text{Accuracy} = \frac{k_1 + TN}{k_1 + TN} = \frac{k_1 + TN - k_1 + k_2}{k_1 + TN} = 1 - \frac{k_1 - k_2}{k_1 + TN}$$

$$F_1 = \frac{k_1 + TN}{k_1 + TN} = \frac{k_1 + TN - k_1 + k_2}{k_1 + TN} = 1 - \frac{k_1 - k_2}{k_1 + TN}$$

$1 - \frac{k_1 - k_2}{k_1 + TN}$

$1 - \frac{k_1 - k_2}{k_1 + TN}$

$\frac{k_1 - k_2}{k_1 + TN}, \frac{k_1 - k_2}{k_1 + TN}$

accuracy equals to F1: $TN = TP \text{ #}$

accuracy less than F1: $TN > TP \text{ #}$

accuracy more than F1: $TN < TP \text{ #}$

Homework1

January 9, 2026

0.1 K-Means

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

class KMeans:
    def __init__(self, init):
        self.data = np.array([(1, 2), (3, 3), (2, 2), (8, 8), (6, 6), (7, 7),
        (-3, -3), (-2, -4), (-7, -7)])
        self.init = init
        self.old_centroids = []
        self.new_centroids = init.copy()
        self.group = {}

    def cluster(self, display=True):
        i = 0
        while (self.new_centroids != self.old_centroids):
            self.old_centroids = self.new_centroids

            centroids = np.array(self.old_centroids)
            self.group = {}
            for each_data in self.data:
                # Vectorization Optimized
                point = np.array(each_data)
                dist = np.sqrt(np.sum((centroids - point) ** 2, axis=1))
                group_of_each_data = np.argmin(dist)

                if (group_of_each_data not in self.group):
                    self.group[group_of_each_data] = []
                self.group[group_of_each_data].append(each_data)

            # Traditional Way
            # min_distance = 1e9
            # group_of_each_data = -1
            # for i in range(len(self.old_centroids)):
            #     current_distance = ( (self.old_centroids[i][0] -
            #     each_data[0]) ** 2 + (self.old_centroids[i][1] - each_data[1]) ** 2 ) ** 0.5
```

```

        #     if (current_distance < min_distance):
        #         group_of_each_data = i
        #         min_distance = current_distance

        # if (group_of_each_data not in self.group):
        #     self.group[group_of_each_data] = []

        # self.group[group_of_each_data].append(each_data)

        self.new_centroids = []
        for each in self.group.keys():
            centroid_x = sum(t[0] for t in self.group[each]) / len(self.
        ↪group[each])
            centroid_y = sum(t[1] for t in self.group[each]) / len(self.
        ↪group[each])
            self.new_centroids.append((centroid_x, centroid_y))

        i += 1

        if (display):
            print(f"--- Round {i} ---")
            print(f"Centroids: {self.new_centroids}")

        if (display):
            for cluster_id, points in self.group.items():
                xs, ys = zip(*points)
                plt.scatter(xs, ys, label=f"Cluster {cluster_id}")

            plt.xlabel("X")
            plt.ylabel("Y")
            plt.show()

```

0.1.1 T5. starting points are (3,3), (2,2), and (-3,-3).

[2]: init_t5 = [(3, 3), (2, 2), (-3, -3)]

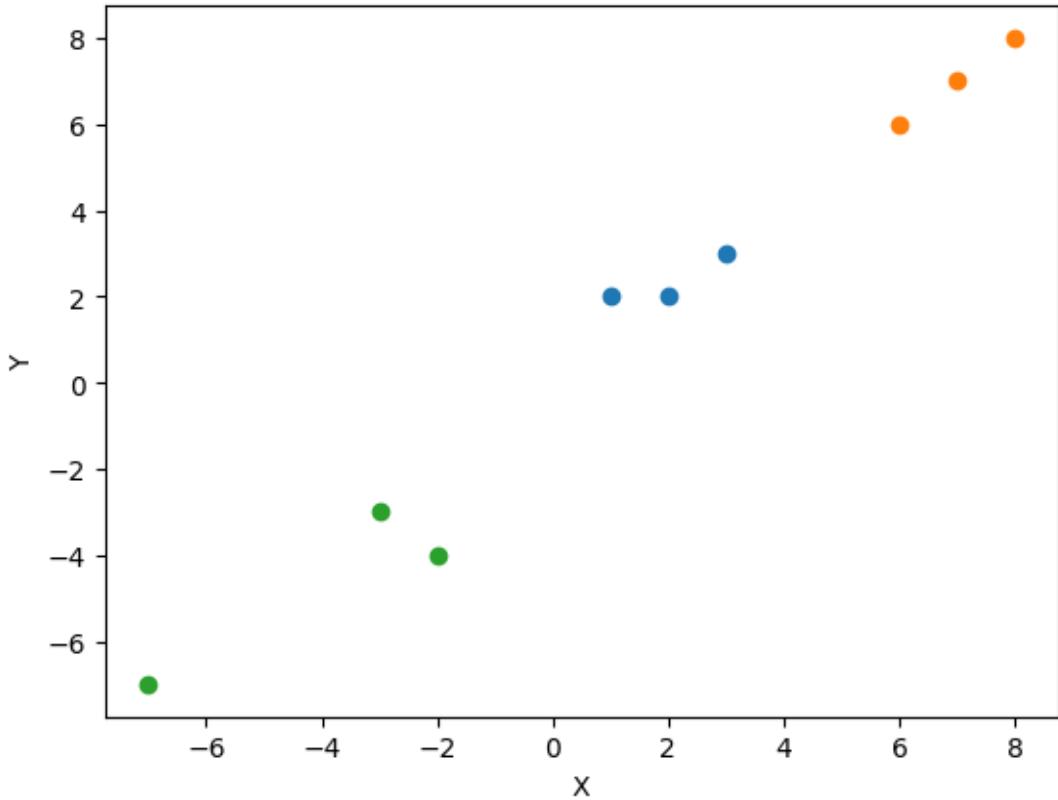
```
km_t5 = KMeans(init_t5)
km_t5.cluster()
```

```

--- Round 1 ---
Centroids: [(np.float64(1.5), np.float64(2.0)), (np.float64(6.0),
np.float64(6.0)), (np.float64(-4.0), np.float64(-4.66666666666667))]
--- Round 2 ---
Centroids: [(np.float64(2.0), np.float64(2.33333333333335)), (np.float64(7.0),
np.float64(7.0)), (np.float64(-4.0), np.float64(-4.66666666666667))]
--- Round 3 ---
Centroids: [(np.float64(2.0), np.float64(2.33333333333335)), (np.float64(7.0),

```

```
np.float64(7.0)), (np.float64(-4.0), np.float64(-4.66666666666667))]
```



0.1.2 T6. starting points are (-3,-3), (2,2), and (-7,-7),

```
[3]: init_t6 = [(-3, -3), (2, 2), (-7, -7)]
```

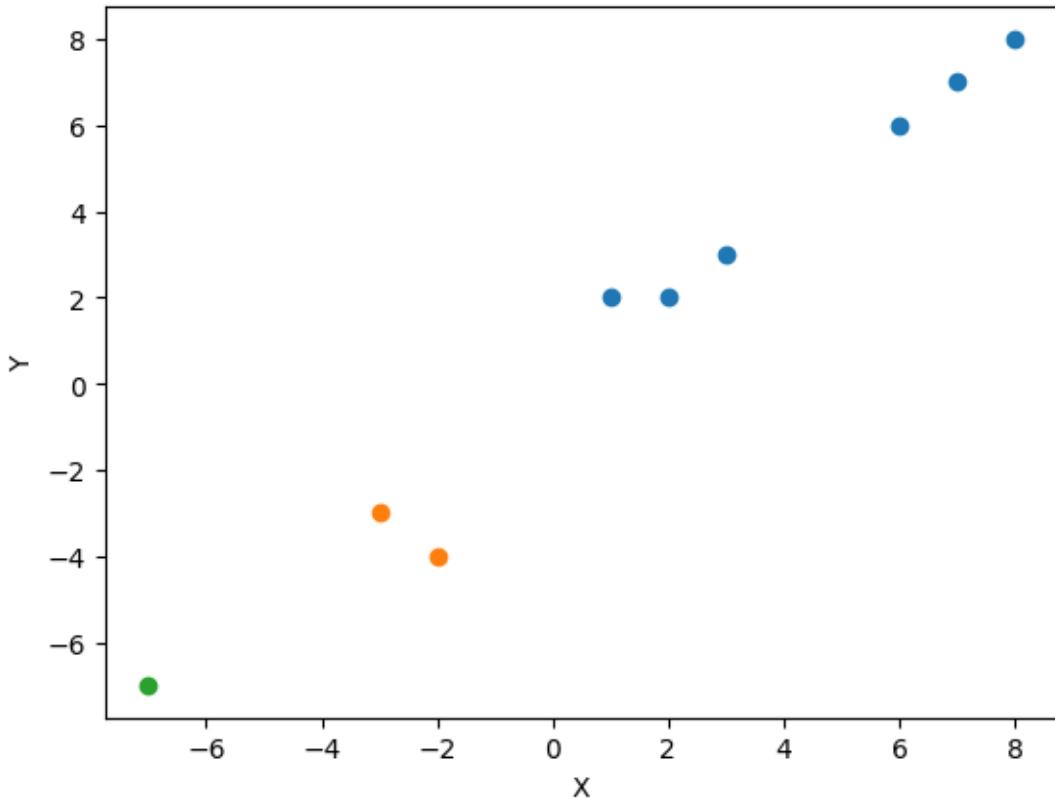
```
km_t6 = KMeans(init_t6)
km_t6.cluster()
```

```
--- Round 1 ---
```

```
Centroids: [(np.float64(4.5), np.float64(4.66666666666667)), (np.float64(-2.5),
np.float64(-3.5)), (np.float64(-7.0), np.float64(-7.0))]
```

```
--- Round 2 ---
```

```
Centroids: [(np.float64(4.5), np.float64(4.66666666666667)), (np.float64(-2.5),
np.float64(-3.5)), (np.float64(-7.0), np.float64(-7.0))]
```



- 0.1.3 T7. Between the two starting set of points in the previous two questions, which one do you think is better? How would you measure the 'goodness' quality of a set of starting points?
- 0.1.4 In general, it is important to try different sets of starting points when doing k-means.

For me, the better one is the first one. The 'goodness' may measure with the variance in the g

- 0.1.5 OT2. What would be the best K for this question? Describe your reasoning.

For me, it's 4 because the distance between (-7, -7) to (-3, -3) and (-2, -4) is too large to l

0.2 Regression

```
[4]: train_url = "http://s3.amazonaws.com/assets.datacamp.com/course/Kaggle/train.csv"
train = pd.read_csv(train_url)

test_url = "http://s3.amazonaws.com/assets.datacamp.com/course/Kaggle/test.csv"
test = pd.read_csv(test_url)
```

```
[5]: train.describe()
```

```
[5]:    PassengerId  Survived  Pclass      Age      SibSp  \
count  891.000000  891.000000  891.000000  714.000000  891.000000
mean   446.000000  0.383838   2.308642  29.699118   0.523008
std    257.353842  0.486592   0.836071  14.526497   1.102743
min    1.000000  0.000000   1.000000  0.420000   0.000000
25%   223.500000  0.000000   2.000000  20.125000   0.000000
50%   446.000000  0.000000   3.000000  28.000000   0.000000
75%   668.500000  1.000000   3.000000  38.000000   1.000000
max   891.000000  1.000000   3.000000  80.000000   8.000000

          Parch      Fare
count  891.000000  891.000000
mean   0.381594  32.204208
std    0.806057  49.693429
min    0.000000  0.000000
25%   0.000000  7.910400
50%   0.000000  14.454200
75%   0.000000  31.000000
max   6.000000  512.329200
```

0.2.1 T8. Median of the training datasets?

```
[6]: train["Age"].median()
```

```
[6]: np.float64(28.0)
```

28

0.2.2 T9. Mode of Embarked

```
[7]: train["Embarked"].mode().iloc[0]
```

```
[7]: 'S'
```

"S" which is 0

0.2.3 T10.

```
[8]: test
```

```
[8]:    PassengerId  Pclass          Name
0            892      3  Kelly, Mr. James
1            893      3  Wilkes, Mrs. James (Ellen Needs)
2            894      2  Myles, Mr. Thomas Francis
3            895      3  Wirz, Mr. Albert
4            896      3  Hirvonen, Mrs. Alexander (Helga E Lindqvist)
..          ...    ...

```

413	1305	3	Spector, Mr. Woolf					
414	1306	1	Oliva y Ocana, Dona. Fermina					
415	1307	3	Saether, Mr. Simon Sivertsen					
416	1308	3	Ware, Mr. Frederick					
417	1309	3	Peter, Master. Michael J					
	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	male	34.5	0	0	330911	7.8292	NaN	Q
1	female	47.0	1	0	363272	7.0000	NaN	S
2	male	62.0	0	0	240276	9.6875	NaN	Q
3	male	27.0	0	0	315154	8.6625	NaN	S
4	female	22.0	1	1	3101298	12.2875	NaN	S
..
413	male	NaN	0	0	A.5. 3236	8.0500	NaN	S
414	female	39.0	0	0	PC 17758	108.9000	C105	C
415	male	38.5	0	0	SOTON/O.Q. 3101262	7.2500	NaN	S
416	male	NaN	0	0	359309	8.0500	NaN	S
417	male	NaN	1	1	2668	22.3583	NaN	C

[418 rows x 11 columns]

```
[9]: import pandas as pd
import numpy as np
import os

class LogisticRegression:
    def __init__(self, train, test, feature_cols, target_col, learning_rate=0.001, iters=100000, threshold=0.5):
        self.train = train
        self.test = test
        self.feature_cols = feature_cols
        self.target_col = target_col
        self.learning_rate = learning_rate
        self.iters = iters
        self.threshold = threshold
        self.theta = np.zeros((len(self.feature_cols), 1)) # column vector

    def fit(self):
        self.process()
        x = np.array(self.train[self.feature_cols].values)
        y = np.array(self.train[self.target_col].values).reshape(-1, 1)
        for i in range(self.iters):
            self.theta += self.learning_rate / x.shape[0] * (x.T @ (y - self.h(x @ self.theta)))

        print(f"trained with learning rate: {self.learning_rate} and iterations: {self.iters}")

    def predict(self, x):
        return self.h(x @ self.theta)
```

```

def predict(self, title="submission", directory="submission"):
    self.test["Survived"] = self.h(np.array(self.test[self.feature_cols].
    ↪values) @ self.theta)
    self.test.loc[self.test["Survived"] < self.threshold, "Survived"] = 0
    self.test.loc[self.test["Survived"] >= self.threshold, "Survived"] = 1
    self.test["Survived"] = self.test["Survived"].astype(int)

    os.makedirs(directory, exist_ok=True)
    self.test[["PassengerId", "Survived"]].to_csv(f"{directory}/{title}.

    ↪.csv", index=False)
    print(f"Saved to {directory}/{title}.csv")

def params(self):
    print(self.theta)

def h(self, x):
    return 1 / (1 + np.exp(-x))

def process(self):
    self.train["Age"] = self.train["Age"].fillna(self.train["Age"].median())
    self.test["Age"] = self.test["Age"].fillna(self.train["Age"].median())

    self.train.loc[self.train["Embarked"] == "S", "Embarked"] = 0
    self.train.loc[self.train["Embarked"] == "C", "Embarked"] = 1
    self.train.loc[self.train["Embarked"] == "Q", "Embarked"] = 2

    self.test.loc[self.test["Embarked"] == "S", "Embarked"] = 0
    self.test.loc[self.test["Embarked"] == "C", "Embarked"] = 1
    self.test.loc[self.test["Embarked"] == "Q", "Embarked"] = 2

    self.train["Embarked"] = self.train["Embarked"].fillna(self.

    ↪train["Embarked"].mode().iloc[0])
    self.test["Embarked"] = self.test["Embarked"].fillna(self.

    ↪train["Embarked"].mode().iloc[0])

    self.train.loc[self.train["Sex"] == "male", "Sex"] = 0
    self.train.loc[self.train["Sex"] == "female", "Sex"] = 1

    self.test.loc[self.test["Sex"] == "male", "Sex"] = 0
    self.test.loc[self.test["Sex"] == "female", "Sex"] = 1
    # print(self.test[self.feature_cols].dtypes)

    self.train[self.feature_cols] = self.train[self.feature_cols].

    ↪astype(float)
    self.test[self.feature_cols] = self.test[self.feature_cols].

    ↪astype(float)

```

```
[10]: lr = LogisticRegression(train, test, ["Pclass", "Sex", "Age", "Embarked"], ↴
    ↴"Survived",iters=100000)
lr.fit()
lr.params()
lr.predict(title="T10_100k_iters")
```

C:\Users\chyut\AppData\Local\Temp\ipykernel_22544\3690440150.py:53:
 FutureWarning: Downcasting object dtype arrays on .fillna, .ffill, .bfill is
 deprecated and will change in a future version. Call
 result.infer_objects(copy=False) instead. To opt-in to the future behavior, set
 `pd.set_option('future.no_silent_downcasting', True)`
 self.train["Embarked"] =
 self.train["Embarked"].fillna(self.train["Embarked"].mode().iloc[0])
C:\Users\chyut\AppData\Local\Temp\ipykernel_22544\3690440150.py:54:
 FutureWarning: Downcasting object dtype arrays on .fillna, .ffill, .bfill is
 deprecated and will change in a future version. Call
 result.infer_objects(copy=False) instead. To opt-in to the future behavior, set
 `pd.set_option('future.no_silent_downcasting', True)`
 self.test["Embarked"] =
 self.test["Embarked"].fillna(self.train["Embarked"].mode().iloc[0])
trained with learning rate: 0.001 and iterations: 100000
[[-0.70281072]
 [2.59431049]
 [-0.00447441]
 [0.35126993]]
Saved to submission/T10_100k_iters.csv

```
[11]: lr = LogisticRegression(train, test, ["Pclass", "Sex", "Age", "Embarked"], ↴
    ↴"Survived",iters=2000000)
lr.fit()
lr.params()
lr.predict(title="T10_2m_iters")
```

trained with learning rate: 0.001 and iterations: 2000000
[[-0.71333809]
 [2.65049353]
 [-0.00447566]
 [0.35092113]]
Saved to submission/T10_2m_iters.csv

```
[12]: print(pd.read_csv("submission/T10_2m_iters.csv").set_index("PassengerId"))
```

PassengerId	Survived
892	0
893	1
894	0
895	0

```

896          1
...
1305          0
1306          1
1307          0
1308          0
1309          0

```

[418 rows x 1 columns]

0.2.4 T11.

Submissions

			Recent ▾
	Submission and Description	Public Score ⓘ	
 T10_2m_iters.csv	Complete · now	0.76555	

0.2.5 T12.

```

[13]: train_high = train.copy()
train_high["Embarked**2"] = train_high["Embarked"] ** 2
train_high["Age*Pclass"] = train_high["Age"] * train_high["Pclass"]
train_high["Age**2"] = train_high["Age"] ** 2
train_high["Age*Sex"] = train_high["Age"] * train_high["Sex"]

test_high = test.copy()
test_high["Embarked**2"] = test_high["Embarked"] ** 2
test_high["Age*Pclass"] = test_high["Age"] * test_high["Pclass"]
test_high["Age**2"] = test_high["Age"] ** 2
test_high["Age*Sex"] = test_high["Age"] * test_high["Sex"]

```

```

[14]: lr = LogisticRegression(train_high, test_high,
    ↴ ["Pclass", "Sex", "Age", "Embarked", "Embarked**2", "Age*Pclass", "Age**2", ↴
    ↴ "Age*Sex"], "Survived", iters=100000)
lr.fit()
lr.params()
lr.predict(title="T12_100k_iters")

```

C:\Users\chyut\AppData\Local\Temp\ipykernel_22544\3690440150.py:39:

RuntimeWarning: overflow encountered in exp

 return 1 / (1 + np.exp(-x))

trained with learning rate: 0.001 and iterations: 100000

[[-0.52522612]

[-2.49243163]

[2.86031408]

```
[ 0.8270162 ]  
[ 0.7425793 ]  
[-1.47554321]  
[ 0.01012895]  
[ 7.54272117]  
Saved to submission/T12_100k_iters.csv
```

Submission and Description	Public Score
 T13_100k_iters.csv Complete · 1h ago	0.75358
 T12_100k_iters.csv Complete · 1h ago	0.58851

0.2.6 T13.

```
[15]: lr = LogisticRegression(train_high, test_high, ["Age", "Sex"],  
    ↪ "Survived", iters=100000)  
lr.fit()  
lr.params()  
lr.predict(title="T13_100k_iters")
```

```
trained with learning rate: 0.001 and iterations: 100000  
[[-0.04000972]  
 [ 2.21648735]]
```

```
Saved to submission/T13_100k_iters.csv
```

Submission and Description	Public Score
 T13_100k_iters.csv Complete · 1h ago	0.75358
 T12_100k_iters.csv Complete · 1h ago	0.58851

0.2.7 OT3.

```
[16]: import pandas as pd  
import numpy as np  
import os  
  
class LinearRegression:  
    def __init__(self, train, test, feature_cols, target_col, learning_rate=0.  
    ↪ 001, iters=100000, threshold=0.5):  
        self.train = train  
        self.test = test  
        self.feature_cols = feature_cols  
        self.target_col = target_col  
        self.learning_rate = learning_rate
```

```

        self.iters = iters
        self.threshold = threshold
        self.theta = np.zeros((len(self.feature_cols), 1)) # column vector

    def fit(self):
        self.process()
        x = np.array(self.train[self.feature_cols].values)
        y = np.array(self.train[self.target_col].values).reshape(-1, 1)
        for i in range(self.iters):
            self.theta += self.learning_rate / x.shape[0] * (x.T @ (y - (x @ self.theta)))

            print(f"trained with learning rate: {self.learning_rate} and iterations: {self.iters}")

    def params(self):
        print(self.theta)
        return self.theta

    def process(self):
        self.train["Age"] = self.train["Age"].fillna(self.train["Age"].median())
        self.test["Age"] = self.test["Age"].fillna(self.train["Age"].median())

        self.train.loc[self.train["Embarked"] == "S", "Embarked"] = 0
        self.train.loc[self.train["Embarked"] == "C", "Embarked"] = 1
        self.train.loc[self.train["Embarked"] == "Q", "Embarked"] = 2

        self.test.loc[self.test["Embarked"] == "S", "Embarked"] = 0
        self.test.loc[self.test["Embarked"] == "C", "Embarked"] = 1
        self.test.loc[self.test["Embarked"] == "Q", "Embarked"] = 2

        self.train["Embarked"] = self.train["Embarked"].fillna(self.
        ↪train["Embarked"].mode().iloc[0])
        self.test["Embarked"] = self.test["Embarked"].fillna(self.
        ↪train["Embarked"].mode().iloc[0])

        self.train.loc[self.train["Sex"] == "male", "Sex"] = 0
        self.train.loc[self.train["Sex"] == "female", "Sex"] = 1

        self.test.loc[self.test["Sex"] == "male", "Sex"] = 0
        self.test.loc[self.test["Sex"] == "female", "Sex"] = 1
        # print(self.test[self.feature_cols].dtypes)

        self.train[self.feature_cols] = self.train[self.feature_cols].
        ↪astype(float)
        self.test[self.feature_cols] = self.test[self.feature_cols].
        ↪astype(float)

```

```
[17]: lr = LinearRegression(train, test, ["Pclass", "Sex", "Age", "Embarked"],  
    ↴ "Survived", iters=1000000)  
lr.fit()  
theta_ot3 = lr.params()
```

trained with learning rate: 0.001 and iterations: 1000000
[[-0.01411427]
 [0.60420619]
 [0.00501483]
 [0.06116326]]

0.2.8 OT4.

```
[18]: x = np.array(train[["Pclass", "Sex", "Age", "Embarked"]])  
y = np.array(train["Survived"]).reshape((-1, 1))  
theta_ot4 = np.linalg.inv(x.T @ x) @ (x.T @ y)  
theta_ot4
```

```
[18]: array([[ -0.01411427],  
 [ 0.60420619],  
 [ 0.00501483],  
 [ 0.06116326]])
```

```
[19]: mse = np.sum((theta_ot3 - theta_ot4) ** 2) / theta_ot3.shape[1]  
mse
```

```
[19]: np.float64(5.471194131828056e-26)
```

$$\text{OT5. } \nabla_A \text{tr}(AB) = B^T$$

$$\text{ก็ } (AB)_{ij} = \sum_m A_{im} B_{mj}$$

$$\text{as trace is diagonal sum, so it is } \sum_k (AB)_{kk} = \sum_k \sum_m A_{km} B_{mk} = \sum_i \sum_j A_{ij} B_{ji}$$

ก็ $f_i(\sum_j A_{ij} B_{ji})$

$$\nabla_A \text{tr}(AB) = \frac{\partial}{\partial A} (\sum_i \sum_j A_{ij} B_{ji}) = \frac{\partial f}{\partial A} = \begin{bmatrix} \frac{\partial f}{\partial A_{11}} & \frac{\partial f}{\partial A_{12}} & \cdots & \frac{\partial f}{\partial A_{1N}} \\ \vdots & \vdots & \ddots & \vdots \\ \vdots & & \ddots & \vdots \\ \frac{\partial f}{\partial A_{N1}} & \cdots & \frac{\partial f}{\partial A_{NN}} \end{bmatrix}$$

↑ หมายความว่า (i,j) นั้นหมายความว่า ได้ผลลัพธ์ของ A_{ij} และ B_{ji}

$$= \begin{bmatrix} B_{11} & B_{21} & \cdots & B_{N1} \\ \vdots & \ddots & \ddots & \vdots \\ B_{1N} & \cdots & \cdots & B_{NN} \end{bmatrix} = B^T \text{ #}$$

$$\text{OT6. } \nabla_A^T f(A) = (\nabla_A f(A))^T$$

$$\text{ก็ } G = \nabla_A f(A) \text{ ก็ } G_{ij} = \frac{\partial f}{\partial A_{ij}}$$

$$\text{ก็ } G^T = (\nabla_A f(A))^T \text{ ก็ } G^T_{ij} = (\nabla_A f(A))^T_{ij} = \frac{\partial f}{\partial A_{ji}}$$

$$\text{หมายความ } B = A^T ; B_{ji} = A_{ij}$$

$$\nabla_B f_{ji} \cdot \nabla_A f_{ji} = \frac{\partial f}{\partial B_{ji}} \cdot \frac{\partial f}{\partial A_{ij}}$$

$$\therefore \nabla_A^T f_{ij} = \frac{\partial f}{\partial A_{ji}}$$

$$\therefore \text{ก็ } \nabla_A^T f = (\nabla_A f(A))^T \text{ #}$$