

**04\_Neural\_Networks\cattern\neural\_net.py**

```
1 from __future__ import print_function
2
3 import numpy as np
4 import matplotlib.pyplot as plt
5
6 class TwoLayerNet(object):
7     """
8         A two-layer fully-connected neural network. The net has an input dimension of
9             N, a hidden layer dimension of H, and performs classification over C classes.
10            We train the network with a softmax loss function and L2 regularization on the
11            weight matrices. The network uses a ReLU nonlinearity after the first fully
12            connected layer.
13
14     In other words, the network has the following architecture:
15
16     input - fully connected layer - ReLU - fully connected layer - softmax
17
18     The outputs of the second fully-connected layer are the scores for each class.
19     """
20
21     def __init__(self, input_size, hidden_size, output_size, std=1e-4):
22         """
23             Initialize the model. Weights are initialized to small random values and
24             biases are initialized to zero. Weights and biases are stored in the
25             variable self.params, which is a dictionary with the following keys:
26
27             W1: First layer weights; has shape (D, H)
28             b1: First layer biases; has shape (H,)
29             W2: Second layer weights; has shape (H, C)
30             b2: Second layer biases; has shape (C,)
31
32             Inputs:
33             - input_size: The dimension D of the input data.
34             - hidden_size: The number of neurons H in the hidden layer.
35             - output_size: The number of classes C.
36         """
37
38         self.params = {}
39         self.params['W1'] = std * np.random.randn(input_size, hidden_size)
40         self.params['b1'] = np.zeros(hidden_size)
41         self.params['W2'] = std * np.random.randn(hidden_size, output_size)
42         self.params['b2'] = np.zeros(output_size)
43
44     def loss(self, X, y=None, reg=0.0):
45         """
46             Compute the loss and gradients for a two layer fully connected neural
47             network.
48
49             Inputs:
```

```

49     - X: Input data of shape (N, D). Each X[i] is a training sample.
50     - y: Vector of training labels. y[i] is the label for X[i], and each y[i] is
51       an integer in the range  $0 \leq y[i] < C$ . This parameter is optional; if it
52       is not passed then we only return scores, and if it is passed then we
53       instead return the loss and gradients.
54     - reg: Regularization strength.
55
56     Returns:
57     If y is None, return a matrix scores of shape (N, C) where scores[i, c] is
58     the score for class c on input X[i].
59
60     If y is not None, instead return a tuple of:
61     - loss: Loss (data loss and regularization loss) for this batch of training
62       samples.
63     - grads: Dictionary mapping parameter names to gradients of those parameters
64       with respect to the loss function; has the same keys as self.params.
65     """
66
67     # Unpack variables from the params dictionary
68     W1, b1 = self.params['W1'], self.params['b1']
69     W2, b2 = self.params['W2'], self.params['b2']
70     N, D = X.shape
71
72     # Compute the forward pass
73     scores = None
74     ##### T#5: Perform the forward pass, computing the class scores for the      #
75     # input.                                         #
76     # Store the result in the scores variable, which should be an array of      #
77     # shape (N, C). Note that this does not include the softmax                 #
78     # HINT: This is just a series of matrix multiplication.                      #
79     #####
80     h1 = X @ W1 + b1
81     h1_relu = np.maximum(0, h1)
82     scores = h1_relu @ W2 + b2
83     #####
84     #                                     END OF T#5                         #
85     #####
86
87     # If the targets are not given then jump out, we're done
88     if y is None:
89         return scores
90
91     # Compute the loss
92     loss = None
93     #####
94     # T#6: Finish the forward pass, and compute the loss. This should include#
95     # both the data loss and L2 regularization for W1 and W2. Store the result  #
96     # in the variable loss, which should be a scalar. Use the Softmax           #
97     # classifier loss.                                         #
98     #####

```

```

99     shift_scores = scores - np.max(scores, axis=1, keepdims=True) # prevent overflow
100    softmax = np.exp(shift_scores)
101    divider = np.sum(softmax, axis=1, keepdims=True)
102    softmax = softmax / divider
103
104    loss = -np.sum(np.log(softmax[np.arange(N), y]))
105
106    # Average for each data
107    loss /= y.shape[0]
108    # Regularization 0.5 * lambda * sigma (every weight square)
109    loss += 0.5 * reg * (np.sum(W1**2) + np.sum(W2**2))
110    #####
111    #                                     END OF T#6
112    #####
113
114    # Backward pass: compute gradients
115    grads = {}
116    #####
117    # T#7: Compute the backward pass, computing derivatives of the weights #
118    # and biases. Store the results in the grads dictionary. For example, #
119    # grads['W1'] should store the gradient on W1, and be a matrix of same size #
120    # don't forget about the regularization term
121    #####
122
123    # init with gradient of regularization term
124    grads["W1"] = reg * self.params["W1"]
125    grads["W2"] = reg * self.params["W2"]
126    grads["b1"] = np.zeros(b1.shape)
127    grads["b2"] = np.zeros(b2.shape)
128
129    dscores = softmax.copy()
130    dscores[np.arange(N), y] -= 1
131    dscores /= N # if not divide here then divide at the the last step of calculation which
132    # doing it here is better
133    grads["b2"] += 1 * np.sum(dscores, axis=0)
134
135    grads["W2"] += h1_relu.T @ dscores
136
137    dh1_relu = dscores @ W2.T
138    dh1_relu[h1 <= 0] = 0
139    grads["b1"] += np.sum(dh1_relu * 1, axis=0)
140
141    grads["W1"] += X.T @ dh1_relu
142
143    #####
144    #                                     END OF T#7
145    #####
146
147    return loss, grads

```

```
148
149 def train(self, X, y, X_val, y_val,
150         learning_rate=1e-3, learning_rate_decay=0.95,
151         reg=5e-6, num_iters=100,
152         batch_size=200, verbose=False):
153     """
154     Train this neural network using stochastic gradient descent.
155
156     Inputs:
157     - X: A numpy array of shape (N, D) giving training data.
158     - y: A numpy array f shape (N,) giving training labels; y[i] = c means that
159       X[i] has label c, where 0 <= c < C.
160     - X_val: A numpy array of shape (N_val, D) giving validation data.
161     - y_val: A numpy array of shape (N_val,) giving validation labels.
162     - learning_rate: Scalar giving learning rate for optimization.
163     - learning_rate_decay: Scalar giving factor used to decay the learning rate
164       after each epoch.
165     - reg: Scalar giving regularization strength.
166     - num_iters: Number of steps to take when optimizing.
167     - batch_size: Number of training examples to use per step.
168     - verbose: boolean; if true print progress during optimization.
169     """
170     num_train = X.shape[0]
171     iterations_per_epoch = max(num_train / batch_size, 1)
172
173     # Use SGD to optimize the parameters in self.model
174     loss_history = []
175     train_acc_history = []
176     val_acc_history = []
177
178     for it in range(num_iters):
179         X_batch = None
180         y_batch = None
181
182         #####
183         # T#8: Create a random minibatch of training data and labels, storing#
184         # them in X_batch and y_batch respectively.                                #
185         # You might find np.random.choice() helpful.                            #
186         #####
187         batch_index = np.random.choice(num_train, batch_size, replace=True)
188         X_batch = X[batch_index]
189         y_batch = y[batch_index]
190         #####
191         # END OF YOUR T#8                                              #
192         #####
193
194         # Compute loss and gradients using the current minibatch
195         loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
196         loss_history.append(loss)
197
```

```

198 ##### T#9: Use the gradients in the grads dictionary to update the #####
199 # parameters of the network (stored in the dictionary self.params) #
200 # using stochastic gradient descent. You'll need to use the gradients #
201 # stored in the grads dictionary defined above. #
202 #####
203 #####
204 self.params['W1'] -= learning_rate * grads["W1"]
205 self.params['W2'] -= learning_rate * grads["W2"]
206 self.params['b1'] -= learning_rate * grads["b1"]
207 self.params['b2'] -= learning_rate * grads["b2"]
208 #####
209 # END OF YOUR T#9 #
210 #####
211 #####
212 if verbose and it % 100 == 0:
213     print('iteration %d / %d: loss %f' % (it, num_iters, loss))
214 #####
215 # Every epoch, check train and val accuracy and decay learning rate.
216 if it % iterations_per_epoch == 0:
217     # Check accuracy
218     train_acc = (self.predict(X_batch) == y_batch).mean()
219     val_acc = (self.predict(X_val) == y_val).mean()
220     train_acc_history.append(train_acc)
221     val_acc_history.append(val_acc)
222 #####
223     # Decay learning rate
224 ##### T#10: Decay learning rate (exponentially) after each epoch #
225 ##### learning_rate *= learning_rate_decay
226 #####
227     learning_rate *= learning_rate_decay
228 #####
229 # END OF YOUR T#10 #
230 #####
231 #####
232 #####
233 return {
234     'loss_history': loss_history,
235     'train_acc_history': train_acc_history,
236     'val_acc_history': val_acc_history,
237 }
238 #####
239 def predict(self, X):
240     """
241     Use the trained weights of this two-layer network to predict labels for
242     data points. For each data point we predict scores for each of the C
243     classes, and assign each data point to the class with the highest score.
244     Inputs:
245     - X: A numpy array of shape (N, D) giving N D-dimensional data points to
246       classify.
247 
```

```
248
249     Returns:
250     - y_pred: A numpy array of shape (N,) giving predicted labels for each of
251       the elements of X. For all i, y_pred[i] = c means that X[i] is predicted
252       to have class c, where 0 <= c < C.
253     """
254     y_pred = None
255
256     ##### T#11: Implement this function; it should be VERY simple! #####
257
258     h1_relu = np.maximum(0, X @ self.params['W1'] + self.params['b1'])
259     scores = h1_relu @ self.params['W2'] + self.params['b2']
260     y_pred = np.argmax(scores, axis=1)
261
262     ##### END OF YOUR T#11 #####
263
264
265
266     return y_pred
267
268
269
```