

Homework 2 MLE and Naive Bayes

Instructions

Answer the questions and upload your answers to courseville. Answers can be in Thai or English. Answers can be either typed or handwritten and scanned. the assignment is divided into several small tasks. Each task is weighted equally (marked with **T**). For this assignment, each task is awarded 1 points. There are also optional tasks (marked with **OT**) counts for 0.5 points each.

MLE

Consider the following very simple model for stock pricing. The price at the end of each day is the price of the previous day multiplied by a fixed, but unknown, rate of return, α , with some noise, w . For a two-day period, we can observe the following sequence

$$y_2 = \alpha y_1 + w_1$$

$$y_1 = \alpha y_0 + w_0$$

where the noises w_0, w_1 are iid with the distribution $N(0, \sigma^2)$, $y_0 \sim N(0, \lambda)$ is independent of the noise sequence. σ^2 and λ are known, while α is unknown.

T1. Find the MLE of the rate of return, α , given the observed price at the end of each day y_2, y_1, y_0 . In other words, compute for the value of α that maximizes $p(y_2, y_1, y_0 | \alpha)$

Hint: This is a Markov process, e.g. y_2 is independent of y_0 given y_1 . In general, a process is Markov if $p(y_n | y_{n-1}, y_{n-2}, \dots) = p(y_n | y_{n-1})$. In other words, the present is independent of the past (y_{n-2}, y_{n-3}, \dots), conditioned on the immediate past y_{n-1} . You may also find the steps of the proof for logistic regression we did in class useful.

OT1. Consider the general case, where

$$y_{n+1} = \alpha y_n + w_n, n = 0, 1, 2, \dots$$

Find the MLE given the observed price y_{N+1}, y_N, \dots, y_0

Simple Bayes Classifier

A student in Pattern Recognition course had finally built the ultimate classifier for cat emotions. He used one input features: the amount of food the cat ate that day, x (Being a good student he already normalized x to standard Normal). He proposed the following likelihood probabilities for class 1 (happy cat) and 2 (sad cat)

$$P(x | w_1) = N(4, 2)$$

$$P(x | w_2) = N(0, 2)$$

$$\delta \sim N(0, 1)$$

MLE

Consider the following very simple model for stock pricing. The price at the end of each day is the price of the previous day multiplied by a fixed, but unknown, rate of return, α , with some noise, w . For a two-day period, we can observe the following sequence

$$y_2 = \alpha y_1 + w_1$$

$$y_1 = \alpha y_0 + w_0$$

where the noises w_0, w_1 are iid with the distribution $N(0, \sigma^2)$, $y_0 \sim N(0, \lambda)$ is independent of the noise sequence. σ^2 and λ are known, while α is unknown.

T1. Find the MLE of the rate of return, α , given the observed price at the end of each day y_2, y_1, y_0 . In other words, compute for the value of α that maximizes $p(y_2, y_1, y_0 | \alpha)$

Hint: This is a Markov process, e.g. y_2 is independent of y_0 given y_1 . In general, a process is Markov if $p(y_n | y_{n-1}, y_{n-2}, \dots) = p(y_n | y_{n-1})$. In other words, the present is independent of the past (y_{n-2}, y_{n-3}, \dots), conditioned on the immediate past y_{n-1} . You may also find the steps of the proof for logistic regression we did in class useful.

OT1. Consider the general case, where

$$y_{n+1} = \alpha y_n + w_n, n = 0, 1, 2, \dots$$

Find the MLE given the observed price y_{N+1}, y_N, \dots, y_0

T1. We want to optimize $p(y_2, y_1, y_0 | \alpha)$

$$p(y_2, y_1, y_0 | \alpha) = p(y_2, y_1 | y_0, \alpha) \cdot p(y_0 | \alpha)$$

$$= p(y_2 | y_1, y_0, \alpha) \cdot p(y_1 | y_0, \alpha) \cdot p(y_0 | \alpha)$$

$$= p(y_2 | y_1, \alpha) \cdot p(y_1 | y_0, \alpha) \cdot p(y_0)$$

$$= N(\alpha y_1, \sigma^2) \cdot N(\alpha y_0, \sigma^2) \cdot N(0, \lambda)$$

; as y_n only dependent to y_{n-1}

and y_0 is independent to α

$$\log L(\alpha) = \log k - \frac{(y_2 - \alpha y_1)^2}{\sigma^2} \cdot \log e + \log k - \frac{(y_1 - \alpha y_0)^2}{\sigma^2} \cdot \log e + \log C ; k \text{ and } C \text{ is constant}$$

$$\text{take } \frac{\partial}{\partial \alpha} ; \quad = -\frac{2}{\sigma^2} (y_2 - \alpha y_1)(-y_1) - \frac{2}{\sigma^2} (y_1 - \alpha y_0)(-y_0)$$

$$0 = \frac{2}{\sigma^2} [y_1 y_2 - \alpha y_1^2 + y_0 y_1 - \alpha y_0^2]$$

$$\alpha = \frac{y_1 y_2 + y_0 y_1}{y_1^2 + y_0^2}$$

OT1. Consider the general case, where

$$y_{n+1} = \alpha y_n + w_n, n = 0, 1, 2, \dots$$

Find the MLE given the observed price y_{N+1}, y_N, \dots, y_0

We want to $\arg\max_{\alpha} (y_{N+1}, y_N, \dots, y_0 | \alpha)$

$$L(\alpha) = p(y_{N+1} | y_N, \alpha) \cdot p(y_N | y_{N-1}, \alpha) \cdot \dots \cdot p(y_1 | y_0, \alpha) \cdot p(y_0 | \alpha) ; \text{ As } y_N \text{ only dependent to } y_{N-1}$$

$$= N(\alpha y_N, \sigma^2) \cdot N(\alpha y_{N-1}, \sigma^2) \cdot \dots \cdot N(\alpha y_0, \sigma^2) \cdot N(0, \lambda)$$

$$\log(L(\alpha)) = \log k - \frac{(y_{N+1} - \alpha y_N)^2}{\sigma^2} + \log k - \frac{(y_N - \alpha y_{N-1})^2}{\sigma^2} + \dots + \log k - \frac{(y_1 - \alpha y_0)^2}{\sigma^2} + \log k - \frac{(y_0)^2}{\lambda}$$

to be $\frac{\partial}{\partial \alpha}$:

$$= -\frac{2}{\sigma^2} (y_{N+1} y_N - \alpha y_N^2) (-1) - \frac{2}{\sigma^2} (y_N y_{N-1} - \alpha y_{N-1}^2) (-1) - \dots - \frac{2}{\sigma^2} (y_1 y_0 - \alpha y_0^2) (-1)$$

$$\alpha (y_0^2 + y_1^2 + \dots + y_{N-1}^2 + y_N^2) = y_{N+1} y_N + y_N y_{N-1} + \dots + y_1 y_0$$

$$\alpha = \frac{y_{N+1} y_N + y_N y_{N-1} + \dots + y_1 y_0}{y_N^2 + y_{N-1}^2 + \dots + y_1^2 + y_0^2} *$$

Simple Bayes Classifier

A student in Pattern Recognition course had finally built the ultimate classifier for cat emotions. He used one input features: the amount of food the cat ate that day, x (Being a good student he already normalized x to standard Normal). He proposed the following likelihood probabilities for class 1 (happy cat) and 2 (sad cat)

$$P(x|w_1) = N(4, 2)$$

$$P(x|w_2) = N(0, 2)$$

$$x \sim N(0, 1)$$



Figure 1: The sad cat and the happy cat used in training

posterior = likelihood · prior

$$p(w_i|x) = \frac{p(x|w_i) \cdot p(w_i)}{p(x)}$$

evidence

T2. Plot the posteriors values of the two classes on the same axis. Using the likelihood ratio test, what is the decision boundary for this classifier? Assume equal prior probabilities.

T3. What happen to the decision boundary if the cat is happy with a prior of 0.75?

OT2. For the ordinary case of $P(x|w_1) = N(\mu_1, \sigma^2)$, $P(x|w_2) = N(\mu_2, \sigma^2)$, $p(w_1) = p(w_2) = 0.5$, prove that the decision boundary is at $x = \frac{\mu_1 + \mu_2}{2}$


OT3. If the student changed his model to

$$P(x|w_1) = N(4, 2)$$

$$P(x|w_2) = N(0, 4)$$

Plot the posteriors values of the two classes on the same axis. What is the decision boundary for this classifier? Assume equal prior probabilities.

$$As N(\mu, \sigma^2) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$



```

1 import random as rnd
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import numpy as np
5 from scipy import stats
6
7
8 class SimpleBayesClassifier:
9
10     def __init__(self, n_pos, n_neg):
11         """
12         Initializes the SimpleBayesClassifier with prior probabilities.
13
14         Parameters:
15         n_pos (int): The number of positive samples.
16         n_neg (int): The number of negative samples.
17
18         Returns:
19         None: This method does not return anything as it is a constructor.
20         """
21
22         self.n_pos = n_pos
23         self.n_neg = n_neg
24         self.prior_pos = n_pos / (n_pos + n_neg)
25         self.prior_neg = n_neg / (n_pos + n_neg)
26
27     def fit_params(self, x, y, n_bins=10):
28         """
29         Computes histogram-based parameters for each feature in the dataset.
30
31         Parameters:
32         x (np.ndarray): The feature matrix, where rows are samples and columns are
33         features.
34         y (np.ndarray): The target array, where each element corresponds to the
35         label of a sample.
36         n_bins (int): Number of bins to use for histogram calculation.
37
38         Returns:
39         (stay_params, leave_params): A tuple containing two lists of tuples,
40         one for 'stay' parameters and one for 'leave' parameters.
41         Each tuple in the list contains the bins and edges of the histogram for a
42         feature.
43         """
44
45         self.stay_params = []
46         self.leave_params = []
47
48         # INSERT CODE HERE
49         for i in range(x.shape[1]): # for each attributes in x_train
50             stay = x[y == 0, i]
51             stay = stay[~np.isnan(stay)]
52             Shist, Sedges = np.histogram(stay, n_bins)
53             Sedges[0] = -np.inf
54             Sedges[-1] = np.inf
55             Shist = Shist / np.sum(Shist)
56             self.stay_params.append((Shist, Sedges))
57
58             leave = x[y == 1, i]
59             leave = leave[~np.isnan(leave)]
60             Lhist, Ledges = np.histogram(leave, n_bins)

```

```

58         Ledges[0] = -np.inf
59         Ledges[-1] = np.inf
60         Lhist = Lhist / np.sum(Lhist)
61         self.leave_params.append((Lhist, Ledges))
62
63     return self.stay_params, self.leave_params
64
65     def predict(self, x, thresh=0):
66         """
67         Predicts the class labels for the given samples using the non-parametric
68         model.
69
70         Parameters:
71         x (np.ndarray): The feature matrix for which predictions are to be made.
72         thresh (float): The threshold for log probability to decide between classes.
73
74         Returns:
75         result (list): A list of predicted class labels (0 or 1) for each sample in
76         the feature matrix.
77         """
78
79         y_pred = []
80
81         # INSERT CODE HERE
82         init = np.log(self.prior_pos) - np.log(self.prior_neg)
83
84         for i in range(x.shape[0]):
85             LH = init
86             for j in range(x.shape[1]):
87                 if np.isnan(x[i][j]):
88                     continue
89                 leave_index = (
90                     np.searchsorted(self.leave_params[j][1], x[i][j], side="right")
91                     - 1
92                 )
93                 stay_index = (
94                     np.searchsorted(self.stay_params[j][1], x[i][j], side="right") -
95                     1
96                 )
97
98                 leave_value = self.leave_params[j][0][leave_index]
99                 if leave_value == 0:
100                     leave_value += 1e-6
101                 stay_value = self.stay_params[j][0][stay_index]
102                 if stay_value == 0:
103                     stay_value += 1e-6
104
105                 LH += np.log(leave_value) - np.log(stay_value)
106
107             if LH > thresh:
108                 y_pred.append(1)
109             else:
110                 y_pred.append(0)
111
112         return y_pred
113
114     def fit_gaussian_params(self, x, y):
115         """
116         Computes mean and standard deviation for each feature in the dataset.

```

```

114     Parameters:
115     x (np.ndarray): The feature matrix, where rows are samples and columns are
features.
116     y (np.ndarray): The target array, where each element corresponds to the
label of a sample.
117
118     Returns:
119     (gaussian_stay_params, gaussian_leave_params): A tuple containing two lists
of tuples,
120     one for 'stay' parameters and one for 'leave' parameters.
121     Each tuple in the list contains the mean and standard deviation for a
feature.
122     """
123
124     self.gaussian_stay_params = [(0, 0) for _ in range(x.shape[1])]
125     self.gaussian_leave_params = [(0, 0) for _ in range(x.shape[1])]
126
127     # INSERT CODE HERE
128     for i in range(x.shape[1]): # for each feature: calculate the parameter
129         stay = x[y == 0, i]
130         stay = stay[~np.isnan(stay)]
131         stay_mean = np.mean(stay)
132         stay_std = max(np.std(stay), 1e-6)
133         self.gaussian_stay_params[i] = (stay_mean, stay_std)
134
135         leave = x[y == 1, i]
136         leave = leave[~np.isnan(leave)]
137         leave_mean = np.mean(leave)
138         leave_std = max(np.std(leave), 1e-6)
139         self.gaussian_leave_params[i] = (leave_mean, leave_std)
140
141     return self.gaussian_stay_params, self.gaussian_leave_params
142
143     def gaussian_predict(self, x, thresh=0):
144         """
145         Predicts the class labels for the given samples using the parametric model.
146
147         Parameters:
148         x (np.ndarray): The feature matrix for which predictions are to be made.
149         thresh (float): The threshold for log probability to decide between classes.
150
151         Returns:
152         result (list): A list of predicted class labels (0 or 1) for each sample in
the feature matrix.
153         """
154
155         y_pred = []
156
157         # INSERT CODE HERE
158         for i in range(x.shape[0]):
159             predict = np.log(self.prior_pos) - np.log(self.prior_neg)
160             for j in range(x.shape[1]):
161                 if np.isnan(x[i][j]):
162                     continue
163
164                 log_leave = max(
165                     stats.norm(
166                         self.gaussian_leave_params[j][0],
167                         self.gaussian_leave_params[j][1],
168

```

```
169         1e-9,
170     )
171     log_stay = max(
172         stats.norm(
173             self.gaussian_stay_params[j][0],
174             self.gaussian_stay_params[j][1]
175         ).pdf(x[i][j]),
176         1e-9,
177     )
178     predict += np.log(log_leave) - np.log(log_stay)
179
180     if predict > thresh:
181         y_pred.append(1)
182     else:
183         y_pred.append(0)
184
185     return np.array(y_pred)
186
```


Pattern_HW2_student_2026

January 28, 2026

1 Employee Attrition Prediction

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

1.0.1 read CSV

```
[2]: df = pd.read_csv('hr-employee-attrition-with-null.csv')
```

1.0.2 Dataset statistic

```
[3]: df.describe()
```

```
[3]:
```

	Unnamed: 0	Age	DailyRate	DistanceFromHome	Education	\
count	1470.000000	1176.000000	1176.000000	1176.000000	1176.000000	
mean	734.500000	37.134354	798.875850	9.37500	2.920918	
std	424.496761	9.190317	406.957684	8.23049	1.028796	
min	0.000000	18.000000	102.000000	1.00000	1.000000	
25%	367.250000	30.000000	457.750000	2.00000	2.000000	
50%	734.500000	36.000000	798.500000	7.00000	3.000000	
75%	1101.750000	43.000000	1168.250000	15.00000	4.000000	
max	1469.000000	60.000000	1499.000000	29.00000	5.000000	

	EmployeeCount	EmployeeNumber	EnvironmentSatisfaction	HourlyRate	\
count	1176.0	1176.000000	1176.000000	1176.000000	
mean	1.0	1031.399660	2.733844	65.821429	
std	0.0	601.188955	1.092992	20.317323	
min	1.0	1.000000	1.000000	30.000000	
25%	1.0	494.750000	2.000000	48.000000	
50%	1.0	1027.500000	3.000000	66.000000	
75%	1.0	1562.250000	4.000000	84.000000	
max	1.0	2068.000000	4.000000	100.000000	

	JobInvolvement	...	RelationshipSatisfaction	StandardHours	\
count	1176.000000	...	1176.000000	1176.0	
mean	2.728741	...	2.694728	80.0	
std	0.705280	...	1.093660	0.0	

min	1.000000	...	1.000000	80.0
25%	2.000000	...	2.000000	80.0
50%	3.000000	...	3.000000	80.0
75%	3.000000	...	4.000000	80.0
max	4.000000	...	4.000000	80.0

	StockOptionLevel	TotalWorkingYears	TrainingTimesLastYear	\
count	1176.000000	1176.000000	1176.000000	
mean	0.752551	11.295068	2.787415	
std	0.822550	7.783376	1.290507	
min	0.000000	0.000000	0.000000	
25%	0.000000	6.000000	2.000000	
50%	1.000000	10.000000	3.000000	
75%	1.000000	15.000000	3.000000	
max	3.000000	40.000000	6.000000	

	WorkLifeBalance	YearsAtCompany	YearsInCurrentRole	\
count	1176.000000	1176.000000	1176.000000	
mean	2.770408	7.067177	4.290816	
std	0.705004	6.127836	3.630901	
min	1.000000	0.000000	0.000000	
25%	2.000000	3.000000	2.000000	
50%	3.000000	5.000000	3.000000	
75%	3.000000	10.000000	7.000000	
max	4.000000	37.000000	18.000000	

	YearsSinceLastPromotion	YearsWithCurrManager
count	1176.000000	1176.000000
mean	2.159014	4.096939
std	3.163524	3.537393
min	0.000000	0.000000
25%	0.000000	2.000000
50%	1.000000	3.000000
75%	2.250000	7.000000
max	15.000000	17.000000

[8 rows x 27 columns]

```
[4]: df.head()
```

```
[4]: Unnamed: 0  Age  Attrition  BusinessTravel  DailyRate  \
0           0  41.0      Yes      Travel_Rarely      NaN
1           1   NaN      No                NaN      279.0
2           2  37.0      Yes                NaN      1373.0
3           3   NaN      No  Travel_Frequently      1392.0
4           4  27.0      No      Travel_Rarely      591.0
```

	Department	DistanceFromHome	Education	EducationField	\
0	NaN	1.0	NaN	Life Sciences	
1	Research & Development	NaN	NaN	Life Sciences	
2	NaN	2.0	2.0	NaN	
3	Research & Development	3.0	4.0	Life Sciences	
4	Research & Development	2.0	1.0	Medical	

	EmployeeCount	...	RelationshipSatisfaction	StandardHours	\
0	1.0	...	1.0	80.0	
1	1.0	...	4.0	NaN	
2	1.0	...	NaN	80.0	
3	NaN	...	3.0	NaN	
4	1.0	...	4.0	80.0	

	StockOptionLevel	TotalWorkingYears	TrainingTimesLastYear	WorkLifeBalance	\
0	0.0	8.0	0.0	NaN	
1	1.0	10.0	NaN	3.0	
2	0.0	7.0	3.0	NaN	
3	NaN	8.0	3.0	NaN	
4	1.0	6.0	NaN	3.0	

	YearsAtCompany	YearsInCurrentRole	YearsSinceLastPromotion	\
0	6.0	NaN	0.0	
1	10.0	NaN	NaN	
2	NaN	0.0	NaN	
3	8.0	NaN	3.0	
4	2.0	2.0	2.0	

	YearsWithCurrManager
0	NaN
1	7.0
2	0.0
3	0.0
4	NaN

[5 rows x 36 columns]

1.0.3 Feature transformation

```
[5]: df.loc[df["Attrition"] == "no", "Attrition"] = 0.0
df.loc[df["Attrition"] == "yes", "Attrition"] = 1.0
string_categorical_col = ['Department', 'Attrition', 'BusinessTravel', 'EducationField', 'Gender', 'JobRole', 'MaritalStatus', 'Over18', 'OverTime']

# ENCODE STRING COLUMNS TO CATEGORICAL COLUMNS
for col in string_categorical_col:
```

```

# INSERT CODE HERE
df[col] = pd.Categorical(df[col]).codes

# HANDLE NULL NUMBERS
# I don't think we need to handle null?
# INSERT CODE HERE

df = df.loc[:, ~df.columns.isin(['EmployeeNumber', 'Unnamed: 0',
↪ 'EmployeeCount', 'StandardHours', 'Over18'])]

```

1.0.4 Splitting data into train and test

```

[6]: from sklearn.model_selection import train_test_split

[7]: X = df.drop(["Attrition"], axis=1)
     Y = df["Attrition"]

     x_train, x_test, y_train, y_test = train_test_split(X, Y, stratify=Y,
↪ test_size=0.1, random_state=12345)

```

1.0.5 Display histogram of each feature

```

[8]: def display_histogram(df, col_name, n_bin = 40):
     # INSERT CODE HERE
     col_nonan = df[col_name][~np.isnan(df[col_name])]

     # col = np.array(df[col_name])
     # col_nonan = np.array( col[~ np.isnan(col)] )


     # hist is the count for each bin
     # bin_edge is the edge values of the bins
     hist, bin_edge = np.histogram(col_nonan, n_bin)
     bin_edge[0] = -np.inf
     bin_edge[-1] = np.inf


     # plot the histogram
     plt.fill_between(bin_edge.repeat(2)[1:-1], hist.repeat(2),
↪ facecolor='steelblue')
     plt.title(col_name + " with " + str(n_bin) + " bins")
     plt.show()


```

1.0.6 T4. Observe the histogram for Age, MonthlyIncome and DistanceFromHome. How many bins have zero counts? Do you think this is a good discretization? Why?

```
[9]: display_histogram(x_train, "Age")  
display_histogram(x_train, "MonthlyIncome")  
display_histogram(x_train, "DistanceFromHome")
```








1.0.7 T5. Can we use a Gaussian to estimate this histogram? Why? What about a Gaussian Mixture Model (GMM)?


As 40 bins, Age and MonthlyIncome has no zero counts bin, but DistanceFromHome have 11 bins with zero counts. I don't think this is a good discretization because there are zero counts bin which lead to probability of 0 due to less data.

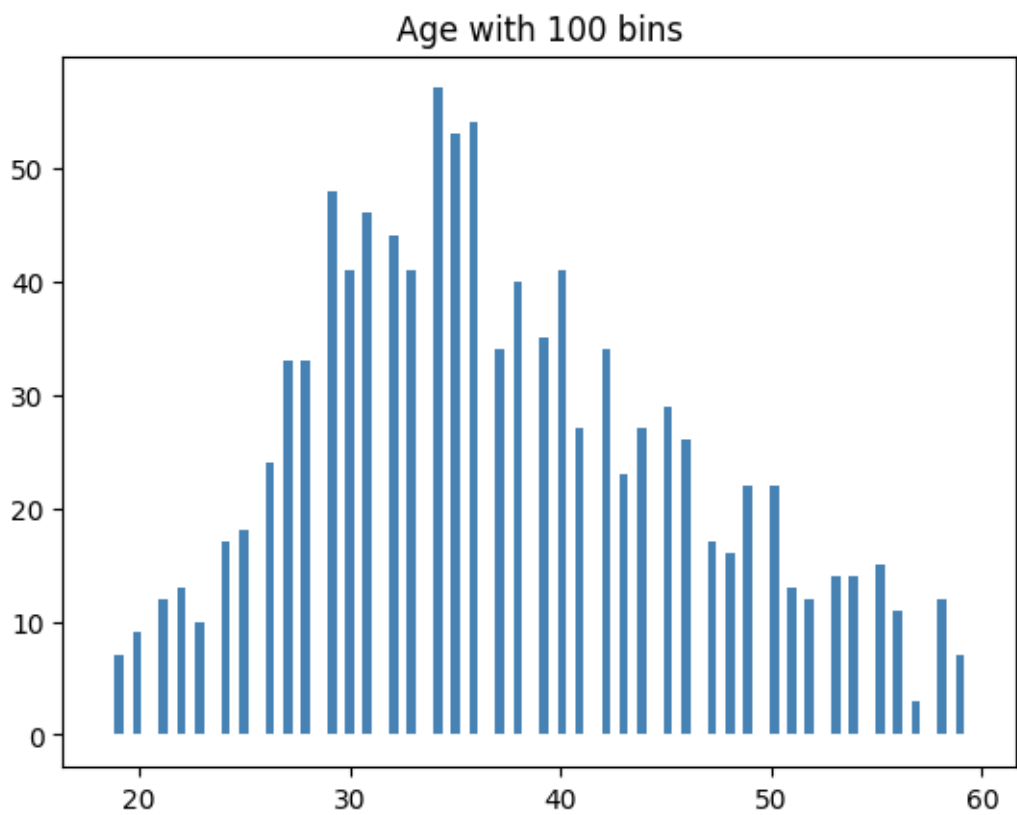
1.0.8 T6. Now plot the histogram according to the method described above (with 10, 40, and 100 bins) and show 3 plots each for Age, MonthlyIncome, and DistanceFromHome. Which bin size is most sensible for each features? Why?


```
[11]: columns = ["Age", "MonthlyIncome", "DistanceFromHome"]
      num_bins = [10, 40, 100]

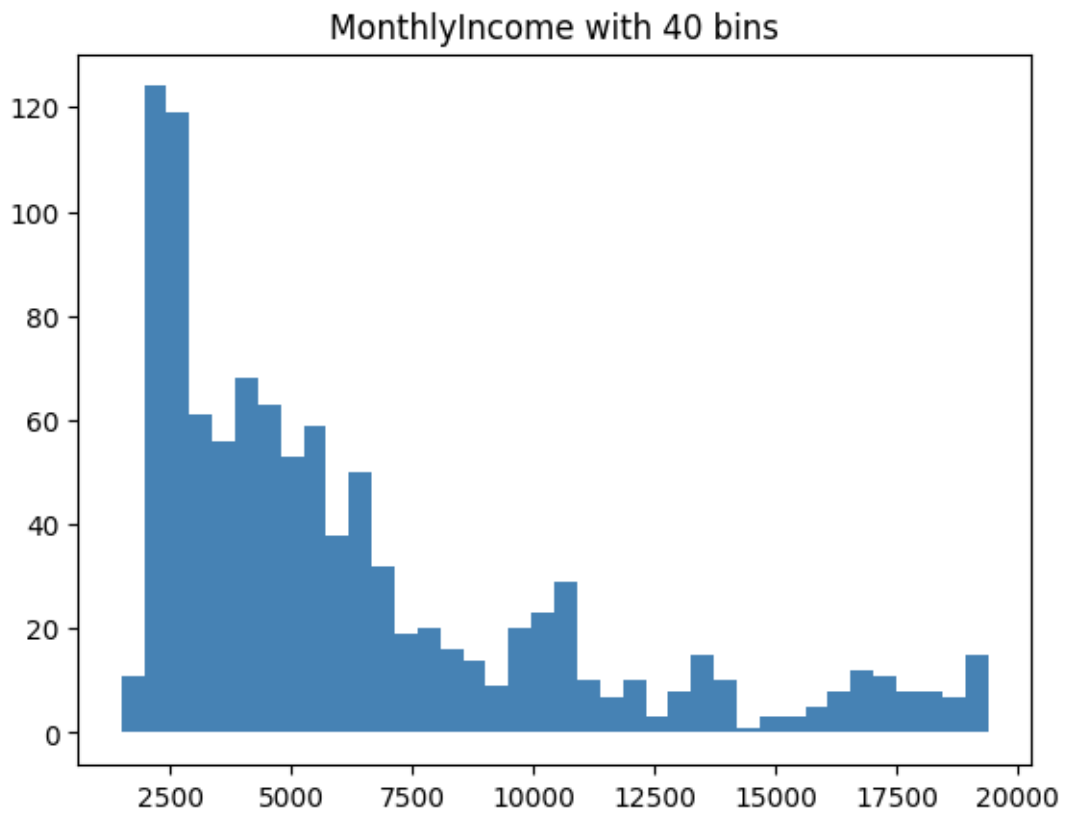
      for col in columns:
          for num in num_bins:
              display_histogram(x_train, col, num)
```










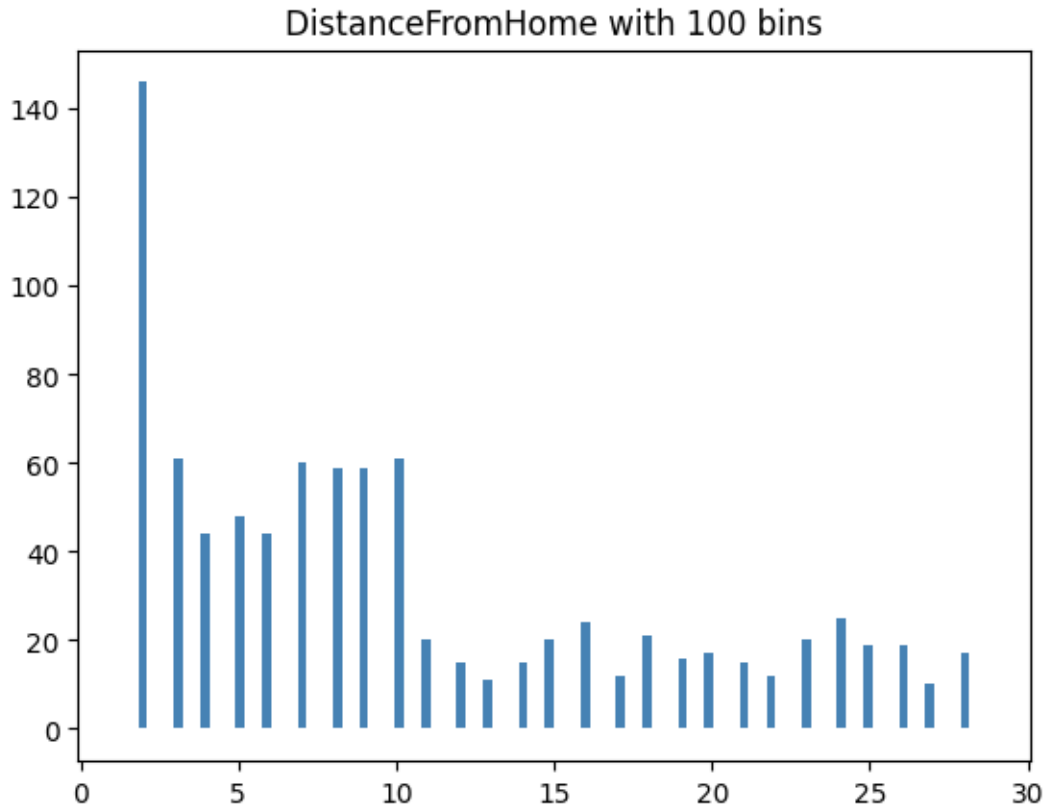












For Age and MonthlyIncome:

I think 40 bins is the best because there are no zero counts. Well, 10 bins also doesn't have zero counts too, but you can elaborate more details in the 40 bins. Anyway, the 100 bins got some zero counts, so it isn't a good discretization.

For DistanceFromHome:

40 bins and 100 bins will both have a zero counts bin. So that, the only good number of bins in these three are 10.

1.0.9 T7. For the rest of the features, which one should be discretized in order to be modeled by histograms? What are the criteria for choosing whether we should discretize a feature or not? Answer this and discretize those features into 10 bins each. In other words, figure out the bin edge for each feature, then use `digitize()` to convert the features to discrete values

First, the encoded from categorical values shouldn't be discretized because the value is already discretized via the encoded function. For me, what should be discretized is the one with continuous range value or have many unique values. The threshold for considering if that features should be discretized or not is > 40 (or any suit number) unique values.


```

[12]: def hist(array, col_name, n_bin=10):
        nonan = array[~np.isnan(array)]

        # hist is the count for each bin
        # bin_edge is the edge values of the bins
        hist, bin_edges = np.histogram(nonan, n_bin)
        bin_edges[0] = -np.inf
        bin_edges[-1] = np.inf

        bin_indices = np.full_like(array, np.nan, dtype=float)
        bin_indices[~np.isnan(array)] = np.digitize(nonan, bin_edges)

        # plot the histogram
        plt.fill_between(bin_edges.repeat(2)[1:-1], hist.repeat(2),
        ↪facecolor='steelblue')
        plt.title(col_name + " with " + str(n_bin) + " bins")
        plt.show()


        return bin_indices, bin_edges

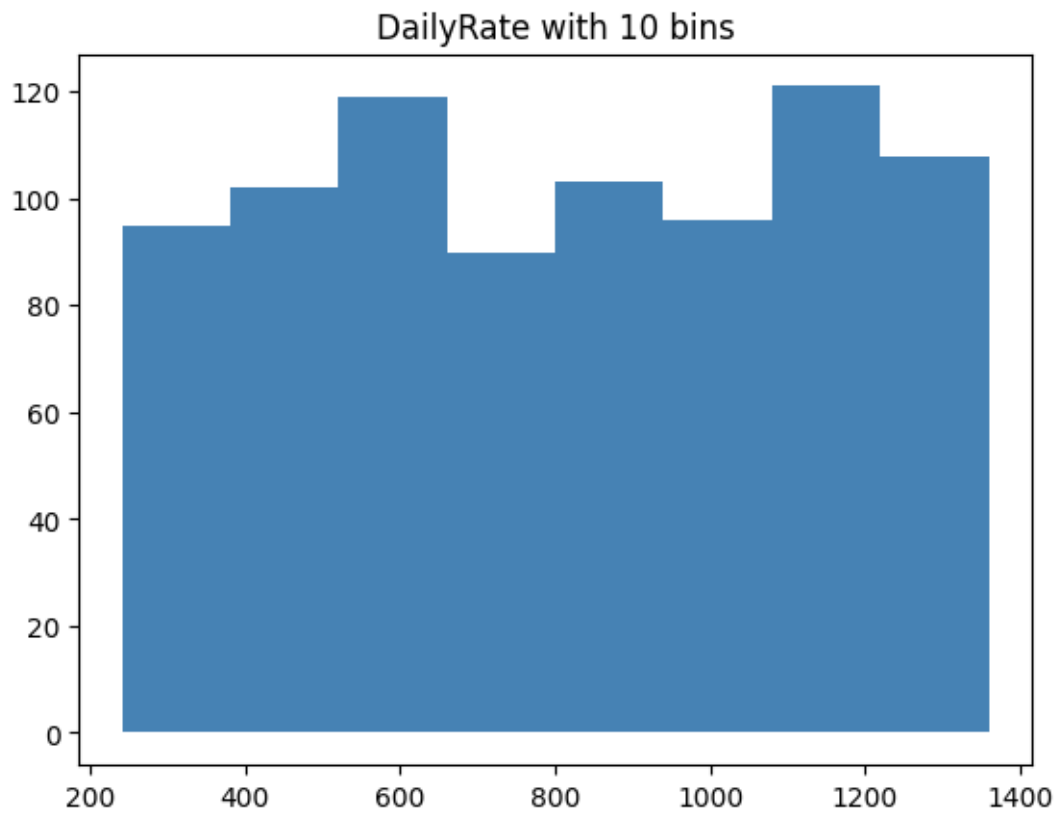
discretize = []


for col in x_train.columns:
    if (x_train[col].nunique() > 10):
        x_train[col], _ = hist(x_train[col], col)
        discretize.append(col)


print(discretize)


```




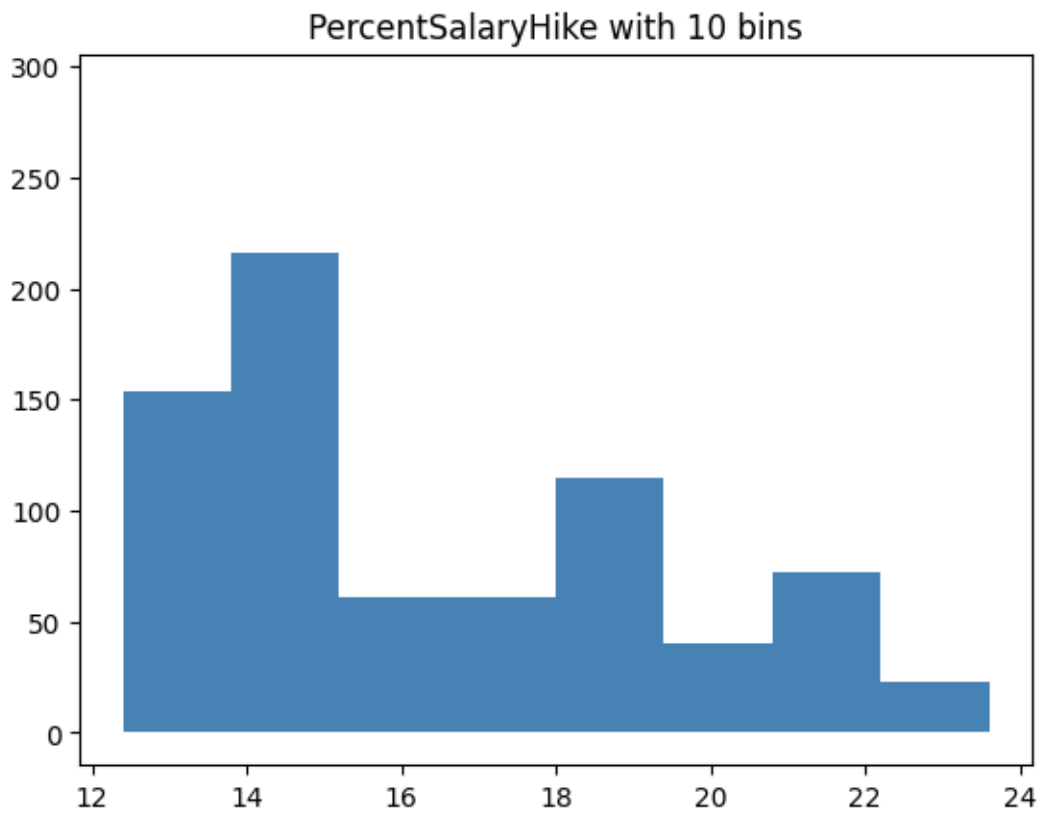


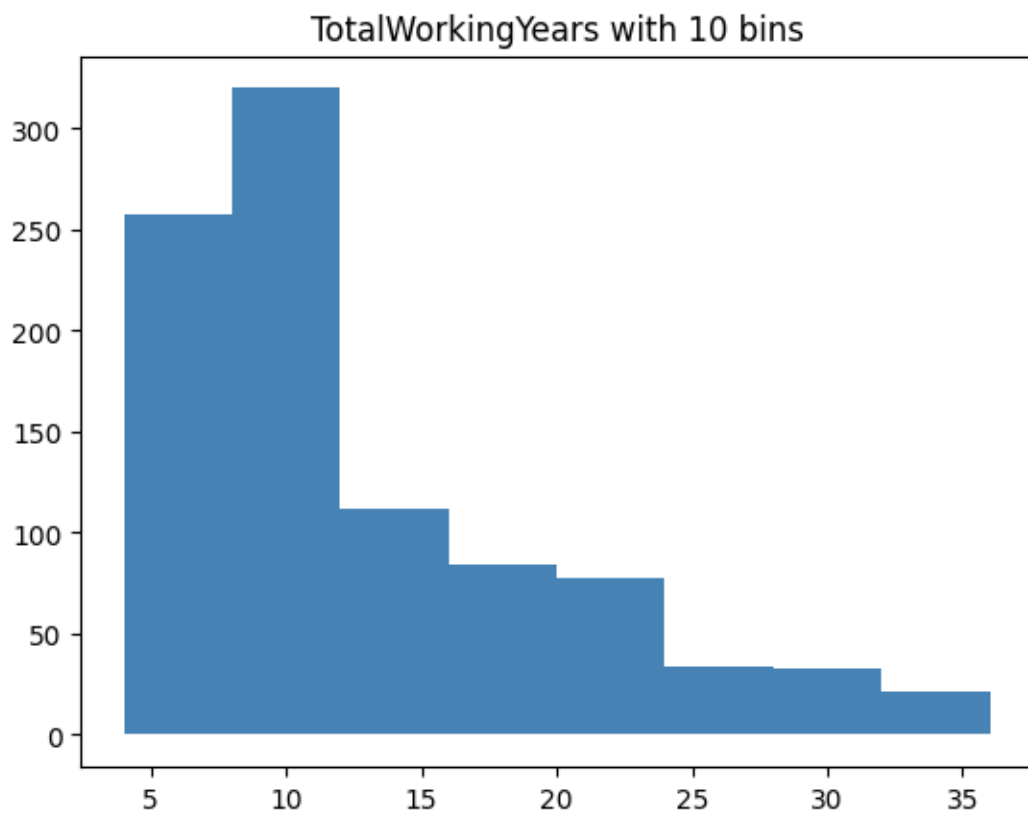
























```
[ 'Age', 'DailyRate', 'DistanceFromHome', 'HourlyRate', 'MonthlyIncome',
  'MonthlyRate', 'PercentSalaryHike', 'TotalWorkingYears', 'YearsAtCompany',
  'YearsInCurrentRole', 'YearsSinceLastPromotion', 'YearsWithCurrManager']
```

1.0.10 T8. What kind of distribution should we use to model histograms? (Answer a distribution name) What is the MLE for the likelihood distribution? (Describe how to do the MLE). Plot the likelihood distributions of MonthlyIncome, JobRole, HourlyRate, and MaritalStatus for different Attrition values.

Multinomial distribution. As we mapped the value into different bins, each bin can be treated as one category. The MLE for the probability of each bin = number of data in that bin / number of all data

```
[13]: def plot_likelihood(x_train, y_train, col, n_bin=10):
        nonan_stay = x_train[y_train == 0][col].dropna()
        nonan_leave = x_train[y_train == 1][col].dropna()

        hist, bin_edges = np.histogram(nonan_stay)
        hist = hist / nonan_stay.shape[0]
        # plot the histogram
```

```


    plt.fill_between(bin_edges.repeat(2)[1:-1], hist.repeat(2),
↳facecolor='steelblue')
    plt.title(f"{col} - Stay")
    plt.show()

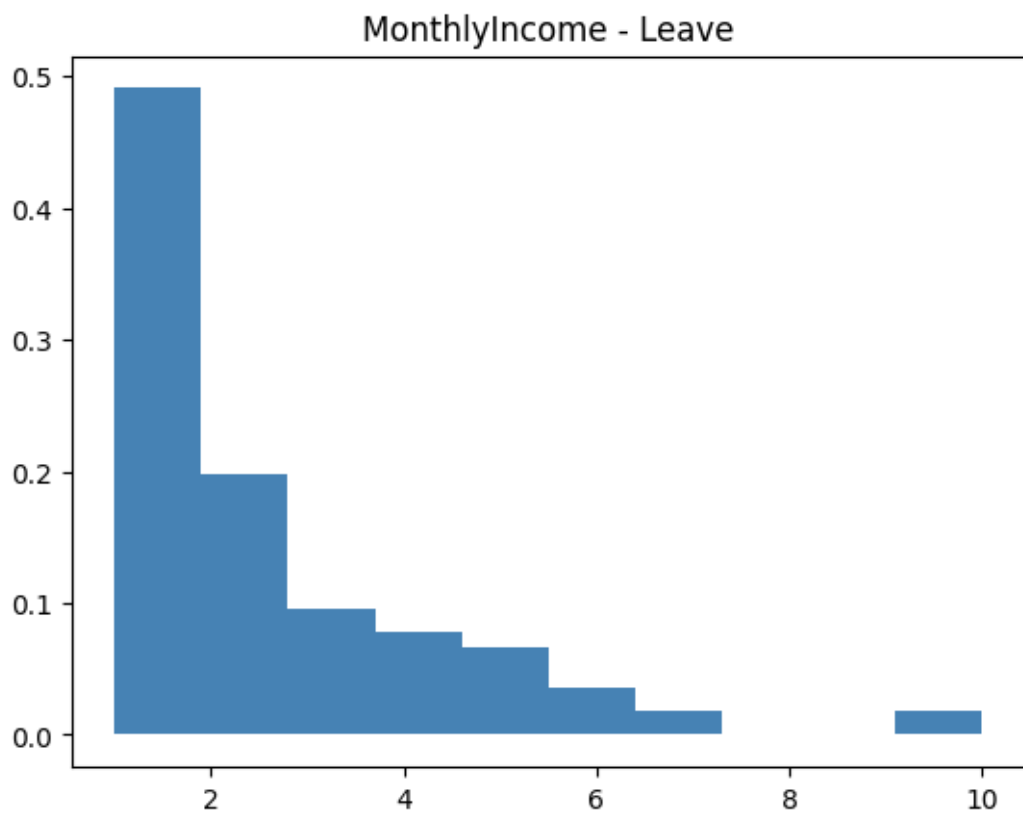
    hist, bin_edges = np.histogram(nonan_leave)
    hist = hist / nonan_leave.shape[0]
    # plot the histogram
    plt.fill_between(bin_edges.repeat(2)[1:-1], hist.repeat(2),
↳facecolor='steelblue')
    plt.title(f"{col} - Leave")
    plt.show()

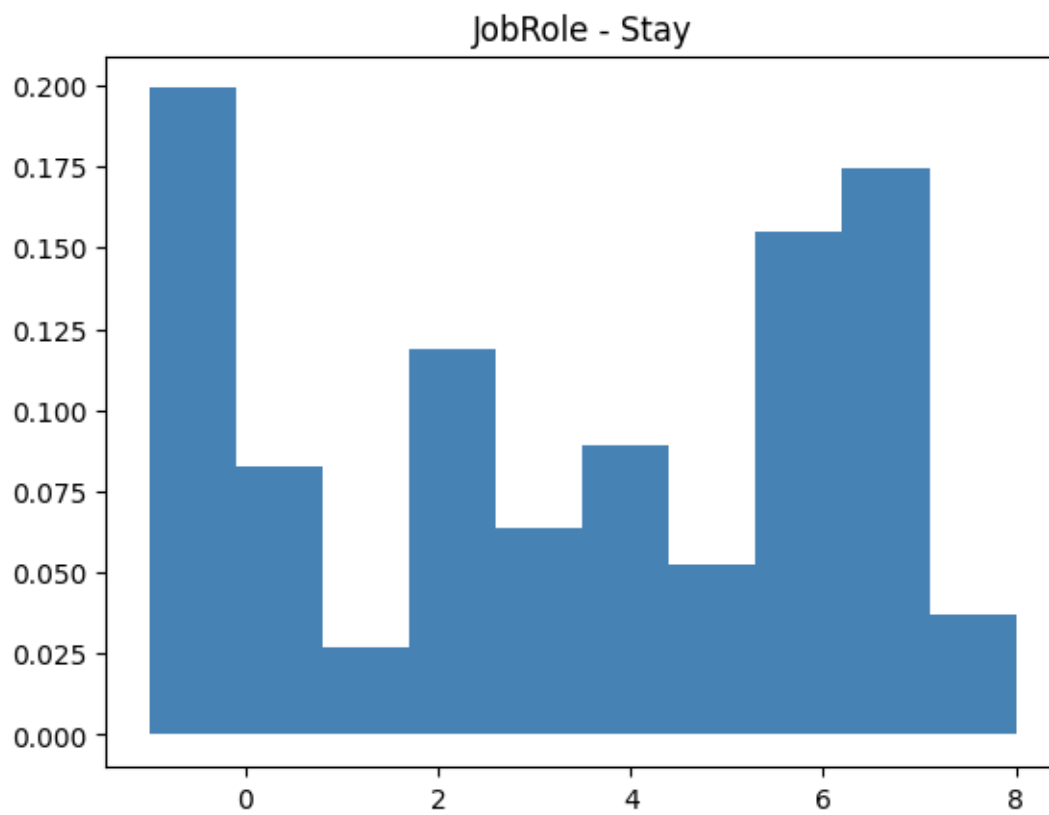
cols = ["MonthlyIncome", "JobRole", "HourlyRate", "MaritalStatus"]


for col in cols:
    plot_likelihood(x_train, y_train, col)


```




















1.0.11 T9. What is the prior distribution of the two classes?

```
[14]: stay = np.sum(y_train == 0) / y_train.shape[0]
      leave = np.sum(y_train == 1) / y_train.shape[0]

      print("Stay:", stay)
      print("Leave:", leave)
```

Stay: 0.8390022675736961

Leave: 0.16099773242630386

1.0.12 T10. If we use the current Naive Bayes with our current Maximum Likelihood Estimates, we will find that some $P(x_i | \text{attrition})$ will be zero and will result in the entire product term to be zero. Propose a method to fix this problem.

Inserting some small values (epsilon) would help preventing the probability to be zero.

1.0.13 T11. Implement your Naive Bayes classifier. Use the learned distributions to classify the test set. Don't forget to allow your classifier to handle missing values in the test set. Report the overall Accuracy. Then, report the Precision, Recall, and F score for detecting attrition. See Lecture 1 for the definitions of each metric.

```
[20]: import importlib, SimpleBayesClassifier
importlib.reload(SimpleBayesClassifier)
from SimpleBayesClassifier import SimpleBayesClassifier
```

```
[21]: model = SimpleBayesClassifier(n_pos = np.sum(y_train == 1), n_neg = np.
    ↪sum(y_train == 0))
```

```
[23]: def check_prior():
    """
    This function designed to test the implementation of the prior probability
    ↪calculation in a Naive Bayes classifier.
    Specifically, it checks if the classifier correctly computes the prior
    ↪probabilities for the
    negative and positive classes based on given input counts.
    """

    # prior_neg = 5/(5 + 5) = 0.5 and # prior_pos = 5/(5 + 5) = 0.5
    assert (SimpleBayesClassifier(5, 5).prior_pos, SimpleBayesClassifier(5, 5).
    ↪prior_neg) == (0.5, 0.5)

    assert (SimpleBayesClassifier(3, 5).prior_pos, SimpleBayesClassifier(3, 5).
    ↪prior_neg) == (0.375, 0.625)
    assert (SimpleBayesClassifier(0, 1).prior_pos, SimpleBayesClassifier(0, 1).
    ↪prior_neg) == (0, 1)
    assert (SimpleBayesClassifier(1, 0).prior_pos, SimpleBayesClassifier(1, 0).
    ↪prior_neg) == (1, 0)

    check_prior()
```

```
[25]: model.fit_params(np.array(x_train), np.array(y_train))
```

```
[25]: ((array([0.02581369, 0.05499439, 0.14253648, 0.17171717, 0.18181818,
    0.16498316, 0.09876543, 0.07182941, 0.05387205, 0.03367003]),
    array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf])),
    (array([0.2027027, 0., 0., 0.08558559, 0., 0., 0.13153153, 0., 0., 0.58018018]),
    array([-inf, -0.7, -0.4, -0.1, 0.2, 0.5, 0.8, 1.1, 1.4, 1.7, inf])),
    (array([0.10961969, 0.08053691, 0.098434, 0.11297539, 0.08501119,
    0.09619687, 0.0950783, 0.11409396, 0.10738255, 0.10067114]),
    array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf])),
    (array([0.19099099, 0., 0., 0.03603604, 0., 0., 0.03603604, 0., 0., 0.03603604, 0.])))
```

```

0.          , 0.54594595, 0.          , 0.          , 0.22702703]],
array([-inf, -0.7, -0.4, -0.1, 0.2, 0.5, 0.8, 1.1, 1.4, 1.7, inf])),
(array([0.34698521, 0.13083049, 0.17406143, 0.08987486, 0.02275313,
0.04778157, 0.05119454, 0.04095563, 0.05005688, 0.04550626])),
array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf])),
(array([0.10946408, 0.          , 0.19156214, 0.          , 0.          ,
0.38312429, 0.          , 0.28164196, 0.          , 0.03420753])),
array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.19459459, 0.01531532, 0.          , 0.33693694, 0.          ,
0.07567568, 0.26666667, 0.          , 0.04234234, 0.06846847])),
array([-inf, -0.4, 0.2, 0.8, 1.4, 2. , 2.6, 3.2, 3.8, 4.4, inf])),
(array([0.18423973, 0.          , 0.          , 0.18201998, 0.          ,
0.          , 0.32741398, 0.          , 0.          , 0.3063263 ])),
array([-inf, 1.3, 1.6, 1.9, 2.2, 2.5, 2.8, 3.1, 3.4, 3.7, inf])),
(array([0.2009009 , 0.          , 0.          , 0.          , 0.          ,
0.33873874, 0.          , 0.          , 0.          , 0.46036036])),
array([-inf, -0.8, -0.6, -0.4, -0.2, 0. , 0.2, 0.4, 0.6, 0.8, inf])),
(array([0.08520179, 0.10426009, 0.0941704 , 0.11098655, 0.09529148,
0.08632287, 0.10426009, 0.09977578, 0.09977578, 0.11995516])),
array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf])),
(array([0.0407701 , 0.          , 0.          , 0.26274066, 0.          ,
0.          , 0.59116648, 0.          , 0.          , 0.10532276])),
array([-inf, 1.3, 1.6, 1.9, 2.2, 2.5, 2.8, 3.1, 3.4, 3.7, inf])),
(array([0.31513083, 0.          , 0.37997725, 0.          , 0.          ,
0.16382253, 0.          , 0.09215017, 0.          , 0.04891923])),
array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.1990991 , 0.08288288, 0.02702703, 0.11891892, 0.06396396,
0.08918919, 0.05225225, 0.15495495, 0.17477477, 0.03693694])),
array([-inf, -0.1, 0.8, 1.7, 2.6, 3.5, 4.4, 5.3, 6.2, 7.1, inf])),
(array([0.18459796, 0.          , 0.          , 0.19705549, 0.          ,
0.          , 0.29331823, 0.          , 0.          , 0.32502831])),
array([-inf, 1.3, 1.6, 1.9, 2.2, 2.5, 2.8, 3.1, 3.4, 3.7, inf])),
(array([0.19099099, 0.          , 0.          , 0.1954955 , 0.          ,
0.          , 0.38108108, 0.          , 0.          , 0.23243243])),
array([-inf, -0.7, -0.4, -0.1, 0.2, 0.5, 0.8, 1.1, 1.4, 1.7, inf])),
(array([0.2073991 , 0.24103139, 0.20627803, 0.08295964, 0.06165919,
0.05605381, 0.03699552, 0.01345291, 0.04372197, 0.05044843])),
array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf])),
(array([0.09775281, 0.11460674, 0.10449438, 0.09662921, 0.0988764 ,
0.09438202, 0.09438202, 0.10674157, 0.11348315, 0.07865169])),
array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf])),
(array([0.13718821, 0.33446712, 0.10544218, 0.12131519, 0.1031746 ,
0.03741497, 0.04421769, 0.04875283, 0.03061224, 0.03741497])),
array([-inf, 0.9, 1.8, 2.7, 3.6, 4.5, 5.4, 6.3, 7.2, 8.1, inf])),
(array([0.19189189, 0.          , 0.          , 0.          , 0.          ,
0.62162162, 0.          , 0.          , 0.          , 0.18648649])),
array([-inf, -0.8, -0.6, -0.4, -0.2, 0. , 0.2, 0.4, 0.6, 0.8, inf])),

```

```

(array([0.26131222, 0.14253394, 0.21266968, 0.05769231, 0.05656109,
        0.11312217, 0.03846154, 0.07126697, 0.02036199, 0.0260181 ]),
 array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf])),
(array([0.85310734, 0.          , 0.          , 0.          , 0.          ,
        0.          , 0.          , 0.          , 0.          , 0.14689266])),
 array([-inf, 3.1, 3.2, 3.3, 3.4, 3.5, 3.6, 3.7, 3.8, 3.9, inf])),
(array([0.1868743 , 0.          , 0.          , 0.20244716, 0.          ,
        0.          , 0.30700779, 0.          , 0.          , 0.30367075])),
 array([-inf, 1.3, 1.6, 1.9, 2.2, 2.5, 2.8, 3.1, 3.4, 3.7, inf])),
(array([0.39595051, 0.          , 0.          , 0.44769404, 0.          ,
        0.          , 0.10686164, 0.          , 0.          , 0.04949381])),
 array([-inf, 0.3, 0.6, 0.9, 1.2, 1.5, 1.8, 2.1, 2.4, 2.7, inf])),
(array([0.08017817, 0.23273942, 0.31514477, 0.11024499, 0.08240535,
        0.08017817, 0.03452116, 0.03452116, 0.02227171, 0.0077951 ]),
 array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf])),
(array([0.03703704, 0.05274972, 0.          , 0.36026936, 0.          ,
        0.332211  , 0.07856341, 0.          , 0.09539843, 0.04377104])),
 array([-inf, 0.6, 1.2, 1.8, 2.4, 3.  , 3.6, 4.2, 4.8, 5.4, inf])),
(array([0.04519774, 0.          , 0.          , 0.23050847, 0.          ,
        0.          , 0.62146893, 0.          , 0.          , 0.10282486])),
 array([-inf, 1.3, 1.6, 1.9, 2.2, 2.5, 2.8, 3.1, 3.4, 3.7, inf])),
(array([0.27160494, 0.34006734, 0.20089787, 0.06734007, 0.03142536,
        0.04938272, 0.02132435, 0.00448934, 0.00561167, 0.00785634])),
 array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf])),
(array([0.17502787, 0.34225195, 0.09587514, 0.18617614, 0.07246377,
        0.07134894, 0.02564103, 0.01672241, 0.01003344, 0.00445931])),
 array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf])),
(array([0.62107623, 0.11210762, 0.07959641, 0.03363229, 0.07847534,
        0.01457399, 0.01345291, 0.02017937, 0.01345291, 0.01345291])),
 array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf])),
(array([0.18459796, 0.34994337, 0.08720272, 0.01812005, 0.2400906 ,
        0.06568516, 0.01585504, 0.02491506, 0.00792752, 0.00566251])),
 array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf]]),
[(array([0.13294798, 0.11560694, 0.16184971, 0.20231214, 0.10982659,
        0.07514451, 0.06358382, 0.05202312, 0.04046243, 0.04624277])),
 array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf])),
(array([0.19248826, 0.          , 0.          , 0.03286385, 0.          ,
        0.          , 0.23004695, 0.          , 0.          , 0.54460094])),
 array([-inf, -0.7, -0.4, -0.1, 0.2, 0.5, 0.8, 1.1, 1.4, 1.7, inf])),
(array([0.08917197, 0.14649682, 0.08917197, 0.11464968, 0.08917197,
        0.10828025, 0.07006369, 0.12101911, 0.07643312, 0.0955414 ]),
 array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf])),
(array([0.22535211, 0.          , 0.          , 0.05164319, 0.          ,
        0.          , 0.41314554, 0.          , 0.          , 0.30985915])),
 array([-inf, -0.7, -0.4, -0.1, 0.2, 0.5, 0.8, 1.1, 1.4, 1.7, inf])),
(array([0.25988701, 0.11864407, 0.14124294, 0.0960452 , 0.03389831,
        0.07909605, 0.05084746, 0.06214689, 0.10734463, 0.05084746])),

```



```

array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf])),
(array([0.13068182, 0.          , 0.19886364, 0.          , 0.          ,
        0.41477273, 0.          , 0.22727273, 0.          , 0.02840909])),
array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.18309859, 0.02816901, 0.          , 0.29577465, 0.          ,
        0.12676056, 0.22535211, 0.          , 0.03755869, 0.10328638])),
array([-inf, -0.4, 0.2, 0.8, 1.4, 2. , 2.6, 3.2, 3.8, 4.4, inf])),
(array([0.26875, 0.          , 0.          , 0.2          , 0.          , 0.          ,
        0.          , 0.          , 0.29375])),
array([-inf, 1.3, 1.6, 1.9, 2.2, 2.5, 2.8, 3.1, 3.4, 3.7, inf])),
(array([0.19248826, 0.          , 0.          , 0.          , 0.          ,
        0.29577465, 0.          , 0.          , 0.          , 0.51173709])),
array([-inf, -0.8, -0.6, -0.4, -0.2, 0. , 0.2, 0.4, 0.6, 0.8, inf])),
(array([0.11242604, 0.08284024, 0.13017751, 0.10650888, 0.09467456,
        0.10650888, 0.07100592, 0.0887574 , 0.0887574 , 0.1183432 ])),
array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf])),
(array([0.1091954 , 0.          , 0.          , 0.31034483, 0.          ,
        0.          , 0.52873563, 0.          , 0.          , 0.05172414])),
array([-inf, 1.3, 1.6, 1.9, 2.2, 2.5, 2.8, 3.1, 3.4, 3.7, inf])),
(array([0.60795455, 0.          , 0.23295455, 0.          , 0.          ,
        0.11931818, 0.          , 0.01704545, 0.          , 0.02272727])),
array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.18779343, 0.02816901, 0.04694836, 0.19248826, 0.01877934,
        0.03755869, 0.00469484, 0.17370892, 0.18779343, 0.12206573])),
array([-inf, -0.1, 0.8, 1.7, 2.6, 3.5, 4.4, 5.3, 6.2, 7.1, inf])),
(array([0.24          , 0.          , 0.          , 0.2          , 0.          ,
        0.          , 0.30857143, 0.          , 0.          , 0.25142857])),
array([-inf, 1.3, 1.6, 1.9, 2.2, 2.5, 2.8, 3.1, 3.4, 3.7, inf])),
(array([0.22065728, 0.          , 0.          , 0.11267606, 0.          ,
        0.          , 0.2629108 , 0.          , 0.          , 0.40375587])),
array([-inf, -0.7, -0.4, -0.1, 0.2, 0.5, 0.8, 1.1, 1.4, 1.7, inf])),
(array([0.49101796, 0.19760479, 0.09580838, 0.07784431, 0.06586826,
        0.03592814, 0.01796407, 0.          , 0.          , 0.01796407])),
array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf])),
(array([0.13333333, 0.06666667, 0.08484848, 0.13939394, 0.08484848,
        0.09090909, 0.11515152, 0.10909091, 0.08484848, 0.09090909])),
array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf])),
(array([0.0960452 , 0.38983051, 0.07909605, 0.06779661, 0.08474576,
        0.06779661, 0.05649718, 0.06779661, 0.02824859, 0.06214689])),
array([-inf, 0.9, 1.8, 2.7, 3.6, 4.5, 5.4, 6.3, 7.2, 8.1, inf])),
(array([0.23474178, 0.          , 0.          , 0.          , 0.          ,
        0.36150235, 0.          , 0.          , 0.          , 0.40375587])),
array([-inf, -0.8, -0.6, -0.4, -0.2, 0. , 0.2, 0.4, 0.6, 0.8, inf])),
(array([0.33519553, 0.15642458, 0.15642458, 0.05586592, 0.06145251,
        0.08379888, 0.03351955, 0.05027933, 0.02793296, 0.03910615])),
array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf])),
(array([0.83529412, 0.          , 0.          , 0.          , 0.          ,

```

```

0.          , 0.          , 0.          , 0.          , 0.16470588]],
array([-inf, 3.1, 3.2, 3.3, 3.4, 3.5, 3.6, 3.7, 3.8, 3.9, inf])),
(array([0.26219512, 0.          , 0.          , 0.20121951, 0.          ,
0.          , 0.2804878 , 0.          , 0.          , 0.25609756])),
array([-inf, 1.3, 1.6, 1.9, 2.2, 2.5, 2.8, 3.1, 3.4, 3.7, inf])),
(array([0.65116279, 0.          , 0.          , 0.22674419, 0.          ,
0.          , 0.06395349, 0.          , 0.          , 0.05813953])),
array([-inf, 0.3, 0.6, 0.9, 1.2, 1.5, 1.8, 2.1, 2.4, 2.7, inf])),
(array([0.25454545, 0.2969697 , 0.22424242, 0.07878788, 0.06060606,
0.03636364, 0.01818182, 0.01212121, 0.00606061, 0.01212121])),
array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf])),
(array([0.05202312, 0.05202312, 0.          , 0.42196532, 0.          ,
0.28901734, 0.10404624, 0.          , 0.06358382, 0.01734104])),
array([-inf, 0.6, 1.2, 1.8, 2.4, 3. , 3.6, 4.2, 4.8, 5.4, inf])),
(array([0.10588235, 0.          , 0.          , 0.23529412, 0.          ,
0.          , 0.52352941, 0.          , 0.          , 0.13529412])),
array([-inf, 1.3, 1.6, 1.9, 2.2, 2.5, 2.8, 3.1, 3.4, 3.7, inf])),
(array([0.48837209, 0.23837209, 0.18023256, 0.03488372, 0.01744186,
0.01744186, 0.01162791, 0.          , 0.00581395, 0.00581395])),
array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf])),
(array([0.31137725, 0.37125749, 0.05389222, 0.1497006 , 0.          ,
0.04191617, 0.04790419, 0.00598802, 0.01197605, 0.00598802])),
array([-inf, 1.8, 2.6, 3.4, 4.2, 5. , 5.8, 6.6, 7.4, 8.2, inf])),
(array([0.66470588, 0.11764706, 0.05882353, 0.01176471, 0.09411765,
0.          , 0.01764706, 0.01176471, 0.01176471, 0.01176471])),
array([-inf, 1.9, 2.8, 3.7, 4.6, 5.5, 6.4, 7.3, 8.2, 9.1, inf])),
(array([0.40462428, 0.28901734, 0.04624277, 0.01734104, 0.          ,
0.17919075, 0.04624277, 0.00578035, 0.          , 0.01156069])),
array([-inf, 1.8, 2.6, 3.4, 4.2, 5. , 5.8, 6.6, 7.4, 8.2, inf]]))

```

```
[26]: def check_fit_params():
```

```

    """
    This function is designed to test the fit_params method of a
    SimpleBayesClassifier.
    This method is presumably responsible for computing parameters for a Naive
    Bayes classifier
    based on the provided training data. The parameters in this context is bins
    and edges from each histogram.
    """

    T = SimpleBayesClassifier(2, 2)
    X_TRAIN_CASE_1 = np.array([
        [0, 1, 2, 3],
        [1, 2, 3, 4],
        [2, 3, 4, 5],
        [3, 4, 5, 6]
    ])

```

```

])
Y_TRAIN_CASE_1 = np.array([0, 1, 0, 1])
STAY_PARAMS_1, LEAVE_PARAMS_1 = T.fit_params(X_TRAIN_CASE_1, Y_TRAIN_CASE_1)

print("STAY PARAMETERS")
for f_idx in range(len(STAY_PARAMS_1)):
    print(f"Feature : {f_idx}")
    print(f"BINS : {STAY_PARAMS_1[f_idx][0]}")
    print(f"EDGES : {STAY_PARAMS_1[f_idx][1]}")
print("")
print("LEAVE PARAMETERS")
for f_idx in range(len(STAY_PARAMS_1)):
    print(f"Feature : {f_idx}")
    print(f"BINS : {LEAVE_PARAMS_1[f_idx][0]}")
    print(f"EDGES : {LEAVE_PARAMS_1[f_idx][1]}")

check_fit_params()

```

STAY PARAMETERS

```

Feature : 0
BINS : [0.5 0.  0.  0.  0.  0.  0.  0.  0.  0.5]
EDGES : [-inf  0.2  0.4  0.6  0.8  1.   1.2  1.4  1.6  1.8  inf]
Feature : 1
BINS : [0.5 0.  0.  0.  0.  0.  0.  0.  0.  0.5]
EDGES : [-inf  1.2  1.4  1.6  1.8  2.   2.2  2.4  2.6  2.8  inf]
Feature : 2
BINS : [0.5 0.  0.  0.  0.  0.  0.  0.  0.  0.5]
EDGES : [-inf  2.2  2.4  2.6  2.8  3.   3.2  3.4  3.6  3.8  inf]
Feature : 3
BINS : [0.5 0.  0.  0.  0.  0.  0.  0.  0.  0.5]
EDGES : [-inf  3.2  3.4  3.6  3.8  4.   4.2  4.4  4.6  4.8  inf]

```

LEAVE PARAMETERS

```

Feature : 0
BINS : [0.5 0.  0.  0.  0.  0.  0.  0.  0.  0.5]
EDGES : [-inf  1.2  1.4  1.6  1.8  2.   2.2  2.4  2.6  2.8  inf]
Feature : 1
BINS : [0.5 0.  0.  0.  0.  0.  0.  0.  0.  0.5]
EDGES : [-inf  2.2  2.4  2.6  2.8  3.   3.2  3.4  3.6  3.8  inf]
Feature : 2
BINS : [0.5 0.  0.  0.  0.  0.  0.  0.  0.  0.5]
EDGES : [-inf  3.2  3.4  3.6  3.8  4.   4.2  4.4  4.6  4.8  inf]
Feature : 3
BINS : [0.5 0.  0.  0.  0.  0.  0.  0.  0.  0.5]
EDGES : [-inf  4.2  4.4  4.6  4.8  5.   5.2  5.4  5.6  5.8  inf]

```

```
[27]: x_test_np = np.array(x_test)
      y_pred = np.array(model.predict(x = x_test_np))
```

```
[28]: y_pred.shape
```

```
[28]: (147,)
```

```
[29]: def evaluate(y_true, y_pred, show_result = True):
      if (y_true.shape[0] != y_pred.shape[0]):
          return -1, -1, -1, -1

      tp = np.sum((y_pred == 1) & (y_true == 1))
      fn = np.sum((y_pred == 0) & (y_true == 1))
      fp = np.sum((y_pred == 1) & (y_true == 0))
      tn = np.sum((y_pred == 0) & (y_true == 0))

      accuracy = (tp + tn) / (tp + tn + fp + fn)
      precision = (tp) / (tp + fp)
      recall = (tp) / (tp + fn)
      F1 = 2 / (1 / precision + 1 / recall)
      fpr = (fp) / (fp + tn)

      if (show_result):
          print(f"Accuracy: {accuracy * 100}%\nPrecision: {precision}\nRecall: {recall}\nF1: {F1}\nFPR: {fpr}\n")

      return accuracy, precision, recall, F1, fpr
```

```
[30]: evaluate(np.array(y_test), y_pred)
```

```
Accuracy: 80.95238095238095%
Precision: 0.375
Recall: 0.25
F1: 0.30000000000000004
FPR: 0.08130081300813008
```

```
[30]: (np.float64(0.8095238095238095),
      np.float64(0.375),
      np.float64(0.25),
      np.float64(0.30000000000000004),
      np.float64(0.08130081300813008))
```

1.0.14 T12. Use the learned distributions to classify the test set. Report the results using the same metric as the previous question.

```
[32]: df = pd.read_csv('hr-employee-attrition-with-null.csv')

df.loc[df["Attrition"] == "no", "Attrition"] = 0.0
df.loc[df["Attrition"] == "yes", "Attrition"] = 1.0
string_categorical_col = ['Department', 'Attrition', 'BusinessTravel',
    ↪ 'EducationField', 'Gender', 'JobRole',
    ↪ 'MaritalStatus', 'Over18', 'OverTime']

# ENCODE STRING COLUMNS TO CATEGORICAL COLUMNS
for col in string_categorical_col:
    # INSERT CODE HERE
    df[col] = pd.Categorical(df[col]).codes
# HANDLE NULL NUMBERS
# I don't think we need to handle null?

# INSERT CODE HERE
df = df.loc[:, ~df.columns.isin(['EmployeeNumber', 'Unnamed: 0',
    ↪ 'EmployeeCount', 'StandardHours', 'Over18'])] # drop these columns

X = df.drop(["Attrition"], axis=1)
Y = df["Attrition"]

x_train, x_test, y_train, y_test = train_test_split(X, Y, stratify=Y,
    ↪ test_size=0.1, random_state=12345)
```

```
[34]: model.fit_gaussian_params(np.array(x_train), np.array(y_train))
```

```
[34]: ((np.float64(37.809203142536475), np.float64(8.896433952739484)),
(np.float64(1.0891891891891892), np.float64(1.2118670987811953)),
(np.float64(808.728187919463), np.float64(407.2251752337319)),
(np.float64(0.809009009009009), np.float64(0.9952635170583247)),
(np.float64(9.0773606370876), np.float64(8.112178479155704)),
(np.float64(2.9395667046750287), np.float64(1.0218798132530393)),
(np.float64(1.6054054054054054), np.float64(1.7453071098573252)),
(np.float64(2.755826859045505), np.float64(1.0796687888712027)),
(np.float64(0.2594594594594595), np.float64(0.7706763588940992)),
(np.float64(65.53699551569507), np.float64(20.404589827789692)),
(np.float64(2.7610419026047563), np.float64(0.6885077018805345)),
(np.float64(2.179749715585893), np.float64(1.120498601502818)),
(np.float64(3.324324324324324), np.float64(3.076684076050504)),
(np.float64(2.7587768969422424), np.float64(1.0964885160093172)),
(np.float64(0.6549549549549549), np.float64(1.0357778761789396)),
(np.float64(6847.662556053811), np.float64(4798.092360725341)),
(np.float64(14266.125842696629), np.float64(7124.572193733051)),
```

```
(np.float64(2.697278911564626), np.float64(2.475772279106189)),
(np.float64(-0.005405405405405406), np.float64(0.6151009347828872)),
(np.float64(15.266968325791856), np.float64(3.6403649988412696)),
(np.float64(3.146892655367232), np.float64(0.35399887452701784)),
(np.float64(2.727474972191324), np.float64(1.0859766282649876)),
(np.float64(0.8098987626546682), np.float64(0.8152600328328238)),
(np.float64(11.946547884187082), np.float64(7.805074751024999)),
(np.float64(2.823793490460157), np.float64(1.314400760052736)),
(np.float64(2.781920903954802), np.float64(0.6830562214440479)),
(np.float64(7.47250280583614), np.float64(6.18731880041615)),
(np.float64(4.498327759197324), np.float64(3.630701412165684)),
(np.float64(2.280269058295964), np.float64(3.22057788303618)),
(np.float64(4.394110985277464), np.float64(3.552047784335249))],
[(np.float64(33.947976878612714), np.float64(10.265417501739252)),
(np.float64(1.1267605633802817), np.float64(1.1538412368896775)),
(np.float64(769.3248407643312), np.float64(399.78765409130403)),
(np.float64(0.8075117370892019), np.float64(1.1071851998485391)),
(np.float64(11.344632768361581), np.float64(8.707415308937323)),
(np.float64(2.8238636363636362), np.float64(1.0156160600369408)),
(np.float64(1.7089201877934272), np.float64(1.8101864241960348)),
(np.float64(2.55625), np.float64(1.171254002127634)),
(np.float64(0.3192488262910798), np.float64(0.7760834613779912)),
(np.float64(63.875739644970416), np.float64(20.622901977931885)),
(np.float64(2.5229885057471266), np.float64(0.755850917220641)),
(np.float64(1.6136363636363635), np.float64(0.9223465403501375)),
(np.float64(3.807511737089202), np.float64(3.2653451249559375)),
(np.float64(2.5714285714285716), np.float64(1.1080411102665897)),
(np.float64(0.8497652582159625), np.float64(1.1732393426854273)),
(np.float64(4590.413173652694), np.float64(3511.341093996613)),
(np.float64(14242.339393939394), np.float64(7085.770015252696)),
(np.float64(3.0282485875706215), np.float64(2.7193165029633763)),
(np.float64(0.16901408450704225), np.float64(0.7809813645794668)),
(np.float64(14.966480446927374), np.float64(3.8168978034257677)),
(np.float64(3.164705882352941), np.float64(0.3709148887161046)),
(np.float64(2.5304878048780486), np.float64(1.133867655243493)),
(np.float64(0.5290697674418605), np.float64(0.8519971447530724)),
(np.float64(8.587878787878788), np.float64(7.430476047221869)),
(np.float64(2.601156069364162), np.float64(1.2056700551978288)),
(np.float64(2.6882352941176473), np.float64(0.8348175939838075)),
(np.float64(5.325581395348837), np.float64(5.541475457256379)),
(np.float64(3.2095808383233533), np.float64(3.2991240632602064)),
(np.float64(1.9529411764705882), np.float64(3.084709471887354)),
(np.float64(2.9653179190751446), np.float64(3.2769927595779498))])
```

```
[35]: def check_fit_gaussian_params():
```

```
    """
```

This function is designed to test the fit_gaussian_params method of a SimpleBayesClassifier.

This method is presumably responsible for computing parameters for a Naive Bayes classifier based on the provided training data. The parameters in this context is mean and STD.

"""

```
T = SimpleBayesClassifier(2, 2)
X_TRAIN_CASE_1 = np.array([
    [0, 1, 2, 3],
    [1, 2, 3, 4],
    [2, 3, 4, 5],
    [3, 4, 5, 6]
])
Y_TRAIN_CASE_1 = np.array([0, 1, 0, 1])
STAY_PARAMS_1, LEAVE_PARAMS_1 = T.fit_gaussian_params(X_TRAIN_CASE_1,
Y_TRAIN_CASE_1)

print("STAY PARAMETERS")
for f_idx in range(len(STAY_PARAMS_1)):
    print(f"Feature : {f_idx}")
    print(f"Mean : {STAY_PARAMS_1[f_idx][0]}")
    print(f"STD. : {STAY_PARAMS_1[f_idx][1]}")
print("")
print("LEAVE PARAMETERS")
for f_idx in range(len(STAY_PARAMS_1)):
    print(f"Feature : {f_idx}")
    print(f"Mean : {LEAVE_PARAMS_1[f_idx][0]}")
    print(f"STD. : {LEAVE_PARAMS_1[f_idx][1]}")

check_fit_gaussian_params()
```

STAY PARAMETERS

```
Feature : 0
Mean : 1.0
STD. : 1.0
Feature : 1
Mean : 2.0
STD. : 1.0
Feature : 2
Mean : 3.0
STD. : 1.0
Feature : 3
Mean : 4.0
STD. : 1.0
```

```
LEAVE PARAMETERS
```

```
Feature : 0  
Mean : 2.0  
STD. : 1.0  
Feature : 1  
Mean : 3.0  
STD. : 1.0  
Feature : 2  
Mean : 4.0  
STD. : 1.0  
Feature : 3  
Mean : 5.0  
STD. : 1.0
```

```
[36]: y_pred = model.gaussian_predict(np.array(x_test))
```

```
[37]: evaluate(y_test, y_pred)
```

```
Accuracy: 81.63265306122449%  
Precision: 0.45161290322580644  
Recall: 0.5833333333333334  
F1: 0.509090909090909  
FPR: 0.13821138211382114
```

```
[37]: (np.float64(0.8163265306122449),  
      np.float64(0.45161290322580644),  
      np.float64(0.5833333333333334),  
      np.float64(0.509090909090909),  
      np.float64(0.13821138211382114))
```

1.0.15 T13 : The random choice baseline is the accuracy if you make a random guess for each test sample. Give random guess (50% leaving, and 50% staying) to the test samples. Report the overall Accuracy. Then, report the Precision, Recall, and F score for attrition prediction using the random choice baseline.

```
[38]: y_random_pred = np.random.default_rng(seed=12345).random(y_test.shape)
```

```
[39]: y_random_pred
```

```
[39]: array([0.22733602, 0.31675834, 0.79736546, 0.67625467, 0.39110955,  
          0.33281393, 0.59830875, 0.18673419, 0.67275604, 0.94180287,  
          0.24824571, 0.94888115, 0.66723745, 0.09589794, 0.44183967,  
          0.88647992, 0.6974535 , 0.32647286, 0.73392816, 0.22013496,  
          0.08159457, 0.1598956 , 0.34010018, 0.46519315, 0.26642103,  
          0.8157764 , 0.19329439, 0.12946908, 0.09166475, 0.59856801,  
          0.8547419 , 0.60162124, 0.93198836, 0.72478136, 0.86055132,  
          0.9293378 , 0.54618601, 0.93767296, 0.49498794, 0.27377318,
```



```

0.45177871, 0.66503892, 0.33089093, 0.90345401, 0.25707418,
0.33982834, 0.2588534 , 0.35544648, 0.00502233, 0.62860454,
0.28238271, 0.06808769, 0.61682898, 0.17632632, 0.30438839,
0.44088681, 0.15020234, 0.21792886, 0.47433312, 0.47636886,
0.25523235, 0.29756527, 0.27906712, 0.26057921, 0.48276159,
0.21197904, 0.4956306 , 0.24626133, 0.83848265, 0.18013059,
0.86215629, 0.17829944, 0.75053133, 0.6111204 , 0.20915503,
0.75987242, 0.24926057, 0.08557173, 0.61805672, 0.53696833,
0.63452671, 0.17437411, 0.24816449, 0.68482298, 0.08087165,
0.8750736 , 0.42869438, 0.6183942 , 0.3131055 , 0.17896286,
0.00971213, 0.21004296, 0.87000068, 0.9728298 , 0.44179234,
0.37874949, 0.27594708, 0.96610411, 0.05820261, 0.4087339 ,
0.16862884, 0.24014406, 0.78000786, 0.2037676 , 0.55205095,
0.36699414, 0.50728172, 0.3334378 , 0.28272167, 0.2818303 ,
0.08538129, 0.48181366, 0.88334289, 0.94722777, 0.02738372,
0.91775224, 0.12152453, 0.74784776, 0.89652074, 0.1679298 ,
0.33146322, 0.37815663, 0.34684896, 0.5162557 , 0.00899403,
0.4226782 , 0.87765773, 0.08740515, 0.48408482, 0.48122773,
0.78257149, 0.96455958, 0.70709644, 0.27373672, 0.6701133 ,
0.3475348 , 0.76812784, 0.67577142, 0.97753203, 0.86670979,
0.04610801, 0.29032371, 0.8623911 , 0.60084783, 0.34425796,
0.05560258, 0.76287267])

```

```

[40]: y_random_pred[y_random_pred >= 0.5] = 1
      y_random_pred[y_random_pred < 0.5] = 0

```

```

[41]: y_random_pred

```

```

[41]: array([0., 0., 1., 1., 0., 0., 1., 0., 1., 1., 0., 1., 1., 0., 0., 1., 1.,
          0., 1., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 1., 1., 1., 1., 1.,
          1., 1., 1., 1., 0., 0., 0., 1., 0., 1., 0., 0., 0., 0., 0., 1., 0.,
          0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.,
          1., 0., 1., 0., 1., 1., 0., 1., 0., 0., 1., 1., 1., 0., 0., 1., 0.,
          1., 0., 1., 0., 0., 0., 0., 1., 1., 0., 0., 0., 1., 0., 0., 0., 0.,
          1., 0., 1., 0., 1., 0., 0., 0., 0., 0., 1., 1., 0., 1., 0., 1., 1.,
          0., 0., 0., 0., 1., 0., 0., 1., 0., 0., 0., 1., 1., 1., 0., 1., 0.,
          1., 1., 1., 1., 0., 0., 1., 1., 0., 0., 1.])

```

```

[42]: evaluate(y_test, y_random_pred)

```

```

Accuracy: 54.421768707483%
Precision: 0.13559322033898305
Recall: 0.3333333333333333
F1: 0.1927710843373494
FPR: 0.4146341463414634

```

```
[42]: (np.float64(0.54421768707483),
      np.float64(0.13559322033898305),
      np.float64(0.3333333333333333),
      np.float64(0.1927710843373494),
      np.float64(0.4146341463414634))
```

1.0.16 T14. The majority rule is the accuracy if you use the most frequent class from the training set as the classification decision. Report the overall Accuracy. Then, report the Precision, Recall, and F score for attrition prediction using the majority rule baseline.

```
[43]: print("Leave:", np.sum(y_train == 1))
      print("Stay:", np.sum(y_train == 0))
      print("Stay (0) is the majority class")
      y_major_pred = np.zeros(y_pred.shape)
```

```

Leave: 213
Stay: 1110
Stay (0) is the majority class

```

```
[44]: y_major_pred
```

[illegible]

```
[45]: evaluate(y_test, y_major_pred)
```

```
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0
```

```
C:\Users\chyut\AppData\Local\Temp\ipykernel_24836\297337688.py:11:
RuntimeWarning: invalid value encountered in scalar divide
    precision = (tp) / (tp + fp)
C:\Users\chyut\AppData\Local\Temp\ipykernel_24836\297337688.py:13:
RuntimeWarning: divide by zero encountered in scalar divide
    F1 = 2 / (1 / precision + 1 / recall)
```

```
[45]: (np.float64(0.8367346938775511),
      np.float64(nan),
      np.float64(0.0),
      np.float64(nan),
      np.float64(0.0))
```

1.0.17 T15. Compare the two baselines with your Naive Bayes classifier.

Mine's accuracy is more than the random baseline, but a little bit less than the stupid baseline. But, the stupid baseline precision is 0 which mine's is better.

1.0.18 T16. Use the following threshold values

\$ t = np.arange(-5,5,0.05) \$ ### find the best accuracy, and F score (and the corresponding thresholds)

```
[46]: t = np.arange(-5, 5, 0.05)
maxaccuracy = -1
maxf1 = -1
maxthresholdacc = -10
maxthresholdf1 = -10

history = {}
history["Accuracy"] = []
history["Precision"] = []
history["Recall"] = []
history["F1"] = []
history["FPR"] = []

for each in t:
    print("--- Threshold =", each, "---")
    y_pred = model.gaussian_predict(np.array(x_test), thresh=each)

    accuracy, precision, recall, f1, fpr = evaluate(y_test, y_pred)

    history["Accuracy"].append(accuracy)
    history["Precision"].append(precision)
    history["Recall"].append(recall)
    history["F1"].append(f1)
    history["FPR"].append(fpr)

    if (accuracy > maxaccuracy):
        maxaccuracy = accuracy
        maxthresholdacc = each
    if (f1 > maxf1):
        maxf1 = f1
        maxthresholdf1 = each
```

```
print("Best Accuracy:", maxaccuracy, "with responding threshold:",  
      ↪maxthresholdacc)  
print("Best F1:", maxf1, "with responding threshold:", maxthresholdf1)
```

--- Threshold = -5.0 ---

Accuracy: 27.2108843537415%
Precision: 0.1732283464566929
Recall: 0.9166666666666666
F1: 0.2913907284768212
FPR: 0.8536585365853658

--- Threshold = -4.95 ---

Accuracy: 27.2108843537415%
Precision: 0.1732283464566929
Recall: 0.9166666666666666
F1: 0.2913907284768212
FPR: 0.8536585365853658

--- Threshold = -4.9 ---

Accuracy: 27.2108843537415%
Precision: 0.1732283464566929
Recall: 0.9166666666666666
F1: 0.2913907284768212
FPR: 0.8536585365853658

--- Threshold = -4.8500000000000005 ---

Accuracy: 27.2108843537415%
Precision: 0.1732283464566929
Recall: 0.9166666666666666
F1: 0.2913907284768212
FPR: 0.8536585365853658

--- Threshold = -4.8000000000000001 ---

Accuracy: 27.2108843537415%
Precision: 0.1732283464566929
Recall: 0.9166666666666666
F1: 0.2913907284768212
FPR: 0.8536585365853658

--- Threshold = -4.7500000000000001 ---

Accuracy: 27.2108843537415%
Precision: 0.1732283464566929
Recall: 0.9166666666666666
F1: 0.2913907284768212
FPR: 0.8536585365853658

--- Threshold = -4.7000000000000001 ---

Accuracy: 27.2108843537415%

Precision: 0.1732283464566929
Recall: 0.9166666666666666
F1: 0.2913907284768212
FPR: 0.8536585365853658

--- Threshold = -4.6500000000000001 ---
Accuracy: 27.2108843537415%
Precision: 0.1732283464566929
Recall: 0.9166666666666666
F1: 0.2913907284768212
FPR: 0.8536585365853658

--- Threshold = -4.6000000000000001 ---
Accuracy: 28.57142857142857%
Precision: 0.176
Recall: 0.9166666666666666
F1: 0.29530201342281875
FPR: 0.8373983739837398

--- Threshold = -4.5500000000000002 ---
Accuracy: 28.57142857142857%
Precision: 0.176
Recall: 0.9166666666666666
F1: 0.29530201342281875
FPR: 0.8373983739837398

--- Threshold = -4.5000000000000002 ---
Accuracy: 29.25170068027211%
Precision: 0.1774193548387097
Recall: 0.9166666666666666
F1: 0.2972972972972973
FPR: 0.8292682926829268

--- Threshold = -4.4500000000000002 ---
Accuracy: 29.931972789115648%
Precision: 0.17886178861788618
Recall: 0.9166666666666666
F1: 0.29931972789115646
FPR: 0.8211382113821138

--- Threshold = -4.4000000000000002 ---
Accuracy: 29.931972789115648%
Precision: 0.17886178861788618
Recall: 0.9166666666666666
F1: 0.29931972789115646
FPR: 0.8211382113821138

--- Threshold = -4.3500000000000002 ---

Accuracy: 29.931972789115648%
 Precision: 0.17886178861788618
 Recall: 0.9166666666666666
 F1: 0.29931972789115646
 FPR: 0.8211382113821138

--- Threshold = -4.3000000000000025 ---
 Accuracy: 29.931972789115648%
 Precision: 0.17886178861788618
 Recall: 0.9166666666666666
 F1: 0.29931972789115646
 FPR: 0.8211382113821138

--- Threshold = -4.250000000000003 ---
 Accuracy: 30.612244897959183%
 Precision: 0.18032786885245902
 Recall: 0.9166666666666666
 F1: 0.30136986301369867
 FPR: 0.8130081300813008

--- Threshold = -4.200000000000003 ---
 Accuracy: 30.612244897959183%
 Precision: 0.18032786885245902
 Recall: 0.9166666666666666
 F1: 0.30136986301369867
 FPR: 0.8130081300813008

--- Threshold = -4.150000000000003 ---
 Accuracy: 30.612244897959183%
 Precision: 0.18032786885245902
 Recall: 0.9166666666666666
 F1: 0.30136986301369867
 FPR: 0.8130081300813008

--- Threshold = -4.100000000000003 ---
 Accuracy: 30.612244897959183%
 Precision: 0.18032786885245902
 Recall: 0.9166666666666666
 F1: 0.30136986301369867
 FPR: 0.8130081300813008

--- Threshold = -4.050000000000003 ---
 Accuracy: 30.612244897959183%
 Precision: 0.18032786885245902
 Recall: 0.9166666666666666
 F1: 0.30136986301369867
 FPR: 0.8130081300813008

--- Threshold = -4.0000000000000036 ---
Accuracy: 31.292517006802722%
Precision: 0.18181818181818182
Recall: 0.9166666666666666
F1: 0.30344827586206896
FPR: 0.8048780487804879

--- Threshold = -3.9500000000000037 ---
Accuracy: 31.292517006802722%
Precision: 0.18181818181818182
Recall: 0.9166666666666666
F1: 0.30344827586206896
FPR: 0.8048780487804879

--- Threshold = -3.9000000000000004 ---
Accuracy: 31.292517006802722%
Precision: 0.18181818181818182
Recall: 0.9166666666666666
F1: 0.30344827586206896
FPR: 0.8048780487804879

--- Threshold = -3.8500000000000004 ---
Accuracy: 31.292517006802722%
Precision: 0.18181818181818182
Recall: 0.9166666666666666
F1: 0.30344827586206896
FPR: 0.8048780487804879

--- Threshold = -3.8000000000000043 ---
Accuracy: 31.97278911564626%
Precision: 0.1833333333333332
Recall: 0.9166666666666666
F1: 0.3055555555555555
FPR: 0.7967479674796748

--- Threshold = -3.7500000000000044 ---
Accuracy: 31.97278911564626%
Precision: 0.1833333333333332
Recall: 0.9166666666666666
F1: 0.3055555555555555
FPR: 0.7967479674796748

--- Threshold = -3.7000000000000046 ---
Accuracy: 32.6530612244898%
Precision: 0.18487394957983194
Recall: 0.9166666666666666
F1: 0.3076923076923077
FPR: 0.7886178861788617

--- Threshold = -3.6500000000000005 ---

Accuracy: 32.6530612244898%

Precision: 0.18487394957983194

Recall: 0.9166666666666666

F1: 0.3076923076923077

FPR: 0.7886178861788617

--- Threshold = -3.6000000000000005 ---

Accuracy: 34.01360544217687%

Precision: 0.18803418803418803

Recall: 0.9166666666666666

F1: 0.3120567375886525

FPR: 0.7723577235772358

--- Threshold = -3.5500000000000005 ---

Accuracy: 34.01360544217687%

Precision: 0.18803418803418803

Recall: 0.9166666666666666

F1: 0.3120567375886525

FPR: 0.7723577235772358

--- Threshold = -3.50000000000000053 ---

Accuracy: 34.69387755102041%

Precision: 0.1896551724137931

Recall: 0.9166666666666666

F1: 0.3142857142857143

FPR: 0.7642276422764228

--- Threshold = -3.45000000000000055 ---

Accuracy: 34.01360544217687%

Precision: 0.1826086956521739

Recall: 0.875

F1: 0.302158273381295

FPR: 0.7642276422764228

--- Threshold = -3.40000000000000057 ---

Accuracy: 34.69387755102041%

Precision: 0.18421052631578946

Recall: 0.875

F1: 0.30434782608695654

FPR: 0.7560975609756098

--- Threshold = -3.3500000000000006 ---

Accuracy: 34.69387755102041%

Precision: 0.18421052631578946

Recall: 0.875

F1: 0.30434782608695654

FPR: 0.7560975609756098

--- Threshold = -3.3000000000000006 ---

Accuracy: 34.69387755102041%

Precision: 0.18421052631578946

Recall: 0.875

F1: 0.30434782608695654

FPR: 0.7560975609756098

--- Threshold = -3.2500000000000006 ---

Accuracy: 34.69387755102041%

Precision: 0.18421052631578946

Recall: 0.875

F1: 0.30434782608695654

FPR: 0.7560975609756098

--- Threshold = -3.20000000000000064 ---

Accuracy: 34.69387755102041%

Precision: 0.18421052631578946

Recall: 0.875

F1: 0.30434782608695654

FPR: 0.7560975609756098

--- Threshold = -3.15000000000000066 ---

Accuracy: 35.374149659863946%

Precision: 0.18584070796460178

Recall: 0.875

F1: 0.30656934306569344

FPR: 0.7479674796747967

--- Threshold = -3.10000000000000068 ---

Accuracy: 36.054421768707485%

Precision: 0.1875

Recall: 0.875

F1: 0.3088235294117647

FPR: 0.7398373983739838

--- Threshold = -3.0500000000000007 ---

Accuracy: 36.054421768707485%

Precision: 0.1875

Recall: 0.875

F1: 0.3088235294117647

FPR: 0.7398373983739838

--- Threshold = -3.0000000000000007 ---

Accuracy: 37.41496598639456%

Precision: 0.19090909090909092

Recall: 0.875

F1: 0.3134328358208955
FPR: 0.7235772357723578

--- Threshold = -2.9500000000000073 ---
Accuracy: 37.41496598639456%
Precision: 0.19090909090909092
Recall: 0.875
F1: 0.3134328358208955
FPR: 0.7235772357723578

--- Threshold = -2.9000000000000075 ---
Accuracy: 37.41496598639456%
Precision: 0.19090909090909092
Recall: 0.875
F1: 0.3134328358208955
FPR: 0.7235772357723578

--- Threshold = -2.8500000000000076 ---
Accuracy: 37.41496598639456%
Precision: 0.19090909090909092
Recall: 0.875
F1: 0.3134328358208955
FPR: 0.7235772357723578

--- Threshold = -2.8000000000000008 ---
Accuracy: 38.095238095238095%
Precision: 0.1926605504587156
Recall: 0.875
F1: 0.31578947368421056
FPR: 0.7154471544715447

--- Threshold = -2.7500000000000008 ---
Accuracy: 40.136054421768705%
Precision: 0.19811320754716982
Recall: 0.875
F1: 0.3230769230769231
FPR: 0.6910569105691057

--- Threshold = -2.7000000000000008 ---
Accuracy: 40.816326530612244%
Precision: 0.2
Recall: 0.875
F1: 0.32558139534883723
FPR: 0.6829268292682927

--- Threshold = -2.6500000000000083 ---
Accuracy: 41.49659863945578%
Precision: 0.20192307692307693

Recall: 0.875
F1: 0.328125
FPR: 0.6747967479674797

--- Threshold = -2.6000000000000085 ---
Accuracy: 41.49659863945578%
Precision: 0.20192307692307693
Recall: 0.875
F1: 0.328125
FPR: 0.6747967479674797

--- Threshold = -2.5500000000000087 ---
Accuracy: 41.49659863945578%
Precision: 0.20192307692307693
Recall: 0.875
F1: 0.328125
FPR: 0.6747967479674797

--- Threshold = -2.5000000000000009 ---
Accuracy: 42.857142857142854%
Precision: 0.20588235294117646
Recall: 0.875
F1: 0.3333333333333333
FPR: 0.6585365853658537

--- Threshold = -2.4500000000000009 ---
Accuracy: 44.89795918367347%
Precision: 0.21212121212121213
Recall: 0.875
F1: 0.3414634146341463
FPR: 0.6341463414634146

--- Threshold = -2.40000000000000092 ---
Accuracy: 44.89795918367347%
Precision: 0.21212121212121213
Recall: 0.875
F1: 0.3414634146341463
FPR: 0.6341463414634146

--- Threshold = -2.35000000000000094 ---
Accuracy: 45.57823129251701%
Precision: 0.21428571428571427
Recall: 0.875
F1: 0.34426229508196715
FPR: 0.6260162601626016

--- Threshold = -2.30000000000000096 ---
Accuracy: 47.61904761904761%

Precision: 0.22105263157894736
Recall: 0.875
F1: 0.35294117647058826
FPR: 0.6016260162601627

--- Threshold = -2.2500000000000098 ---
Accuracy: 48.29931972789115%
Precision: 0.22340425531914893
Recall: 0.875
F1: 0.35593220338983056
FPR: 0.5934959349593496

--- Threshold = -2.200000000000001 ---
Accuracy: 48.97959183673469%
Precision: 0.22580645161290322
Recall: 0.875
F1: 0.358974358974359
FPR: 0.5853658536585366

--- Threshold = -2.150000000000001 ---
Accuracy: 49.65986394557823%
Precision: 0.22826086956521738
Recall: 0.875
F1: 0.3620689655172414
FPR: 0.5772357723577236

--- Threshold = -2.1000000000000103 ---
Accuracy: 51.02040816326531%
Precision: 0.23333333333333334
Recall: 0.875
F1: 0.3684210526315789
FPR: 0.5609756097560976

--- Threshold = -2.0500000000000105 ---
Accuracy: 53.06122448979592%
Precision: 0.2413793103448276
Recall: 0.875
F1: 0.37837837837837845
FPR: 0.5365853658536586

--- Threshold = -2.0000000000000107 ---
Accuracy: 53.74149659863946%
Precision: 0.2441860465116279
Recall: 0.875
F1: 0.3818181818181817
FPR: 0.5284552845528455

--- Threshold = -1.9500000000000108 ---

Accuracy: 55.78231292517006%
Precision: 0.25301204819277107
Recall: 0.875
F1: 0.39252336448598135
FPR: 0.5040650406504065

--- Threshold = -1.9000000000000011 ---
Accuracy: 55.78231292517006%
Precision: 0.25301204819277107
Recall: 0.875
F1: 0.39252336448598135
FPR: 0.5040650406504065

--- Threshold = -1.8500000000000012 ---
Accuracy: 55.78231292517006%
Precision: 0.25301204819277107
Recall: 0.875
F1: 0.39252336448598135
FPR: 0.5040650406504065

--- Threshold = -1.8000000000000014 ---
Accuracy: 56.4625850340136%
Precision: 0.25609756097560976
Recall: 0.875
F1: 0.39622641509433965
FPR: 0.4959349593495935

--- Threshold = -1.7500000000000015 ---
Accuracy: 57.14285714285714%
Precision: 0.25925925925925924
Recall: 0.875
F1: 0.4
FPR: 0.4878048780487805

--- Threshold = -1.7000000000000017 ---
Accuracy: 57.82312925170068%
Precision: 0.2625
Recall: 0.875
F1: 0.40384615384615385
FPR: 0.4796747967479675

--- Threshold = -1.6500000000000012 ---
Accuracy: 57.14285714285714%
Precision: 0.25316455696202533
Recall: 0.8333333333333334
F1: 0.38834951456310685
FPR: 0.4796747967479675

```

--- Threshold = -1.6000000000000012 ---
Accuracy: 58.50340136054422%
Precision: 0.2597402597402597
Recall: 0.8333333333333334
F1: 0.396039603960396
FPR: 0.4634146341463415

--- Threshold = -1.55000000000000123 ---
Accuracy: 58.50340136054422%
Precision: 0.2597402597402597
Recall: 0.8333333333333334
F1: 0.396039603960396
FPR: 0.4634146341463415

--- Threshold = -1.50000000000000124 ---
Accuracy: 59.863945578231295%
Precision: 0.2666666666666666
Recall: 0.8333333333333334
F1: 0.40404040404040403
FPR: 0.44715447154471544

--- Threshold = -1.45000000000000126 ---
Accuracy: 59.863945578231295%
Precision: 0.2666666666666666
Recall: 0.8333333333333334
F1: 0.40404040404040403
FPR: 0.44715447154471544

--- Threshold = -1.40000000000000128 ---
Accuracy: 61.224489795918366%
Precision: 0.273972602739726
Recall: 0.8333333333333334
F1: 0.41237113402061853
FPR: 0.43089430894308944

--- Threshold = -1.3500000000000013 ---
Accuracy: 61.904761904761905%
Precision: 0.2777777777777778
Recall: 0.8333333333333334
F1: 0.4166666666666667
FPR: 0.42276422764227645

--- Threshold = -1.30000000000000131 ---
Accuracy: 61.904761904761905%
Precision: 0.2777777777777778
Recall: 0.8333333333333334
F1: 0.4166666666666667
FPR: 0.42276422764227645

```

--- Threshold = -1.2500000000000133 ---

Accuracy: 63.94557823129252%

Precision: 0.2898550724637681

Recall: 0.8333333333333334

F1: 0.4301075268817205

FPR: 0.3983739837398374

--- Threshold = -1.2000000000000135 ---

Accuracy: 65.3061224489796%

Precision: 0.29850746268656714

Recall: 0.8333333333333334

F1: 0.43956043956043955

FPR: 0.3821138211382114

--- Threshold = -1.1500000000000137 ---

Accuracy: 66.6666666666666%

Precision: 0.3076923076923077

Recall: 0.8333333333333334

F1: 0.449438202247191

FPR: 0.36585365853658536

--- Threshold = -1.1000000000000139 ---

Accuracy: 68.02721088435374%

Precision: 0.31746031746031744

Recall: 0.8333333333333334

F1: 0.4597701149425287

FPR: 0.34959349593495936

--- Threshold = -1.050000000000014 ---

Accuracy: 68.70748299319727%

Precision: 0.3225806451612903

Recall: 0.8333333333333334

F1: 0.46511627906976744

FPR: 0.34146341463414637

--- Threshold = -1.0000000000000142 ---

Accuracy: 68.70748299319727%

Precision: 0.3225806451612903

Recall: 0.8333333333333334

F1: 0.46511627906976744

FPR: 0.34146341463414637

--- Threshold = -0.9500000000000144 ---

Accuracy: 69.38775510204081%

Precision: 0.32786885245901637

Recall: 0.8333333333333334

F1: 0.47058823529411764

FPR: 0.3333333333333333

--- Threshold = -0.900000000000000146 ---

Accuracy: 70.74829931972789%

Precision: 0.3389830508474576

Recall: 0.8333333333333334

F1: 0.48192771084337344

FPR: 0.3170731707317073

--- Threshold = -0.850000000000000147 ---

Accuracy: 72.10884353741497%

Precision: 0.3508771929824561

Recall: 0.8333333333333334

F1: 0.4938271604938272

FPR: 0.3008130081300813

--- Threshold = -0.800000000000000149 ---

Accuracy: 74.82993197278913%

Precision: 0.37735849056603776

Recall: 0.8333333333333334

F1: 0.5194805194805195

FPR: 0.2682926829268293

--- Threshold = -0.750000000000000151 ---

Accuracy: 76.19047619047619%

Precision: 0.39215686274509803

Recall: 0.8333333333333334

F1: 0.5333333333333333

FPR: 0.25203252032520324

--- Threshold = -0.700000000000000153 ---

Accuracy: 76.19047619047619%

Precision: 0.39215686274509803

Recall: 0.8333333333333334

F1: 0.5333333333333333

FPR: 0.25203252032520324

--- Threshold = -0.650000000000000155 ---

Accuracy: 77.55102040816327%

Precision: 0.40425531914893614

Recall: 0.7916666666666666

F1: 0.5352112676056338

FPR: 0.22764227642276422

--- Threshold = -0.600000000000000156 ---

Accuracy: 78.2312925170068%

Precision: 0.41304347826086957

Recall: 0.7916666666666666

F1: 0.5428571428571428
FPR: 0.21951219512195122

--- Threshold = -0.55000000000000158 ---
Accuracy: 79.59183673469387%
Precision: 0.4318181818181818
Recall: 0.7916666666666666
F1: 0.5588235294117646
FPR: 0.2032520325203252

--- Threshold = -0.5000000000000016 ---
Accuracy: 79.59183673469387%
Precision: 0.42857142857142855
Recall: 0.75
F1: 0.5454545454545454
FPR: 0.1951219512195122

--- Threshold = -0.450000000000001616 ---
Accuracy: 80.27210884353741%
Precision: 0.43902439024390244
Recall: 0.75
F1: 0.5538461538461539
FPR: 0.18699186991869918

--- Threshold = -0.400000000000001634 ---
Accuracy: 80.27210884353741%
Precision: 0.43902439024390244
Recall: 0.75
F1: 0.5538461538461539
FPR: 0.18699186991869918

--- Threshold = -0.35000000000000165 ---
Accuracy: 80.95238095238095%
Precision: 0.45
Recall: 0.75
F1: 0.5625
FPR: 0.17886178861788618

--- Threshold = -0.30000000000000167 ---
Accuracy: 80.27210884353741%
Precision: 0.4358974358974359
Recall: 0.7083333333333334
F1: 0.5396825396825398
FPR: 0.17886178861788618

--- Threshold = -0.25000000000000169 ---
Accuracy: 79.59183673469387%
Precision: 0.42105263157894735

Recall: 0.6666666666666666
F1: 0.5161290322580645
FPR: 0.17886178861788618

--- Threshold = -0.20000000000001705 ---
Accuracy: 80.95238095238095%
Precision: 0.4444444444444444
Recall: 0.6666666666666666
F1: 0.5333333333333333
FPR: 0.16260162601626016

--- Threshold = -0.15000000000001723 ---
Accuracy: 81.63265306122449%
Precision: 0.45714285714285713
Recall: 0.6666666666666666
F1: 0.5423728813559322
FPR: 0.15447154471544716

--- Threshold = -0.10000000000001741 ---
Accuracy: 82.31292517006803%
Precision: 0.46875
Recall: 0.625
F1: 0.5357142857142857
FPR: 0.13821138211382114

--- Threshold = -0.050000000000017586 ---
Accuracy: 82.31292517006803%
Precision: 0.46875
Recall: 0.625
F1: 0.5357142857142857
FPR: 0.13821138211382114

--- Threshold = -1.7763568394002505e-14 ---
Accuracy: 81.63265306122449%
Precision: 0.45161290322580644
Recall: 0.5833333333333334
F1: 0.509090909090909
FPR: 0.13821138211382114

--- Threshold = 0.04999999999998206 ---
Accuracy: 81.63265306122449%
Precision: 0.45161290322580644
Recall: 0.5833333333333334
F1: 0.509090909090909
FPR: 0.13821138211382114

--- Threshold = 0.09999999999998188 ---
Accuracy: 81.63265306122449%

Precision: 0.4482758620689655
Recall: 0.5416666666666666
F1: 0.4905660377358491
FPR: 0.13008130081300814

--- Threshold = 0.1499999999999817 ---
Accuracy: 82.31292517006803%
Precision: 0.4642857142857143
Recall: 0.5416666666666666
F1: 0.5
FPR: 0.12195121951219512

--- Threshold = 0.19999999999998153 ---
Accuracy: 82.99319727891157%
Precision: 0.48148148148148145
Recall: 0.5416666666666666
F1: 0.5098039215686274
FPR: 0.11382113821138211

--- Threshold = 0.24999999999998135 ---
Accuracy: 82.31292517006803%
Precision: 0.46153846153846156
Recall: 0.5
F1: 0.48000000000000001
FPR: 0.11382113821138211

--- Threshold = 0.29999999999998117 ---
Accuracy: 82.31292517006803%
Precision: 0.46153846153846156
Recall: 0.5
F1: 0.48000000000000001
FPR: 0.11382113821138211

--- Threshold = 0.349999999999981 ---
Accuracy: 82.31292517006803%
Precision: 0.45454545454545453
Recall: 0.4166666666666667
F1: 0.4347826086956522
FPR: 0.0975609756097561

--- Threshold = 0.3999999999999808 ---
Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.4166666666666667
F1: 0.45454545454545453
FPR: 0.08130081300813008

--- Threshold = 0.44999999999998064 ---

Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.4166666666666667
F1: 0.45454545454545453
FPR: 0.08130081300813008

--- Threshold = 0.49999999999998046 ---
Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.375
F1: 0.4285714285714286
FPR: 0.07317073170731707

--- Threshold = 0.5499999999999803 ---
Accuracy: 82.31292517006803%
Precision: 0.4375
Recall: 0.2916666666666667
F1: 0.35000000000000003
FPR: 0.07317073170731707

--- Threshold = 0.5999999999999801 ---
Accuracy: 82.31292517006803%
Precision: 0.4375
Recall: 0.2916666666666667
F1: 0.35000000000000003
FPR: 0.07317073170731707

--- Threshold = 0.6499999999999799 ---
Accuracy: 82.99319727891157%
Precision: 0.4666666666666667
Recall: 0.2916666666666667
F1: 0.358974358974359
FPR: 0.06504065040650407

--- Threshold = 0.6999999999999797 ---
Accuracy: 82.99319727891157%
Precision: 0.46153846153846156
Recall: 0.25
F1: 0.32432432432432434
FPR: 0.056910569105691054

--- Threshold = 0.7499999999999796 ---
Accuracy: 82.99319727891157%
Precision: 0.46153846153846156
Recall: 0.25
F1: 0.32432432432432434
FPR: 0.056910569105691054

--- Threshold = 0.799999999999794 ---
Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.25
F1: 0.3333333333333333
FPR: 0.04878048780487805

--- Threshold = 0.849999999999792 ---
Accuracy: 85.03401360544217%
Precision: 0.6
Recall: 0.25
F1: 0.3529411764705882
FPR: 0.032520325203252036

--- Threshold = 0.89999999999979 ---
Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.16666666666666666
F1: 0.25
FPR: 0.032520325203252036

--- Threshold = 0.949999999999789 ---
Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.16666666666666666
F1: 0.25
FPR: 0.032520325203252036

--- Threshold = 0.999999999999787 ---
Accuracy: 84.35374149659864%
Precision: 0.5714285714285714
Recall: 0.16666666666666666
F1: 0.25806451612903225
FPR: 0.024390243902439025

--- Threshold = 1.049999999999785 ---
Accuracy: 85.03401360544217%
Precision: 0.6666666666666666
Recall: 0.16666666666666666
F1: 0.26666666666666666
FPR: 0.016260162601626018

--- Threshold = 1.099999999999783 ---
Accuracy: 84.35374149659864%
Precision: 0.6
Recall: 0.125
F1: 0.20689655172413796
FPR: 0.016260162601626018

--- Threshold = 1.1499999999999782 ---

Accuracy: 84.35374149659864%

Precision: 0.6

Recall: 0.125

F1: 0.20689655172413796

FPR: 0.016260162601626018

--- Threshold = 1.199999999999978 ---

Accuracy: 84.35374149659864%

Precision: 0.6

Recall: 0.125

F1: 0.20689655172413796

FPR: 0.016260162601626018

--- Threshold = 1.2499999999999778 ---

Accuracy: 84.35374149659864%

Precision: 0.6

Recall: 0.125

F1: 0.20689655172413796

FPR: 0.016260162601626018

--- Threshold = 1.2999999999999776 ---

Accuracy: 84.35374149659864%

Precision: 0.6

Recall: 0.125

F1: 0.20689655172413796

FPR: 0.016260162601626018

--- Threshold = 1.3499999999999774 ---

Accuracy: 83.6734693877551%

Precision: 0.5

Recall: 0.08333333333333333

F1: 0.14285714285714285

FPR: 0.016260162601626018

--- Threshold = 1.3999999999999773 ---

Accuracy: 83.6734693877551%

Precision: 0.5

Recall: 0.08333333333333333

F1: 0.14285714285714285

FPR: 0.016260162601626018

--- Threshold = 1.449999999999977 ---

Accuracy: 83.6734693877551%

Precision: 0.5

Recall: 0.08333333333333333

F1: 0.14285714285714285

FPR: 0.016260162601626018

--- Threshold = 1.499999999999977 ---

Accuracy: 83.6734693877551%

Precision: 0.5

Recall: 0.08333333333333333

F1: 0.14285714285714285

FPR: 0.016260162601626018

--- Threshold = 1.5499999999999767 ---

Accuracy: 83.6734693877551%

Precision: 0.5

Recall: 0.08333333333333333

F1: 0.14285714285714285

FPR: 0.016260162601626018

--- Threshold = 1.5999999999999766 ---

Accuracy: 83.6734693877551%

Precision: 0.5

Recall: 0.08333333333333333

F1: 0.14285714285714285

FPR: 0.016260162601626018

--- Threshold = 1.6499999999999764 ---

Accuracy: 82.99319727891157%

Precision: 0.3333333333333333

Recall: 0.04166666666666664

F1: 0.07407407407407407

FPR: 0.016260162601626018

--- Threshold = 1.6999999999999762 ---

Accuracy: 82.99319727891157%

Precision: 0.3333333333333333

Recall: 0.04166666666666664

F1: 0.07407407407407407

FPR: 0.016260162601626018

--- Threshold = 1.749999999999976 ---

Accuracy: 82.99319727891157%

Precision: 0.3333333333333333

Recall: 0.04166666666666664

F1: 0.07407407407407407

FPR: 0.016260162601626018

--- Threshold = 1.7999999999999758 ---

Accuracy: 82.99319727891157%

Precision: 0.3333333333333333

Recall: 0.04166666666666664

F1: 0.07407407407407407
FPR: 0.016260162601626018

--- Threshold = 1.8499999999999757 ---
Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.04166666666666664
F1: 0.07692307692307693
FPR: 0.008130081300813009

--- Threshold = 1.8999999999999755 ---
Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.04166666666666664
F1: 0.07692307692307693
FPR: 0.008130081300813009

--- Threshold = 1.9499999999999753 ---
Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.04166666666666664
F1: 0.07692307692307693
FPR: 0.008130081300813009

--- Threshold = 1.9999999999999751 ---
Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.04166666666666664
F1: 0.07692307692307693
FPR: 0.008130081300813009

--- Threshold = 2.049999999999975 ---
Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.04166666666666664
F1: 0.07692307692307693
FPR: 0.008130081300813009

--- Threshold = 2.0999999999999748 ---
Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.04166666666666664
F1: 0.07692307692307693
FPR: 0.008130081300813009

--- Threshold = 2.1499999999999746 ---
Accuracy: 83.6734693877551%
Precision: 0.5

Recall: 0.04166666666666664
F1: 0.07692307692307693
FPR: 0.008130081300813009

--- Threshold = 2.1999999999999744 ---

C:\Users\chyut\AppData\Local\Temp\ipykernel_24836\297337688.py:13:
RuntimeWarning: divide by zero encountered in scalar divide
F1 = 2 / (1 / precision + 1 / recall)

Accuracy: 82.99319727891157%
Precision: 0.0
Recall: 0.0
F1: 0.0
FPR: 0.008130081300813009

--- Threshold = 2.2499999999999742 ---

Accuracy: 82.99319727891157%
Precision: 0.0
Recall: 0.0
F1: 0.0
FPR: 0.008130081300813009

--- Threshold = 2.299999999999974 ---

Accuracy: 82.99319727891157%
Precision: 0.0
Recall: 0.0
F1: 0.0
FPR: 0.008130081300813009

--- Threshold = 2.349999999999974 ---

Accuracy: 82.99319727891157%
Precision: 0.0
Recall: 0.0
F1: 0.0
FPR: 0.008130081300813009

--- Threshold = 2.3999999999999737 ---

Accuracy: 82.99319727891157%
Precision: 0.0
Recall: 0.0
F1: 0.0
FPR: 0.008130081300813009

--- Threshold = 2.4499999999999735 ---

Accuracy: 82.99319727891157%
Precision: 0.0
Recall: 0.0
F1: 0.0

FPR: 0.008130081300813009

--- Threshold = 2.4999999999999734 ---

Accuracy: 82.99319727891157%

Precision: 0.0

Recall: 0.0

F1: 0.0

FPR: 0.008130081300813009

--- Threshold = 2.549999999999973 ---

Accuracy: 82.99319727891157%

Precision: 0.0

Recall: 0.0

F1: 0.0

FPR: 0.008130081300813009

--- Threshold = 2.599999999999973 ---

Accuracy: 82.99319727891157%

Precision: 0.0

Recall: 0.0

F1: 0.0

FPR: 0.008130081300813009

--- Threshold = 2.649999999999973 ---

C:\Users\chyut\AppData\Local\Temp\ipykernel_24836\297337688.py:11:

RuntimeWarning: invalid value encountered in scalar divide

precision = (tp) / (tp + fp)

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 2.6999999999999726 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 2.7499999999999725 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 2.7999999999999723 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.849999999999972 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.899999999999972 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.9499999999999718 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.9999999999999716 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 3.0499999999999723 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 3.0999999999999712 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 3.14999999999997 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.199999999999971 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.2499999999999716 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.2999999999999705 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.3499999999999694 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.39999999999997 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.449999999999971 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.49999999999997 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.5499999999999687 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.5999999999999694 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.64999999999997 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.699999999999969 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.749999999999968 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.7999999999999687 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan
FPR: 0.0

--- Threshold = 3.8499999999999694 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 3.8999999999999684 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 3.9499999999999673 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 3.999999999999968 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.049999999999969 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.099999999999968 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.149999999999967 ---
Accuracy: 83.6734693877551%
Precision: nan

Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.199999999999967 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.249999999999968 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.299999999999967 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.349999999999966 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.399999999999967 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.449999999999967 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.499999999999966 ---
Accuracy: 83.6734693877551%

Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.549999999999965 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.599999999999966 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.649999999999967 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.699999999999965 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.749999999999965 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.799999999999965 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.849999999999966 ---

Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0


--- Threshold = 4.899999999999965 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.949999999999964 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

Best Accuracy: 0.8503401360544217 with responding threshold: 0.8499999999999792
Best F1: 0.5625 with responding threshold: -0.35000000000000165

1.0.19 T17. Plot the RoC of your classifier.

```
[47]: plt.plot(history["FPR"], history["Recall"])  
plt.plot([0, 1], [0, 1], color='gray', linestyle='--', linewidth=2)  
plt.xlabel("False Positive Rate")  
plt.ylabel("True Positive Rate")  
plt.title("RoC")  
plt.show()
```



1.0.20 T18. Change the number of discretization bins to 5. What happens to the RoC curve? Which discretization is better? The number of discretization bins can be considered as a hyperparameter, and must be chosen by comparing the final performance.

```
[48]: df = pd.read_csv('hr-employee-attrition-with-null.csv')
df.loc[df["Attrition"] == "no", "Attrition"] = 0.0
df.loc[df["Attrition"] == "yes", "Attrition"] = 1.0
string_categorical_col = ['Department', 'Attrition', 'BusinessTravel', '
↳ 'EducationField', 'Gender', 'JobRole',
                           'MaritalStatus', 'Over18', 'OverTime']

# ENCODE STRING COLUMNS TO CATEGORICAL COLUMNS
for col in string_categorical_col:
    # INSERT CODE HERE
    df[col] = pd.Categorical(df[col]).codes
# HANDLE NULL NUMBERS
# I don't think we need to handle null?
```

```
# INSERT CODE HERE
df = df.loc[:, ~df.columns.isin(['EmployeeNumber', 'Unnamed: 0',
↳'EmployeeCount', 'StandardHours', 'Over18'])] # drop these columns
X = df.drop(["Attrition"], axis=1)
Y = df["Attrition"]

x_train, x_test, y_train, y_test = train_test_split(X, Y, stratify=Y,
↳test_size=0.1, random_state=12345)
```

```
[49]: def hist(array, col_name, n_bin=10):
        nonan = array[~np.isnan(array)]

        # hist is the count for each bin
        # bin_edge is the edge values of the bins
        hist, bin_edges = np.histogram(nonan, n_bin)
        bin_edges[0] = -np.inf
        bin_edges[-1] = np.inf

        bin_indices = np.full_like(array, np.nan, dtype=float)
        bin_indices[~np.isnan(array)] = np.digitize(nonan, bin_edges)

        return bin_indices, bin_edges

discretize = []

for col in x_train.columns:
    if (x_train[col].nunique() > 10):
        x_train[col], _ = hist(x_train[col], col, 5)
        discretize.append(col)

print(discretize)
```

```
['Age', 'DailyRate', 'DistanceFromHome', 'HourlyRate', 'MonthlyIncome',
'MonthlyRate', 'PercentSalaryHike', 'TotalWorkingYears', 'YearsAtCompany',
'YearsInCurrentRole', 'YearsSinceLastPromotion', 'YearsWithCurrManager']
```

```
[50]: import importlib, SimpleBayesClassifier
importlib.reload(SimpleBayesClassifier)
from SimpleBayesClassifier import SimpleBayesClassifier

model = SimpleBayesClassifier(n_pos = np.sum(y_train == 1), n_neg = np.
↳sum(y_train == 0))

model.fit_params(np.array(x_train), np.array(y_train))
```

```
[50]: ([array([0.08080808, 0.          , 0.31425365, 0.          , 0.          ,
0.34680135, 0.          , 0.17059484, 0.          , 0.08754209]),
```

```

array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.2027027 , 0.          , 0.          , 0.08558559, 0.          ,
        0.          , 0.13153153, 0.          , 0.          , 0.58018018])),
array([-inf, -0.7, -0.4, -0.1, 0.2, 0.5, 0.8, 1.1, 1.4, 1.7, inf])),
(array([0.1901566 , 0.          , 0.2114094 , 0.          , 0.          ,
        0.18120805, 0.          , 0.20917226, 0.          , 0.20805369])),
array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.19099099, 0.          , 0.          , 0.03603604, 0.          ,
        0.          , 0.54594595, 0.          , 0.          , 0.22702703])),
array([-inf, -0.7, -0.4, -0.1, 0.2, 0.5, 0.8, 1.1, 1.4, 1.7, inf])),
(array([0.4778157 , 0.          , 0.26393629, 0.          , 0.          ,
        0.0705347 , 0.          , 0.09215017, 0.          , 0.09556314])),
array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.10946408, 0.          , 0.19156214, 0.          , 0.          ,
        0.38312429, 0.          , 0.28164196, 0.          , 0.03420753])),
array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.19459459, 0.01531532, 0.          , 0.33693694, 0.          ,
        0.07567568, 0.26666667, 0.          , 0.04234234, 0.06846847])),
array([-inf, -0.4, 0.2, 0.8, 1.4, 2. , 2.6, 3.2, 3.8, 4.4, inf])),
(array([0.18423973, 0.          , 0.          , 0.18201998, 0.          ,
        0.          , 0.32741398, 0.          , 0.          , 0.3063263 ])),
array([-inf, 1.3, 1.6, 1.9, 2.2, 2.5, 2.8, 3.1, 3.4, 3.7, inf])),
(array([0.2009009 , 0.          , 0.          , 0.          , 0.          ,
        0.33873874, 0.          , 0.          , 0.          , 0.46036036])),
array([-inf, -0.8, -0.6, -0.4, -0.2, 0. , 0.2, 0.4, 0.6, 0.8, inf])),
(array([0.18946188, 0.          , 0.20515695, 0.          , 0.          ,
        0.18161435, 0.          , 0.20403587, 0.          , 0.21973094])),
array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.0407701 , 0.          , 0.          , 0.26274066, 0.          ,
        0.          , 0.59116648, 0.          , 0.          , 0.10532276])),
array([-inf, 1.3, 1.6, 1.9, 2.2, 2.5, 2.8, 3.1, 3.4, 3.7, inf])),
(array([0.31513083, 0.          , 0.37997725, 0.          , 0.          ,
        0.16382253, 0.          , 0.09215017, 0.          , 0.04891923])),
array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.1990991 , 0.08288288, 0.02702703, 0.11891892, 0.06396396,
        0.08918919, 0.05225225, 0.15495495, 0.17477477, 0.03693694])),
array([-inf, -0.1, 0.8, 1.7, 2.6, 3.5, 4.4, 5.3, 6.2, 7.1, inf])),
(array([0.18459796, 0.          , 0.          , 0.19705549, 0.          ,
        0.          , 0.29331823, 0.          , 0.          , 0.32502831])),
array([-inf, 1.3, 1.6, 1.9, 2.2, 2.5, 2.8, 3.1, 3.4, 3.7, inf])),
(array([0.19099099, 0.          , 0.          , 0.1954955 , 0.          ,
        0.          , 0.38108108, 0.          , 0.          , 0.23243243])),
array([-inf, -0.7, -0.4, -0.1, 0.2, 0.5, 0.8, 1.1, 1.4, 1.7, inf])),
(array([0.44843049, 0.          , 0.28923767, 0.          , 0.          ,
        0.117713 , 0.          , 0.05044843, 0.          , 0.0941704 ])),
array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.21235955, 0.          , 0.2011236 , 0.          , 0.          ,

```

```

        0.19325843, 0.          , 0.2011236 , 0.          , 0.19213483]],
    array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
    (array([0.13718821, 0.33446712, 0.10544218, 0.12131519, 0.1031746 ,
        0.03741497, 0.04421769, 0.04875283, 0.03061224, 0.03741497])),
    array([-inf, 0.9, 1.8, 2.7, 3.6, 4.5, 5.4, 6.3, 7.2, 8.1, inf])),
    (array([0.19189189, 0.          , 0.          , 0.          , 0.          ,
        0.62162162, 0.          , 0.          , 0.          , 0.18648649])),
    array([-inf, -0.8, -0.6, -0.4, -0.2, 0. , 0.2, 0.4, 0.6, 0.8, inf])),
    (array([0.40384615, 0.          , 0.27036199, 0.          , 0.          ,
        0.16968326, 0.          , 0.10972851, 0.          , 0.04638009])),
    array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
    (array([0.85310734, 0.          , 0.          , 0.          , 0.          ,
        0.          , 0.          , 0.          , 0.          , 0.14689266])),
    array([-inf, 3.1, 3.2, 3.3, 3.4, 3.5, 3.6, 3.7, 3.8, 3.9, inf])),
    (array([0.1868743 , 0.          , 0.          , 0.20244716, 0.          ,
        0.          , 0.30700779, 0.          , 0.          , 0.30367075])),
    array([-inf, 1.3, 1.6, 1.9, 2.2, 2.5, 2.8, 3.1, 3.4, 3.7, inf])),
    (array([0.39595051, 0.          , 0.          , 0.44769404, 0.          ,
        0.          , 0.10686164, 0.          , 0.          , 0.04949381])),
    array([-inf, 0.3, 0.6, 0.9, 1.2, 1.5, 1.8, 2.1, 2.4, 2.7, inf])),
    (array([0.31291759, 0.          , 0.42538976, 0.          , 0.          ,
        0.16258352, 0.          , 0.06904232, 0.          , 0.03006682])),
    array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
    (array([0.03703704, 0.05274972, 0.          , 0.36026936, 0.          ,
        0.332211 , 0.07856341, 0.          , 0.09539843, 0.04377104])),
    array([-inf, 0.6, 1.2, 1.8, 2.4, 3. , 3.6, 4.2, 4.8, 5.4, inf])),
    (array([0.04519774, 0.          , 0.          , 0.23050847, 0.          ,
        0.          , 0.62146893, 0.          , 0.          , 0.10282486])),
    array([-inf, 1.3, 1.6, 1.9, 2.2, 2.5, 2.8, 3.1, 3.4, 3.7, inf])),
    (array([0.61167228, 0.          , 0.26823793, 0.          , 0.          ,
        0.08080808, 0.          , 0.02581369, 0.          , 0.01346801])),
    array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
    (array([0.51727982, 0.          , 0.28205128, 0.          , 0.          ,
        0.14381271, 0.          , 0.04236343, 0.          , 0.01449275])),
    array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
    (array([0.73318386, 0.          , 0.1132287 , 0.          , 0.          ,
        0.09304933, 0.          , 0.03363229, 0.          , 0.02690583])),
    array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
    (array([0.53454134, 0.          , 0.10532276, 0.          , 0.          ,
        0.30577576, 0.          , 0.0407701 , 0.          , 0.01359003])),
    array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf]]),
    [(array([0.24855491, 0.          , 0.36416185, 0.          , 0.          ,
        0.1849711 , 0.          , 0.11560694, 0.          , 0.0867052 ])),
    array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
    (array([0.19248826, 0.          , 0.          , 0.03286385, 0.          ,
        0.          , 0.23004695, 0.          , 0.          , 0.54460094])),
    array([-inf, -0.7, -0.4, -0.1, 0.2, 0.5, 0.8, 1.1, 1.4, 1.7, inf])),

```

```

(array([0.23566879, 0.          , 0.20382166, 0.          , 0.          ,
        0.19745223, 0.          , 0.1910828 , 0.          , 0.17197452]),
 array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.22535211, 0.          , 0.          , 0.05164319, 0.          ,
        0.          , 0.41314554, 0.          , 0.          , 0.30985915]),
 array([-inf, -0.7, -0.4, -0.1, 0.2, 0.5, 0.8, 1.1, 1.4, 1.7, inf])),
(array([0.37853107, 0.          , 0.23728814, 0.          , 0.          ,
        0.11299435, 0.          , 0.11299435, 0.          , 0.15819209]),
 array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.13068182, 0.          , 0.19886364, 0.          , 0.          ,
        0.41477273, 0.          , 0.22727273, 0.          , 0.02840909]),
 array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.18309859, 0.02816901, 0.          , 0.29577465, 0.          ,
        0.12676056, 0.22535211, 0.          , 0.03755869, 0.10328638]),
 array([-inf, -0.4, 0.2, 0.8, 1.4, 2. , 2.6, 3.2, 3.8, 4.4, inf])),
(array([0.26875, 0.          , 0.          , 0.2          , 0.          , 0.          ,
        0.          , 0.          , 0.29375])),
 array([-inf, 1.3, 1.6, 1.9, 2.2, 2.5, 2.8, 3.1, 3.4, 3.7, inf])),
(array([0.19248826, 0.          , 0.          , 0.          , 0.          ,
        0.29577465, 0.          , 0.          , 0.          , 0.51173709]),
 array([-inf, -0.8, -0.6, -0.4, -0.2, 0. , 0.2, 0.4, 0.6, 0.8, inf])),
(array([0.19526627, 0.          , 0.23668639, 0.          , 0.          ,
        0.20118343, 0.          , 0.15976331, 0.          , 0.20710059]),
 array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.1091954 , 0.          , 0.          , 0.31034483, 0.          ,
        0.          , 0.52873563, 0.          , 0.          , 0.05172414]),
 array([-inf, 1.3, 1.6, 1.9, 2.2, 2.5, 2.8, 3.1, 3.4, 3.7, inf])),
(array([0.60795455, 0.          , 0.23295455, 0.          , 0.          ,
        0.11931818, 0.          , 0.01704545, 0.          , 0.02272727]),
 array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.18779343, 0.02816901, 0.04694836, 0.19248826, 0.01877934,
        0.03755869, 0.00469484, 0.17370892, 0.18779343, 0.12206573]),
 array([-inf, -0.1, 0.8, 1.7, 2.6, 3.5, 4.4, 5.3, 6.2, 7.1, inf])),
(array([0.24          , 0.          , 0.          , 0.2          , 0.          ,
        0.          , 0.30857143, 0.          , 0.          , 0.25142857]),
 array([-inf, 1.3, 1.6, 1.9, 2.2, 2.5, 2.8, 3.1, 3.4, 3.7, inf])),
(array([0.22065728, 0.          , 0.          , 0.11267606, 0.          ,
        0.          , 0.2629108 , 0.          , 0.          , 0.40375587]),
 array([-inf, -0.7, -0.4, -0.1, 0.2, 0.5, 0.8, 1.1, 1.4, 1.7, inf])),
(array([0.68862275, 0.          , 0.17365269, 0.          , 0.          ,
        0.10179641, 0.          , 0.01796407, 0.          , 0.01796407]),
 array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.2          , 0.          , 0.22424242, 0.          , 0.          ,
        0.17575758, 0.          , 0.22424242, 0.          , 0.17575758]),
 array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.0960452 , 0.38983051, 0.07909605, 0.06779661, 0.08474576,
        0.06779661, 0.05649718, 0.06779661, 0.02824859, 0.06214689]),

```

```

array([-inf, 0.9, 1.8, 2.7, 3.6, 4.5, 5.4, 6.3, 7.2, 8.1, inf])),
(array([0.23474178, 0.          , 0.          , 0.          , 0.          ,
        0.36150235, 0.          , 0.          , 0.          , 0.40375587])),
array([-inf, -0.8, -0.6, -0.4, -0.2, 0. , 0.2, 0.4, 0.6, 0.8, inf])),
(array([0.49162011, 0.          , 0.2122905 , 0.          , 0.          ,
        0.1452514 , 0.          , 0.08379888, 0.          , 0.06703911])),
array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.83529412, 0.          , 0.          , 0.          , 0.          ,
        0.          , 0.          , 0.          , 0.          , 0.16470588])),
array([-inf, 3.1, 3.2, 3.3, 3.4, 3.5, 3.6, 3.7, 3.8, 3.9, inf])),
(array([0.26219512, 0.          , 0.          , 0.20121951, 0.          ,
        0.          , 0.2804878 , 0.          , 0.          , 0.25609756])),
array([-inf, 1.3, 1.6, 1.9, 2.2, 2.5, 2.8, 3.1, 3.4, 3.7, inf])),
(array([0.65116279, 0.          , 0.          , 0.22674419, 0.          ,
        0.          , 0.06395349, 0.          , 0.          , 0.05813953])),
array([-inf, 0.3, 0.6, 0.9, 1.2, 1.5, 1.8, 2.1, 2.4, 2.7, inf])),
(array([0.55151515, 0.          , 0.3030303 , 0.          , 0.          ,
        0.0969697 , 0.          , 0.03030303, 0.          , 0.01818182])),
array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.05202312, 0.05202312, 0.          , 0.42196532, 0.          ,
        0.28901734, 0.10404624, 0.          , 0.06358382, 0.01734104])),
array([-inf, 0.6, 1.2, 1.8, 2.4, 3. , 3.6, 4.2, 4.8, 5.4, inf])),
(array([0.10588235, 0.          , 0.          , 0.23529412, 0.          ,
        0.          , 0.52352941, 0.          , 0.          , 0.13529412])),
array([-inf, 1.3, 1.6, 1.9, 2.2, 2.5, 2.8, 3.1, 3.4, 3.7, inf])),
(array([0.72674419, 0.          , 0.21511628, 0.          , 0.          ,
        0.03488372, 0.          , 0.01162791, 0.          , 0.01162791])),
array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.68263473, 0.          , 0.20359281, 0.          , 0.          ,
        0.08982036, 0.          , 0.01796407, 0.          , 0.00598802])),
array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.78235294, 0.          , 0.07058824, 0.          , 0.          ,
        0.09411765, 0.          , 0.02941176, 0.          , 0.02352941])),
array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf])),
(array([0.69364162, 0.          , 0.06358382, 0.          , 0.          ,
        0.22543353, 0.          , 0.00578035, 0.          , 0.01156069])),
array([-inf, 1.4, 1.8, 2.2, 2.6, 3. , 3.4, 3.8, 4.2, 4.6, inf]]))

```

```

[51]: hbin5 = {}
hbin5["Accuracy"] = []
hbin5["Precision"] = []
hbin5["Recall"] = []
hbin5["F1"] = []
hbin5["FPR"] = []

for each in t:
    print("--- Threshold =", each, "---")

```

```

y_pred = np.array(model.predict(np.array(x_test), thresh=each))

accuracy, precision, recall, f1, fpr = evaluate(y_test, y_pred)

hbin5["Accuracy"].append(accuracy)
hbin5["Precision"].append(precision)
hbin5["Recall"].append(recall)
hbin5["F1"].append(f1)
hbin5["FPR"].append(fpr)

plt.plot(hbin5["FPR"], hbin5["Recall"])
plt.plot([0, 1], [0, 1], color='gray', linestyle='--', linewidth=2)
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("RoC with 5 bin")
plt.show()

```

--- Threshold = -5.0 ---

Accuracy: 50.34013605442177%
Precision: 0.24210526315789474
Recall: 0.9583333333333334
F1: 0.3865546218487395
FPR: 0.5853658536585366

--- Threshold = -4.95 ---

Accuracy: 50.34013605442177%
Precision: 0.24210526315789474
Recall: 0.9583333333333334
F1: 0.3865546218487395
FPR: 0.5853658536585366

--- Threshold = -4.9 ---

Accuracy: 51.70068027210885%
Precision: 0.24731182795698925
Recall: 0.9583333333333334
F1: 0.39316239316239315
FPR: 0.5691056910569106

--- Threshold = -4.8500000000000005 ---

Accuracy: 51.70068027210885%
Precision: 0.24731182795698925
Recall: 0.9583333333333334
F1: 0.39316239316239315
FPR: 0.5691056910569106

--- Threshold = -4.8000000000000001 ---

Accuracy: 52.38095238095239%
Precision: 0.25

Recall: 0.9583333333333334
F1: 0.396551724137931
FPR: 0.5609756097560976

--- Threshold = -4.7500000000000001 ---
Accuracy: 54.421768707483%
Precision: 0.25842696629213485
Recall: 0.9583333333333334
F1: 0.40707964601769914
FPR: 0.5365853658536586

--- Threshold = -4.7000000000000001 ---
Accuracy: 55.78231292517006%
Precision: 0.26436781609195403
Recall: 0.9583333333333334
F1: 0.4144144144144144
FPR: 0.5203252032520326

--- Threshold = -4.6500000000000001 ---
Accuracy: 56.4625850340136%
Precision: 0.26744186046511625
Recall: 0.9583333333333334
F1: 0.4181818181818182
FPR: 0.5121951219512195

--- Threshold = -4.6000000000000001 ---
Accuracy: 58.50340136054422%
Precision: 0.27710843373493976
Recall: 0.9583333333333334
F1: 0.42990654205607476
FPR: 0.4878048780487805

--- Threshold = -4.5500000000000002 ---
Accuracy: 59.183673469387756%
Precision: 0.2804878048780488
Recall: 0.9583333333333334
F1: 0.4339622641509434
FPR: 0.4796747967479675

--- Threshold = -4.5000000000000002 ---
Accuracy: 59.183673469387756%
Precision: 0.2804878048780488
Recall: 0.9583333333333334
F1: 0.4339622641509434
FPR: 0.4796747967479675

--- Threshold = -4.4500000000000002 ---
Accuracy: 59.863945578231295%

Precision: 0.2839506172839506
Recall: 0.9583333333333334
F1: 0.4380952380952381
FPR: 0.4715447154471545

--- Threshold = -4.4000000000000002 ---
Accuracy: 59.863945578231295%
Precision: 0.2839506172839506
Recall: 0.9583333333333334
F1: 0.4380952380952381
FPR: 0.4715447154471545

--- Threshold = -4.3500000000000002 ---
Accuracy: 60.544217687074834%
Precision: 0.2875
Recall: 0.9583333333333334
F1: 0.44230769230769224
FPR: 0.4634146341463415

--- Threshold = -4.30000000000000025 ---
Accuracy: 61.224489795918366%
Precision: 0.2911392405063291
Recall: 0.9583333333333334
F1: 0.4466019417475728
FPR: 0.45528455284552843

--- Threshold = -4.2500000000000003 ---
Accuracy: 61.224489795918366%
Precision: 0.2857142857142857
Recall: 0.9166666666666666
F1: 0.43564356435643564
FPR: 0.44715447154471544

--- Threshold = -4.2000000000000003 ---
Accuracy: 60.544217687074834%
Precision: 0.27631578947368424
Recall: 0.875
F1: 0.42000000000000001
FPR: 0.44715447154471544

--- Threshold = -4.1500000000000003 ---
Accuracy: 61.224489795918366%
Precision: 0.28
Recall: 0.875
F1: 0.4242424242424243
FPR: 0.43902439024390244

--- Threshold = -4.1000000000000003 ---

Accuracy: 62.585034013605444%
Precision: 0.2876712328767123
Recall: 0.875
F1: 0.4329896907216495
FPR: 0.42276422764227645

--- Threshold = -4.0500000000000003 ---
Accuracy: 62.585034013605444%
Precision: 0.28169014084507044
Recall: 0.8333333333333334
F1: 0.42105263157894735
FPR: 0.4146341463414634

--- Threshold = -4.00000000000000036 ---
Accuracy: 63.94557823129252%
Precision: 0.2898550724637681
Recall: 0.8333333333333334
F1: 0.4301075268817205
FPR: 0.3983739837398374

--- Threshold = -3.95000000000000037 ---
Accuracy: 63.94557823129252%
Precision: 0.2898550724637681
Recall: 0.8333333333333334
F1: 0.4301075268817205
FPR: 0.3983739837398374

--- Threshold = -3.9000000000000004 ---
Accuracy: 64.62585034013605%
Precision: 0.29411764705882354
Recall: 0.8333333333333334
F1: 0.4347826086956522
FPR: 0.3902439024390244

--- Threshold = -3.8500000000000004 ---
Accuracy: 65.3061224489796%
Precision: 0.29850746268656714
Recall: 0.8333333333333334
F1: 0.43956043956043955
FPR: 0.3821138211382114

--- Threshold = -3.80000000000000043 ---
Accuracy: 65.3061224489796%
Precision: 0.2857142857142857
Recall: 0.75
F1: 0.4137931034482759
FPR: 0.36585365853658536

--- Threshold = -3.7500000000000044 ---
Accuracy: 65.98639455782312%
Precision: 0.2903225806451613
Recall: 0.75
F1: 0.4186046511627907
FPR: 0.35772357723577236

--- Threshold = -3.7000000000000046 ---
Accuracy: 68.02721088435374%
Precision: 0.3050847457627119
Recall: 0.75
F1: 0.4337349397590362
FPR: 0.3333333333333333

--- Threshold = -3.6500000000000005 ---
Accuracy: 68.70748299319727%
Precision: 0.3103448275862069
Recall: 0.75
F1: 0.43902439024390244
FPR: 0.3252032520325203

--- Threshold = -3.6000000000000005 ---
Accuracy: 68.70748299319727%
Precision: 0.2962962962962963
Recall: 0.6666666666666666
F1: 0.41025641025641024
FPR: 0.3089430894308943

--- Threshold = -3.5500000000000005 ---
Accuracy: 68.02721088435374%
Precision: 0.2830188679245283
Recall: 0.625
F1: 0.38961038961038963
FPR: 0.3089430894308943

--- Threshold = -3.5000000000000053 ---
Accuracy: 68.02721088435374%
Precision: 0.27450980392156865
Recall: 0.5833333333333334
F1: 0.3733333333333335
FPR: 0.3008130081300813

--- Threshold = -3.4500000000000055 ---
Accuracy: 70.06802721088435%
Precision: 0.2916666666666667
Recall: 0.5833333333333334
F1: 0.38888888888888895
FPR: 0.2764227642276423

--- Threshold = -3.4000000000000057 ---

Accuracy: 70.06802721088435%

Precision: 0.2916666666666667

Recall: 0.5833333333333334

F1: 0.38888888888888895

FPR: 0.2764227642276423

--- Threshold = -3.3500000000000006 ---

Accuracy: 70.06802721088435%

Precision: 0.2916666666666667

Recall: 0.5833333333333334

F1: 0.38888888888888895

FPR: 0.2764227642276423

--- Threshold = -3.3000000000000006 ---

Accuracy: 71.42857142857143%

Precision: 0.30434782608695654

Recall: 0.5833333333333334

F1: 0.4

FPR: 0.2601626016260163

--- Threshold = -3.2500000000000006 ---

Accuracy: 72.10884353741497%

Precision: 0.3111111111111111

Recall: 0.5833333333333334

F1: 0.40579710144927533

FPR: 0.25203252032520324

--- Threshold = -3.20000000000000064 ---

Accuracy: 72.78911564625851%

Precision: 0.3181818181818182

Recall: 0.5833333333333334

F1: 0.411764705882353

FPR: 0.24390243902439024

--- Threshold = -3.15000000000000066 ---

Accuracy: 72.10884353741497%

Precision: 0.3023255813953488

Recall: 0.5416666666666666

F1: 0.38805970149253727

FPR: 0.24390243902439024

--- Threshold = -3.10000000000000068 ---

Accuracy: 72.78911564625851%

Precision: 0.30952380952380953

Recall: 0.5416666666666666

F1: 0.393939393939394

FPR: 0.23577235772357724

--- Threshold = -3.0500000000000007 ---

Accuracy: 72.10884353741497%

Precision: 0.2926829268292683

Recall: 0.5

F1: 0.3692307692307692

FPR: 0.23577235772357724

--- Threshold = -3.0000000000000007 ---

Accuracy: 72.10884353741497%

Precision: 0.2926829268292683

Recall: 0.5

F1: 0.3692307692307692

FPR: 0.23577235772357724

--- Threshold = -2.95000000000000073 ---

Accuracy: 72.10884353741497%

Precision: 0.2926829268292683

Recall: 0.5

F1: 0.3692307692307692

FPR: 0.23577235772357724

--- Threshold = -2.90000000000000075 ---

Accuracy: 72.10884353741497%

Precision: 0.28205128205128205

Recall: 0.4583333333333333

F1: 0.3492063492063492

FPR: 0.22764227642276422

--- Threshold = -2.85000000000000076 ---

Accuracy: 72.10884353741497%

Precision: 0.28205128205128205

Recall: 0.4583333333333333

F1: 0.3492063492063492

FPR: 0.22764227642276422

--- Threshold = -2.8000000000000008 ---

Accuracy: 72.10884353741497%

Precision: 0.28205128205128205

Recall: 0.4583333333333333

F1: 0.3492063492063492

FPR: 0.22764227642276422

--- Threshold = -2.7500000000000008 ---

Accuracy: 73.46938775510205%

Precision: 0.2972972972972973

Recall: 0.4583333333333333

F1: 0.36065573770491804
FPR: 0.21138211382113822

--- Threshold = -2.7000000000000008 ---
Accuracy: 74.14965986394559%
Precision: 0.28125
Recall: 0.375
F1: 0.32142857142857145
FPR: 0.18699186991869918

--- Threshold = -2.6500000000000003 ---
Accuracy: 74.14965986394559%
Precision: 0.28125
Recall: 0.375
F1: 0.32142857142857145
FPR: 0.18699186991869918

--- Threshold = -2.6000000000000005 ---
Accuracy: 74.82993197278913%
Precision: 0.2903225806451613
Recall: 0.375
F1: 0.32727272727272727
FPR: 0.17886178861788618

--- Threshold = -2.5500000000000007 ---
Accuracy: 74.14965986394559%
Precision: 0.26666666666666666
Recall: 0.3333333333333333
F1: 0.2962962962962963
FPR: 0.17886178861788618

--- Threshold = -2.5000000000000009 ---
Accuracy: 75.51020408163265%
Precision: 0.2857142857142857
Recall: 0.3333333333333333
F1: 0.3076923076923077
FPR: 0.16260162601626016

--- Threshold = -2.4500000000000009 ---
Accuracy: 76.87074829931973%
Precision: 0.3076923076923077
Recall: 0.3333333333333333
F1: 0.32
FPR: 0.14634146341463414

--- Threshold = -2.4000000000000002 ---
Accuracy: 76.87074829931973%
Precision: 0.3076923076923077

Recall: 0.3333333333333333
F1: 0.32
FPR: 0.14634146341463414

--- Threshold = -2.3500000000000094 ---
Accuracy: 77.55102040816327%
Precision: 0.32
Recall: 0.3333333333333333
F1: 0.32653061224489793
FPR: 0.13821138211382114

--- Threshold = -2.3000000000000096 ---
Accuracy: 77.55102040816327%
Precision: 0.32
Recall: 0.3333333333333333
F1: 0.32653061224489793
FPR: 0.13821138211382114

--- Threshold = -2.2500000000000098 ---
Accuracy: 77.55102040816327%
Precision: 0.32
Recall: 0.3333333333333333
F1: 0.32653061224489793
FPR: 0.13821138211382114

--- Threshold = -2.200000000000001 ---
Accuracy: 78.2312925170068%
Precision: 0.3333333333333333
Recall: 0.3333333333333333
F1: 0.3333333333333333
FPR: 0.13008130081300814

--- Threshold = -2.150000000000001 ---
Accuracy: 78.91156462585033%
Precision: 0.34782608695652173
Recall: 0.3333333333333333
F1: 0.3404255319148936
FPR: 0.12195121951219512

--- Threshold = -2.1000000000000103 ---
Accuracy: 78.91156462585033%
Precision: 0.34782608695652173
Recall: 0.3333333333333333
F1: 0.3404255319148936
FPR: 0.12195121951219512

--- Threshold = -2.0500000000000105 ---
Accuracy: 78.91156462585033%

Precision: 0.34782608695652173
Recall: 0.3333333333333333
F1: 0.3404255319148936
FPR: 0.12195121951219512

--- Threshold = -2.0000000000000107 ---
Accuracy: 79.59183673469387%
Precision: 0.363636363636365
Recall: 0.3333333333333333
F1: 0.34782608695652173
FPR: 0.11382113821138211

--- Threshold = -1.9500000000000108 ---
Accuracy: 80.95238095238095%
Precision: 0.4
Recall: 0.3333333333333333
F1: 0.363636363636365
FPR: 0.0975609756097561

--- Threshold = -1.900000000000011 ---
Accuracy: 81.63265306122449%
Precision: 0.42105263157894735
Recall: 0.3333333333333333
F1: 0.37209302325581395
FPR: 0.08943089430894309

--- Threshold = -1.8500000000000112 ---
Accuracy: 82.31292517006803%
Precision: 0.4444444444444444
Recall: 0.3333333333333333
F1: 0.38095238095238093
FPR: 0.08130081300813008

--- Threshold = -1.8000000000000114 ---
Accuracy: 82.99319727891157%
Precision: 0.47058823529411764
Recall: 0.3333333333333333
F1: 0.3902439024390244
FPR: 0.07317073170731707

--- Threshold = -1.7500000000000115 ---
Accuracy: 82.99319727891157%
Precision: 0.47058823529411764
Recall: 0.3333333333333333
F1: 0.3902439024390244
FPR: 0.07317073170731707

--- Threshold = -1.7000000000000117 ---

Accuracy: 82.99319727891157%
Precision: 0.47058823529411764
Recall: 0.3333333333333333
F1: 0.3902439024390244
FPR: 0.07317073170731707

--- Threshold = -1.6500000000000012 ---
Accuracy: 84.35374149659864%
Precision: 0.5333333333333333
Recall: 0.3333333333333333
F1: 0.41025641025641024
FPR: 0.056910569105691054

--- Threshold = -1.6000000000000012 ---
Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.2916666666666667
F1: 0.3684210526315789
FPR: 0.056910569105691054

--- Threshold = -1.55000000000000123 ---
Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.2916666666666667
F1: 0.3684210526315789
FPR: 0.056910569105691054

--- Threshold = -1.50000000000000124 ---
Accuracy: 84.35374149659864%
Precision: 0.5384615384615384
Recall: 0.2916666666666667
F1: 0.3783783783783784
FPR: 0.04878048780487805

--- Threshold = -1.45000000000000126 ---
Accuracy: 84.35374149659864%
Precision: 0.5384615384615384
Recall: 0.2916666666666667
F1: 0.3783783783783784
FPR: 0.04878048780487805

--- Threshold = -1.40000000000000128 ---
Accuracy: 84.35374149659864%
Precision: 0.5384615384615384
Recall: 0.2916666666666667
F1: 0.3783783783783784
FPR: 0.04878048780487805

```

--- Threshold = -1.3500000000000013 ---
Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.25
F1: 0.3333333333333333
FPR: 0.04878048780487805

--- Threshold = -1.30000000000000131 ---
Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.25
F1: 0.3333333333333333
FPR: 0.04878048780487805

--- Threshold = -1.25000000000000133 ---
Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.25
F1: 0.3333333333333333
FPR: 0.04878048780487805

--- Threshold = -1.20000000000000135 ---
Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.25
F1: 0.3333333333333333
FPR: 0.04878048780487805

--- Threshold = -1.15000000000000137 ---
Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.2083333333333334
F1: 0.29411764705882354
FPR: 0.04065040650406504

--- Threshold = -1.10000000000000139 ---
Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.2083333333333334
F1: 0.29411764705882354
FPR: 0.04065040650406504

--- Threshold = -1.0500000000000014 ---
Accuracy: 84.35374149659864%
Precision: 0.5714285714285714
Recall: 0.16666666666666666
F1: 0.25806451612903225
FPR: 0.024390243902439025

```

--- Threshold = -1.0000000000000142 ---

Accuracy: 84.35374149659864%

Precision: 0.5714285714285714

Recall: 0.16666666666666666

F1: 0.25806451612903225

FPR: 0.024390243902439025

--- Threshold = -0.9500000000000144 ---

Accuracy: 84.35374149659864%

Precision: 0.5714285714285714

Recall: 0.16666666666666666

F1: 0.25806451612903225

FPR: 0.024390243902439025

--- Threshold = -0.9000000000000146 ---

Accuracy: 84.35374149659864%

Precision: 0.5714285714285714

Recall: 0.16666666666666666

F1: 0.25806451612903225

FPR: 0.024390243902439025

--- Threshold = -0.8500000000000147 ---

Accuracy: 84.35374149659864%

Precision: 0.5714285714285714

Recall: 0.16666666666666666

F1: 0.25806451612903225

FPR: 0.024390243902439025

--- Threshold = -0.8000000000000149 ---

Accuracy: 84.35374149659864%

Precision: 0.5714285714285714

Recall: 0.16666666666666666

F1: 0.25806451612903225

FPR: 0.024390243902439025

--- Threshold = -0.7500000000000151 ---

Accuracy: 84.35374149659864%

Precision: 0.5714285714285714

Recall: 0.16666666666666666

F1: 0.25806451612903225

FPR: 0.024390243902439025

--- Threshold = -0.7000000000000153 ---

Accuracy: 84.35374149659864%

Precision: 0.5714285714285714

Recall: 0.16666666666666666

F1: 0.25806451612903225

FPR: 0.024390243902439025

--- Threshold = -0.65000000000000155 ---

Accuracy: 84.35374149659864%

Precision: 0.5714285714285714

Recall: 0.16666666666666666

F1: 0.25806451612903225

FPR: 0.024390243902439025

--- Threshold = -0.60000000000000156 ---

Accuracy: 85.03401360544217%

Precision: 0.6666666666666666

Recall: 0.16666666666666666

F1: 0.26666666666666666

FPR: 0.016260162601626018

--- Threshold = -0.55000000000000158 ---

Accuracy: 85.03401360544217%

Precision: 0.6666666666666666

Recall: 0.16666666666666666

F1: 0.26666666666666666

FPR: 0.016260162601626018

--- Threshold = -0.5000000000000016 ---

Accuracy: 85.03401360544217%

Precision: 0.6666666666666666

Recall: 0.16666666666666666

F1: 0.26666666666666666

FPR: 0.016260162601626018

--- Threshold = -0.450000000000001616 ---

Accuracy: 85.03401360544217%

Precision: 0.6666666666666666

Recall: 0.16666666666666666

F1: 0.26666666666666666

FPR: 0.016260162601626018

--- Threshold = -0.400000000000001634 ---

Accuracy: 85.03401360544217%

Precision: 0.6666666666666666

Recall: 0.16666666666666666

F1: 0.26666666666666666

FPR: 0.016260162601626018

--- Threshold = -0.35000000000000165 ---

Accuracy: 85.03401360544217%

Precision: 0.6666666666666666

Recall: 0.16666666666666666

F1: 0.2666666666666666
FPR: 0.016260162601626018

--- Threshold = -0.30000000000000167 ---
Accuracy: 85.03401360544217%
Precision: 0.6666666666666666
Recall: 0.1666666666666666
F1: 0.2666666666666666
FPR: 0.016260162601626018

--- Threshold = -0.25000000000000169 ---
Accuracy: 84.35374149659864%
Precision: 0.6
Recall: 0.125
F1: 0.20689655172413796
FPR: 0.016260162601626018

--- Threshold = -0.200000000000001705 ---
Accuracy: 84.35374149659864%
Precision: 0.6
Recall: 0.125
F1: 0.20689655172413796
FPR: 0.016260162601626018

--- Threshold = -0.150000000000001723 ---
Accuracy: 85.03401360544217%
Precision: 0.75
Recall: 0.125
F1: 0.21428571428571427
FPR: 0.008130081300813009

--- Threshold = -0.100000000000001741 ---
Accuracy: 84.35374149659864%
Precision: 0.6666666666666666
Recall: 0.08333333333333333
F1: 0.14814814814814814
FPR: 0.008130081300813009

--- Threshold = -0.0500000000000017586 ---
Accuracy: 84.35374149659864%
Precision: 0.6666666666666666
Recall: 0.08333333333333333
F1: 0.14814814814814814
FPR: 0.008130081300813009

--- Threshold = -1.7763568394002505e-14 ---
Accuracy: 84.35374149659864%
Precision: 0.6666666666666666

Recall: 0.08333333333333333
F1: 0.14814814814814814
FPR: 0.008130081300813009

--- Threshold = 0.04999999999998206 ---
Accuracy: 84.35374149659864%
Precision: 0.6666666666666666
Recall: 0.08333333333333333
F1: 0.14814814814814814
FPR: 0.008130081300813009

--- Threshold = 0.09999999999998188 ---
Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.041666666666666664
F1: 0.07692307692307693
FPR: 0.008130081300813009

--- Threshold = 0.1499999999999817 ---
Accuracy: 83.6734693877551%
Precision: 0.5
Recall: 0.041666666666666664
F1: 0.07692307692307693
FPR: 0.008130081300813009

--- Threshold = 0.19999999999998153 ---

C:\Users\chyut\AppData\Local\Temp\ipykernel_24836\297337688.py:13:
RuntimeWarning: divide by zero encountered in scalar divide
F1 = 2 / (1 / precision + 1 / recall)

Accuracy: 82.99319727891157%
Precision: 0.0
Recall: 0.0
F1: 0.0
FPR: 0.008130081300813009

--- Threshold = 0.24999999999998135 ---
Accuracy: 82.99319727891157%
Precision: 0.0
Recall: 0.0
F1: 0.0
FPR: 0.008130081300813009

--- Threshold = 0.29999999999998117 ---
Accuracy: 82.99319727891157%
Precision: 0.0
Recall: 0.0
F1: 0.0

FPR: 0.008130081300813009

--- Threshold = 0.349999999999981 ---

Accuracy: 82.99319727891157%

Precision: 0.0

Recall: 0.0

F1: 0.0

FPR: 0.008130081300813009

--- Threshold = 0.3999999999999808 ---

Accuracy: 82.99319727891157%

Precision: 0.0

Recall: 0.0

F1: 0.0

FPR: 0.008130081300813009

--- Threshold = 0.44999999999998064 ---

Accuracy: 82.99319727891157%

Precision: 0.0

Recall: 0.0

F1: 0.0

FPR: 0.008130081300813009

--- Threshold = 0.49999999999998046 ---

Accuracy: 82.99319727891157%

Precision: 0.0

Recall: 0.0

F1: 0.0

FPR: 0.008130081300813009

--- Threshold = 0.5499999999999803 ---

Accuracy: 82.99319727891157%

Precision: 0.0

Recall: 0.0

F1: 0.0

FPR: 0.008130081300813009

--- Threshold = 0.5999999999999801 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 0.6499999999999799 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan
FPR: 0.0

--- Threshold = 0.6999999999999797 ---

C:\Users\chyut\AppData\Local\Temp\ipykernel_24836\297337688.py:11:

RuntimeWarning: invalid value encountered in scalar divide

precision = (tp) / (tp + fp)

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 0.7499999999999796 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 0.7999999999999794 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 0.8499999999999792 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 0.899999999999979 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 0.9499999999999789 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 0.9999999999999787 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 1.0499999999999785 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 1.0999999999999783 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 1.1499999999999782 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 1.199999999999978 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 1.2499999999999778 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 1.2999999999999776 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 1.3499999999999774 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 1.3999999999999773 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 1.449999999999977 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 1.499999999999977 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 1.5499999999999767 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 1.5999999999999766 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 1.6499999999999764 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan
FPR: 0.0

--- Threshold = 1.6999999999999762 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 1.749999999999976 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 1.7999999999999758 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 1.8499999999999757 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 1.8999999999999755 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 1.9499999999999753 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 1.9999999999999751 ---
Accuracy: 83.6734693877551%
Precision: nan

Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.049999999999975 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.0999999999999748 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.1499999999999746 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.1999999999999744 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.2499999999999742 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.299999999999974 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.349999999999974 ---
Accuracy: 83.6734693877551%

Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.3999999999999737 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.4499999999999735 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.4999999999999734 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.549999999999973 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.599999999999973 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.649999999999973 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.6999999999999726 ---

Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.7499999999999725 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.7999999999999723 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.849999999999972 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.899999999999972 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.9499999999999718 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 2.9999999999999716 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 3.0499999999999723 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 3.0999999999999712 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 3.14999999999997 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 3.199999999999971 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 3.2499999999999716 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 3.2999999999999705 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 3.3499999999999694 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 3.39999999999997 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.449999999999971 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.49999999999997 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.5499999999999687 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.5999999999999694 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.64999999999997 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.699999999999969 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.749999999999968 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.7999999999999687 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.8499999999999694 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.8999999999999684 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.9499999999999673 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 3.999999999999968 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 4.049999999999969 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan
FPR: 0.0

--- Threshold = 4.099999999999968 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.149999999999967 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.199999999999967 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.249999999999968 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.299999999999967 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.349999999999966 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.399999999999967 ---
Accuracy: 83.6734693877551%
Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 4.449999999999967 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 4.499999999999966 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 4.549999999999965 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 4.599999999999966 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 4.649999999999967 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 4.699999999999965 ---

Accuracy: 83.6734693877551%

Precision: nan

Recall: 0.0

F1: nan

FPR: 0.0

--- Threshold = 4.7499999999999645 ---

Accuracy: 83.6734693877551%


Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.799999999999965 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.849999999999966 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.899999999999965 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0

--- Threshold = 4.949999999999964 ---
Accuracy: 83.6734693877551%
Precision: nan
Recall: 0.0
F1: nan
FPR: 0.0



1.0.21 OT4

```
[52]: df = pd.read_csv('hr-employee-attrition-with-null.csv')
df.loc[df["Attrition"] == "no", "Attrition"] = 0.0
df.loc[df["Attrition"] == "yes", "Attrition"] = 1.0
string_categorical_col = ['Department', 'Attrition', 'BusinessTravel',
    ↪ 'EducationField', 'Gender', 'JobRole',
    ↪ 'MaritalStatus', 'Over18', 'OverTime']

for col in string_categorical_col:
    df[col] = pd.Categorical(df[col]).codes

df = df.loc[:, ~df.columns.isin(['EmployeeNumber', 'Unnamed: 0',
    ↪ 'EmployeeCount', 'StandardHours', 'Over18'])] # drop these columns
X = df.drop(["Attrition"], axis=1)
Y = df["Attrition"]

accuracies = []
```

```

for i in range(10):
    x_train, x_test, y_train, y_test = train_test_split(X, Y, stratify=Y,
↳test_size=0.1)

    def hist(array, col_name, n_bin=10):
        nonan = array[~np.isnan(array)]

        # hist is the count for each bin
        # bin_edge is the edge values of the bins
        hist, bin_edges = np.histogram(nonan, n_bin)
        bin_edges[0] = -np.inf
        bin_edges[-1] = np.inf

        bin_indices = np.full_like(array, np.nan, dtype=float)
        bin_indices[~np.isnan(array)] = np.digitize(nonan, bin_edges)

        return bin_indices, bin_edges

    discretize = []

    for col in x_train.columns:
        if (x_train[col].nunique() > 10):
            x_train[col], _ = hist(x_train[col], col, 5)
            discretize.append(col)

    import importlib, SimpleBayesClassifier
    importlib.reload(SimpleBayesClassifier)
    from SimpleBayesClassifier import SimpleBayesClassifier

    model = SimpleBayesClassifier(n_pos = np.sum(y_train == 1), n_neg = np.
↳sum(y_train == 0))

    model.fit_gaussian_params(np.array(x_train), np.array(y_train))
    y_pred = model.gaussian_predict(np.array(x_test))

    accuracy, _, _, _ = evaluate(np.array(y_test), y_pred, False)

    accuracies.append(accuracy)

accuracies = np.array(accuracies)
print("Mean:", np.mean(accuracies))
print("Variance:", np.std(accuracies) ** 2)

```

Mean: 0.8251700680272108

Variance: 0.00042621130084686886